

SHOPPAGE v4.1

National Commerce Intelligence & Referral OS

A neutral discovery layer across merchants, marketplaces, social commerce and physical markets

BINDING CATEGORY DECISION

Shoppage is not a cart-led marketplace. It is Zimbabwe's independent commerce search, comparison, referral and merchant-intelligence layer. Existing marketplaces are supply partners and destinations that can list their products on Shoppage alongside retailer websites, informal merchants, social sellers and physical markets.

Confidential operating master | 14 July 2026 | Harare - Sandton

Supersedes Shoppage v3.0 and the rejected cart-led v4.0 draft

Contents - Operating Master

Document status and interpretation

Executive verdict

1. The business Shoppage should build
2. Non-negotiable boundaries
3. Why referral beats another marketplace
4. The 10x strategic thesis
5. Twelve-month target state

Part I - Red-team reconstruction of v3.0

6. What v3.0 got right
7. What v3.0 got wrong
8. Ten defects and their corrections
9. What remains, what changes and what is killed
10. Six mandatory public-launch gates

Part II - Zimbabwe market and ecosystem model

11. Zimbabwe ground truth
12. The market problem
13. Competitive landscape
14. Buyer jobs to be done
15. Merchant jobs to be done
16. Marketplace-partner jobs to be done
17. Ecosystem participant model
18. Launch wedge and coverage cells

Part III - Consumer product and referral experience

19. Product-system overview
20. Local Commerce Search
21. Search result architecture
22. Product comparison page
23. Total acquisition context
24. Merchant Page
25. Marketplace Page
26. Market Page and stall digital twin
27. Request Network
28. Universal Link and QR
29. Referral Ledger
30. Trust Passport
31. Live Trust Seal - the 5th Icon
32. Showcase Shorts
33. Saved Needs and alerts
34. Guides and decision support

35. Language, aliases and input modes

36. Low-data and inclusive design

Part IV - Marketplace federation

37. Why marketplaces should list on Shoppage

38. Partner value proposition

39. Marketplace integration modes

40. Partner Feed Contract

41. Deep-link and destination rules

42. Cross-marketplace ranking neutrality

43. Attribution callbacks

44. Partner analytics

45. Brand and anti-cannibalisation controls

46. No-custody boundary

47. Marketplace pilot design

Part V - Merchant, agency and field operating system

48. Merchant claim workflow

49. Offer-ingestion modes

50. Offer freshness system

51. Business Improvement Score

52. Lead and request inbox

53. Quote and reserve handoff

54. Export and distribution cards

55. Reviews and reputation

56. Agency workspace

57. Agency commercial model

58. Market captains

59. First-party plugins and embeds

Part VI - Commerce Intelligence Graph, AI and technology

60. Core graph model

61. Authoritative-field doctrine

62. Entity resolution

63. Master product catalogue

64. Reference Library

65. Source Rights Register

66. Offer and price history

67. Demand graph

68. AI use cases

69. AI prohibitions

70. Evaluation harness

71. Privacy and data protection

72. Security baseline

73. Recommended technology architecture

74. Independence doctrine

75. Reliability and performance

76. Build-versus-buy doctrine

Part VII - Go-to-market and distribution

77. Coverage-cell launch sequence

78. Anchor-merchant acquisition

79. Marketplace acquisition

80. Buyer acquisition

81. Facebook Groups doctrine

82. WhatsApp doctrine

83. Search and SEO

84. Creator programme

85. Agency channel

86. Market and association partnerships

87. Diaspora distribution

88. Expansion logic

Part VIII - Revenue model and unit economics

89. Revenue principles

90. Launch pricing

91. Merchant subscription economics

92. Cost per click and qualified action

93. Sponsored placements

94. Agency revenue

95. Marketplace and retailer revenue

96. Intelligence products

97. Unit-economics definitions

98. Illustrative month-12 base economics

99. Twenty-four-month scenarios

100. Capital requirement

101. Break-even and kill logic

Part IX - Organisation, governance, compliance and execution

102. Launch team

103. Governance cadence

104. Data operations

105. Trust and safety operations

106. Legal and compliance workstreams

107. Principal risks

108. First 30 days

109. Days 31-90

110. Months 4-6

- 111. Months 7-12
- 112. Months 13-24
- 113. KPI tree
- 114. Investor proof package
- 115. Final verdict
- Part X - Defensibility, network effects and scale
 - 116. Moat doctrine
 - 117. Rights-cleared Commerce Intelligence Graph
 - 118. Zimbabwean aliases and commercial language
 - 119. Compatibility graph
 - 120. Offer-freshness history
 - 121. Demand and zero-result graph
 - 122. Referral-outcome graph
 - 123. Physical-market topology
 - 124. Marketplace federation network
 - 125. Agency maintenance network
 - 126. QR and fifth-icon distribution infrastructure
 - 127. Trust and correction history
 - 128. Data-quality operating system
 - 129. Merchant workflow embedding
 - 130. Independent identity and canonical links
 - 131. Price and availability time series
 - 132. Request liquidity
 - 133. Partner feed infrastructure
 - 134. Features that are not moats
 - 135. Moat scorecard
 - 136. Defence against bulk extraction
 - 137. Regional replication
- Part XI - Detailed product and operating rules
 - 138. Search ranking feature stack
 - 139. Sponsored and organic separation
 - 140. Offer lifecycle
 - 141. Source-confidence model
 - 142. Zero-result workflow
 - 143. Request broadcast controls
 - 144. Duplicate and impersonation controls
 - 145. Conflicting marketplace and merchant data
 - 146. Merchant inactivity and offboarding
 - 147. Marketplace-partner offboarding
 - 148. Complaint triage
 - 149. Correction and suppression process

- 150. Price anomaly controls
- 151. Creator campaign workflow
- 152. Agency service-level agreement
- 153. Market-captain compensation model
- 154. Marketplace partner economics
- 155. Merchant value proof
- 156. Buyer repeat loops
- 157. Weekly operating review
- Part XII - Commercialisation and launch playbooks
- 158. First partner archetypes
- 159. Merchant acquisition offer
- 160. Marketplace-partner pitch
- 161. Agency pitch
- 162. Market-operator pitch
- 163. Buyer launch campaign
- 164. Content operating calendar
- 165. SEO launch architecture
- 166. Diaspora pilot
- 167. Pricing experiments
- 168. Channel attribution standards
- 169. Commercial pipeline stages
- 170. Onboarding-capacity model
- 171. Staffing-capacity assumptions
- 172. Cash-control rules
- 173. Vendor selection
- 174. Data quality sampling
- 175. Launch-day war room
- 176. Post-launch triage
- 177. Metric definitions that prevent self-deception
- 178. Investment narrative
- 179. Board decision calendar
- 180. Detailed 24-month roadmap

Appendices

- Appendix A - Binding decision register
- Appendix B - Minimum product and offer schema
- Appendix C - Freshness matrix
- Appendix D - Source Rights Register template
- Appendix E - Merchant onboarding SOP
- Appendix F - Marketplace partner pilot proposal
- Appendix G - Agency operating pack
- Appendix H - Market captain SOP

[Appendix I - Trust evidence dictionary](#)

[Appendix J - Attribution event taxonomy](#)

[Appendix K - Partner feed fields](#)

[Appendix L - Financial assumptions register](#)

[Appendix M - Extended risk register](#)

[Appendix N - Experiment backlog](#)

[Appendix O - Public launch scorecard](#)

[Appendix P - Competitive positioning matrix](#)

[Appendix Q - Neutral ranking constitution](#)

[Q1. Organic ranking objective](#)

[Q2. Ranking inputs](#)

[Q3. Destination neutrality](#)

[Q4. Sponsored placement](#)

[Q5. Ranking governance](#)

[Q6. Anti-gaming controls](#)

[Q7. Ranking audit scorecard](#)

[Appendix R - Marketplace and retailer listing agreement term sheet](#)

[R1. Service scope](#)

[R2. Responsibilities and transaction boundary](#)

[R3. Data rights schedule](#)

[R4. Feed service level](#)

[R5. Brand and seller presentation](#)

[R6. Attribution and reporting](#)

[R7. Commercial structures](#)

[R8. Competition, neutrality and non-exclusivity](#)

[R9. Consumer complaints and corrections](#)

[R10. Termination and continuity](#)

[Appendix S - Commercial sensitivity, capital discipline and stop rules](#)

[S1. Revenue-engine sensitivities](#)

[S2. Contribution model](#)

[S3. Illustrative operating cases](#)

[S4. Capital allocation order](#)

[S5. Hiring discipline](#)

[S6. Stop, pivot and double-down rules](#)

[S7. Financing posture](#)

[S8. Board financial pack](#)

[Appendix T - Source and evidence register](#)

[Top](#) | [Bottom](#) | [TOC](#)

Document status and interpretation [Back to TOC](#)

This document corrects the strategic error in the earlier v4.0 draft. It preserves the core thesis of Shoppage v3.0 - referral beats marketplace - and rebuilds the business around that decision. The goal is not to create another seller marketplace, warehouse, wallet, courier network or checkout stack. The goal is to become the national layer that helps buyers discover, compare and reach the best available seller or destination, while giving merchants, marketplaces, agencies and market operators a common intelligence and distribution system.

Evidence status. Statements marked Decision are binding operating choices. Targets are management objectives and are not forecasts. Hypotheses require field testing. External market facts are sourced in Appendix T and should be refreshed before investor circulation. Internal asset counts remain subject to rights, duplication and quality audits.

The attached Shoppage v3.0 file is 37 actual PDF pages and approximately 14,900 words, despite a table of contents that describes more than 80 pages. Sections 12 through 23 are substantially repeated rather than being distinct operating chapters. This v4.1 document replaces that artificial length with real operating design, economics, partner architecture, controls and execution gates. [S01]

ONE-SENTENCE MODEL

Shoppage aggregates rights-cleared product, price, stock, location, trust and demand signals; ranks fresh local offers; then routes the buyer to the merchant, marketplace, WhatsApp conversation, telephone, physical stall or partner checkout that can complete the transaction.

Executive verdict [Back to TOC](#)

1. The business Shoppage should build [Back to TOC](#)

Shoppage should build a Zimbabwe-first commerce intelligence and referral network. Its consumer surface is local commerce search and comparison. Its merchant surface is a lightweight shop and lead operating system. Its partner surface is a federated listing and attribution layer for marketplaces, retailer websites, agencies, creators and physical markets. Its data asset is a rights-governed Commerce Intelligence Graph linking products, variants, offers, merchants, branches, marketplaces, places, requests, trust evidence and outcomes.

The system does not need to defeat Ownai, Classifieds, Shop@Zim, Gateway Stream or any future marketplace. It should make them searchable and comparable, send them better-intent traffic, and let them buy intelligence or distribution services. The larger and more fragmented the marketplace landscape becomes, the more valuable the neutral aggregation layer becomes. Current Zimbabwe platforms already demonstrate multiple forms of online shopping and marketplace activity; fragmentation is therefore an input to the thesis, not proof that the category is closed. [S07-S10]

Shoppage owns the canonical product and place graph, search relevance, freshness standards, attribution IDs, request matching and trust evidence. The destination owns price acceptance, contractual sale, payment, fulfilment, returns and transaction liability unless a future separately governed service is launched. This clean boundary preserves speed, capital efficiency and neutrality.

2. Non-negotiable boundaries [Back to TOC](#)

- No Shoppage cart, wallet, stored value, settlement account, warehouse, owned fleet or national delivery promise in the core model.
- No claim that Shoppage sold, delivered, guaranteed or refunded an item when the destination merchant or marketplace did so.
- No unauthorised scraping of Facebook groups, marketplaces, merchant websites, product images or price feeds.
- No default search experience dominated by reference-only products or unclaimed business profiles.

- No pay-to-win organic ranking. Sponsored results are separately labelled and cannot manufacture trust or freshness.
- No field-verification claim unless the evidence predicate is defined, stored, timestamped and auditable.
- No public plugin marketplace, unrestricted bulk export or open intelligence API until extraction, privacy and competition controls pass review.
- No category or city expansion until a coverage cell resolves real demand better than WhatsApp and Facebook alternatives.

3. Why referral beats another marketplace [Back to TOC](#)

A cart-led marketplace would immediately inherit payment integration, reconciliation, disputes, returns, seller settlement, delivery failures, fraud reserves, customer support and consumer-law exposure. It would also need to persuade merchants to maintain another stock ledger and buyers to trust another checkout. Shoppage can create value earlier by organising supply and intent without taking custody of the transaction.

Referral is not a weak directory model when four conditions are met: the search result is materially better than browsing social posts; offers are fresh and branch-linked; the referral is attributable; and Shoppage continues to capture the outcome state without pretending a click equals a sale. The product therefore needs a Referral Ledger and Request Network, not a shopping cart.

The global comparison-shopping category validates the basic pattern. PriceRunner presents millions of products from thousands of stores; PriceCheck positions itself as a comparison service across retailers; Google supports Comparison Shopping Services and programmatic merchant product feeds. These precedents do not prove Zimbabwe economics, but they prove that a neutral discovery layer can coexist with and distribute to multiple stores and marketplaces. [S11-S15]

4. The 10x strategic thesis [Back to TOC](#)

1. Federate all destinations. A product page may contain offers from an independent shop, a physical market stall, a retailer website, a classifieds listing and a marketplace listing. Destination type is explicit.
2. Own the product and demand graph, not the transaction. Canonical identities, compatibility, local aliases, price history, freshness and unresolved requests become the durable asset.
3. Make fresh local offers the default. Reference products support search understanding and catalogue enrichment but never imply local availability.
4. Turn every click into a measurable commercial event. Universal Links preserve source, campaign, product, merchant, branch and partner identity across WhatsApp, web, SMS and QR.
5. Make marketplaces customers. Give them qualified traffic, feed diagnostics, category demand, missing-product signals, price positioning and embeddable intelligence.
6. Make agencies the field distribution layer. Agencies maintain clients, sell managed freshness and intelligence, and earn recurring margin instead of merely posting content.
7. Make physical markets searchable units. Market pages and stall twins capture local topology that pure online platforms cannot reproduce cheaply.
8. Create a Request Network for products that are absent, stale, complex or quote-based. Demand enters the graph even when supply does not yet exist.
9. Enforce evidence-weighted trust. Verification is predicate-specific and expires; response and freshness history matter more than a decorative badge.
10. Expand only through density. Shoppage becomes useful cell by cell, category by category, not by publishing a national shell with no current offers.

5. Twelve-month target state [Back to TOC](#)

At month 12, Shoppage should be a working Harare-first commerce search system with at least three defensible coverage cells, not a national vanity directory. A target cell is useful when the top demand entities have multiple current offers, merchants respond, buyers reach a destination, and unresolved requests generate measurable supply activation.

Measure	Month-12 base target	Why it matters
Active merchants and destination partners	2,500 merchants; 15 marketplace/retailer partners	Supply diversity without requiring seller migration
Fresh branch-linked offers	150,000 offers; 70% refreshed inside category SLA	Search truth rather than catalogue theatre
Monthly intent sessions	150,000	Enough demand to test repeat behaviour and monetisation
Monthly attributable actions	20,000 clicks, calls, WhatsApp starts, quotes, reserves or directions	Commercial value visible to merchants
Resolved Demand Rate	35% of qualified needs	North-star proof that Shoppage solves buying problems
Monthly recurring and repeatable revenue	US\$30,000-US\$40,000	Supports a disciplined Zimbabwe operating team
Owned Demand Share	At least 35% direct, QR, saved need, notification or referral traffic	Reduces dependence on Facebook, Google and TikTok

Part I - Red-team reconstruction of v3.0 [Back to TOC](#)

Keep the category. Remove the noise, repetition, unproven claims and execution sprawl.

6. What v3.0 got right [Back to TOC](#)

The prior model made several high-quality strategic choices that remain binding. It rejected the empty-marketplace cold start; separated reference products from local offers; prioritised Home, Power and Building as a compatibility-heavy wedge; designed for low-data Android use; treated Facebook groups as distribution rather than infrastructure; introduced Merchant Shops, Markets in Markets, Universal Links, the Request Network, Live Trust Seal, agency intelligence and a rights register; and recognised that stock, price and operating claims must be time-bound.

The insight that merchant-specific truth differs from product-master truth is fundamental. Shoppage may know what a Jinko panel model is, but only the merchant or authorised feed can assert that a branch has six units at US\$145 today. The rebuilt architecture preserves that separation in the data model, ranking system and public labels.

7. What v3.0 got wrong [Back to TOC](#)

The largest weakness was not the referral thesis. It was failure to convert the thesis into a disciplined operating system. The document mixed strategy, features, market statistics, global benchmark claims, speculative moats and repeated boilerplate. It named 36-plus moats but did not distinguish a feature, capability, distribution asset, data asset and true switching barrier. It also put agency SaaS, shoppable video, plugins, national intelligence, market captains, one million products and 10,000 claim-ready profiles into one launch story without a resource-constrained sequence.

Several market claims were too precise for the cited evidence, and some global conversion or market-size claims were used as if they proved local willingness to pay. The 200,000-company and 5,200-group files may be valuable, but raw records and membership sums are not active supply or reachable demand. Duplicate members, stale groups, legal rights, false locations and category errors can destroy the apparent scale.

8. Ten defects and their corrections [Back to TOC](#)

Defect	Why it fails	v4.1 correction
Artificial length and repetition	Hides missing operating detail	Every chapter now has a distinct decision, mechanism, metric or control
Five businesses at once	Dilutes capital and accountability	One core: intelligence + discovery + referral; tools serve that core
Feature-as-moat language	Easy to copy	Moats are measured as rights, history, integrations, density and operational learning
Unclear marketplace relationship	Creates unnecessary competition	Marketplaces are first-class destinations and potential paying partners

Defect	Why it fails	v4.1 correction
Weak attribution	Clicks can be worthless	Universal Link + lead events + outcome confidence model
Overloaded preload	Creates liability and bad search	Reference, candidate, corroborated and active-offer layers are visibly separated
Pricing without value proof	US\$19 may still be too high or too low	Fee gates depend on attributable gross-profit evidence
Agency plan ahead of product truth	Agencies can amplify bad data	Certification requires freshness and client outcome thresholds
Plugin marketplace too early	Security and support burden	First-party plugins only until API controls and demand exist
National launch language	Produces thin supply	Cell-based density gates control expansion

9. What remains, what changes and what is killed [Back to TOC](#)

Status	Components	Operating interpretation
Remain	Commerce graph; Merchant Page; Markets in Markets; Universal Link; Request Network; Trust evidence; agency tools; reference library	These are the core system
Change	Showcase Shorts; 5th Icon; plugin architecture; AI recommendations; creator programme	They become controlled distribution or productivity modules, not separate businesses
Defer	Public plugin marketplace; outcome commissions; behavioural personalisation; unrestricted APIs; national data products	Release only after density, attribution and rights gates
Kill	Shoppage checkout, payment custody, owned delivery, warehouses, forced seller migration, unverified national inventory claims	Incompatible with the neutral referral thesis

10. Six mandatory public-launch gates [Back to TOC](#)

1. Source-rights gate: every public dataset and feed has a documented right class, permitted use, attribution rule, retention rule and suppression path.
2. Coverage-density gate: at least 70% of the top 300 demand entities in the launch cell have two current local offers or a working request route.
3. Native-demand gate: in blind tests, users reach a suitable seller faster and with fewer dead ends than through their normal WhatsApp/Facebook process.
4. Attribution gate: at least 90% of instrumented outbound actions retain source, destination and offer identity; duplicate and bot events are excluded.
5. Partner-economics gate: merchants and marketplaces can identify incremental value that exceeds the proposed fee or bid price.
6. Privacy-security gate: independent review confirms lawful data processing, data-controller obligations, access controls, incident response and deletion processes. Zimbabwe's licensing regulations classify controllers by data-subject volume and impose defined obligations; this cannot remain a footnote. [S05-S06]

Part II - Zimbabwe market and ecosystem model [Back to TOC](#)

Design for fragmented commerce, mixed formality, multi-currency pricing and social-first discovery.

11. Zimbabwe ground truth [Back to TOC](#)

The first ZIMSTAT Economic Census identified 204,798 operational establishments and reported that 76.1% were informal. It also reported a very large micro-establishment base. That supports assisted onboarding, simple offer maintenance and place-based discovery; it does not justify publishing every establishment as an active seller. [S02]

The Reserve Bank's payment reports show a substantial and evolving digital-payment system and a multi-currency operating context. Shoppage should display source currency, update time and destination payment methods, but it should not calculate regulatory conclusions or take payment custody from those facts. [S03-S04]

Mobile and social channels are practical commerce infrastructure, but usage, affordability and connectivity vary. The product must load quickly, expose contact and location without compulsory registration, compress media, and preserve actions when a social platform or identity provider is unavailable.

12. The market problem [Back to TOC](#)

Zimbabwean buyers can often find sellers, but they struggle to establish whether the exact item exists, whether the advertised price is current, whether the specification is correct, whether the seller is reachable and where the transaction should happen. The discovery cost is distributed across Facebook groups, WhatsApp statuses, calls, separate marketplaces, Google results, physical visits and word of mouth.

Merchants face the mirror problem. They can publish everywhere but cannot tell which content produced useful demand. A solar shop may repeat the same offer in ten groups, a WhatsApp status, a website and a marketplace, but still have no canonical product identity, no price history, no demand-gap view and no reliable attribution. Shoppage reduces the search and measurement costs on both sides.

13. Competitive landscape [Back to TOC](#)

Ownai operates a seller marketplace; Classifieds describes a marketplace in which the sale contract is directly between buyer and seller; Shop@Zim positions itself as a vendor marketplace; Gateway Stream offers online shopping and delivery. These are not reasons to build another cart. They are evidence that destination supply already exists and can be federated. [S07-S10]

Existing option	Strength	Gap Shoppage targets
Facebook and WhatsApp	Reach, familiarity, rapid posting	Weak structured comparison, freshness and attribution
Classifieds	Broad listings and established audience	Product identity and branch-stock comparability are limited
Cart-led marketplaces	Checkout, seller storefronts, buyer process	Fragmented coverage; each only sees its own supply

Existing option	Strength	Gap Shoppage targets
Retailer websites	Authoritative merchant content and checkout	Single-retailer view; inconsistent product data
Google search	Broad discovery and web indexing	Local informal supply, market-stall topology and stock truth are incomplete
Shoppage	Cross-destination comparison and referral intelligence	Must prove freshness, density and trust before scale

14. Buyer jobs to be done [Back to TOC](#)

- “Show me who actually has this exact product or compatible alternative near me.”
- “Compare the real price, currency, freshness, location, collection and delivery conditions.”
- “I do not know the model number; interpret my photo, voice note, local name or problem.”
- “I need three quotes or a supplier for a bill of quantities, not a generic product page.”
- “I am outside Zimbabwe and need a reliable local destination for my family or project.”
- “Alert me when the item appears, the price changes or a trusted nearby seller becomes active.”

15. Merchant jobs to be done [Back to TOC](#)

- “Make my real offers discoverable without forcing me to rebuild my business on a new marketplace.”
- “Turn Facebook, WhatsApp, marketplace and website posts into attributable leads.”
- “Let me update stock and price in seconds, not maintain another complex store.”
- “Show me what buyers are searching for, where demand is unresolved and how my pricing compares.”
- “Give my existing website a live trust and stock layer.”
- “Let my agency or staff manage my page with controlled permissions and an audit trail.”

16. Marketplace-partner jobs to be done [Back to TOC](#)

- Acquire incremental high-intent traffic without transferring checkout or customer-contract ownership to Shoppage.
- Improve product-feed quality, canonical matching and search visibility across Zimbabwean aliases and local categories.
- Understand category demand, missing products, price position and geographical opportunity outside the partner’s own traffic.
- Receive referrals with preserved campaign and product identity, then return privacy-safe outcome callbacks where commercially agreed.
- Expose marketplace trust policies, seller standards, delivery areas and buyer protection accurately on a partner page.
- Participate in cross-marketplace comparison without losing brand, merchant relationship or destination economics.

17. Ecosystem participant model [Back to TOC](#)

Participant	What Shoppage provides	What participant controls
Independent merchant	Page, offers, comparison visibility, leads, analytics	Price, stock, transaction, fulfilment, remedies

Participant	What Shoppage provides	What participant controls
Marketplace	Feed ingestion, partner page, product matching, referral traffic	Checkout, seller governance, payments, buyer protection
Retailer website	Product indexing, deep links, badge, demand insight	Website, inventory, checkout and delivery
Informal market/stall	Digital twin, QR, assisted updates, location evidence	Offer truth and local fulfilment
Agency	Multi-client workspace, reports, content and embeds	Client service and commercial relationship
Creator/group admin	Campaign links, briefs, disclosure and attribution	Audience and content execution
Buyer	Search, comparison, saved needs, request matching	Destination choice and transaction decision

18. Launch wedge and coverage cells [Back to TOC](#)

Home, Power and Building remains the best first wedge because product compatibility, local availability, bulky fulfilment, quote behaviour, high consideration and diaspora project spending create information value. The wedge should be divided into coherent subgraphs rather than treated as one catalogue: solar and backup power; electrical; tools and hardware; plumbing; building materials; appliances and parts; installation and repair services.

The recommended first cells are Msasa-Graniteside-Workington for formal and industrial supply, Glen View-Gazaland for high-density building and manufacturing activity, and Siyaso-Magaba-Mbare for informal breadth. Exact counts from internal files are hypotheses until entity resolution and field sampling confirm them.

CELL DEFINITION

A coverage cell is a bounded combination of geography, category and demand set with enough current supply to be useful. It is not merely a map polygon or a sales territory.

Part III - Consumer product and referral experience

Back to TOC

Make discovery decisive, comparison honest and handoff measurable.

19. Product-system overview [Back to TOC](#)

The consumer product has seven primary surfaces: Search, Product Comparison, Merchant Page, Marketplace Page, Market Page, Request Network and Saved Needs. Supporting surfaces include guides, shoppable clips, price history, trust evidence and local category pages. Every surface must end in a clear next action: visit destination, WhatsApp, call, request quote, reserve with the destination, get directions, save need or report a correction.

20. Local Commerce Search [Back to TOC](#)

Default search returns active local offers and reachable destinations. Results are grouped by canonical product or service, then ranked using query match, location, freshness, compatibility, merchant response, trust evidence, destination quality and user intent. Reference products may interpret the query but appear in a separate library state when no local offer exists.

The result type must be obvious. “Marketplace offer,” “Retailer checkout,” “WhatsApp seller,” “Call to confirm,” “Physical stall” and “Reference only” cannot share an ambiguous button. The buyer should know where they are going and what protection or process applies before leaving Shoppage.

21. Search result architecture [Back to TOC](#)

Layer	Public meaning	Ranking rule
Fresh offer	Price/stock or service claim updated within category SLA	Eligible for primary results
Recently confirmed destination	Merchant active, but item needs confirmation	Secondary result with “confirm” action
Corroborated profile	Business/place evidence exists; no current offer	Shown for supplier discovery, not as inventory
Reference product	Canonical knowledge only	Separate library; never implies Zimbabwe availability
Demand request	Buyer need awaiting supply	Matched privately or published with consent

22. Product comparison page [Back to TOC](#)

A comparison page represents one canonical product or tightly defined variant. It shows a specification summary, local aliases, compatibility, fresh offer count, price range by currency, locations, destination types and update times. The buyer can filter by city, distance, collection, delivery claim, marketplace, merchant type, trust evidence and freshness.

The lowest displayed price must not automatically rank first. The system should calculate a transparent “comparison context” using base price, mandatory fees disclosed by the destination, delivery estimate where supplied, stock confidence, update age and destination trust. Shoppage should avoid a false single landed-price number when fees or currency conversion are uncertain.

23. Total acquisition context [Back to TOC](#)

For products where shipping, transport, installation or exchange-rate effects are material, Shoppage displays components rather than pretending certainty. Example: item price US\$140; delivery “quoted by merchant”; collection available in Msasa; stock confirmed six hours ago; destination accepts USD and ZiG; warranty handled by merchant. The user can compare the known components and identify unknowns before referral.

24. Merchant Page [Back to TOC](#)

The Merchant Page is the merchant’s canonical Shoppage identity, not necessarily its replacement website. It contains verified name and contacts, branches, market or mall relationships, current offers, payment methods stated by the merchant, collection and delivery claims, response time, trust predicates, complaints path, corrections, social/website links and action buttons.

A merchant with a strong website keeps that website as its destination. A merchant without one receives a lightweight Shoppage Page that can serve as the link in WhatsApp About, Facebook, TikTok bio, email signatures and printed QR material. The page remains useful without requiring the merchant to operate a full website or cart.

25. Marketplace Page [Back to TOC](#)

Every participating marketplace receives a Marketplace Page explaining its categories, seller model, destination checkout, buyer-protection terms, delivery scope, accepted currencies as represented by the partner, response routes, current feed status and fresh offer coverage. Individual marketplace sellers may also have profiles where the partner permits it.

Shoppage must not merge a marketplace’s reputation with an individual seller’s evidence. A destination card should state whether the offer is sold by the marketplace, sold by a third-party seller on the marketplace, or merely advertised. The destination’s terms govern the eventual sale.

26. Market Page and stall digital twin [Back to TOC](#)

Markets in Markets converts physical clusters into searchable structures: market, zone, building, aisle, cooperative, shop and stall. A Market Page contains verified boundaries, landmarks, opening patterns, roster confidence, categories, current clips, active merchants, navigation, safety notices, operator contacts and correction paths.

Stall relationships require evidence from operators, rosters, field visits or merchant proof. Nearby coordinates alone are insufficient. Market captains may maintain rosters, but all changes are logged and conflicts of interest are visible.

27. Request Network [Back to TOC](#)

The Request Network is the system’s strategic answer to missing or stale supply. A buyer submits a structured need by text, voice note, image, bill of quantities or simple form. The request is normalised, location-bounded and matched to qualified merchants, marketplaces or agencies. Responses return through Shoppage as offers or destination links.

Requests create demand data before transactions exist. They reveal zero-result queries, incompatible offers, areas without supply, price sensitivity and quote requirements. Public requests require consent and privacy filtering. High-risk categories are excluded or manually controlled.

28. Universal Link and QR Back to TOC

Every product, offer, merchant, marketplace, market, campaign, creator and request receives a stable Shoppage Universal Link. The resolver retains privacy-safe attribution when copied between Facebook, WhatsApp, SMS, QR and browsers, then routes to the best available destination. If an external app is unavailable, the link falls back to web, phone, directions or a saved need.

Universal IDs are Shoppage-owned. Facebook Page IDs, WhatsApp numbers, Google Business IDs and partner listing IDs are external identifiers attached to the canonical record, never the primary key or only recovery route.

29. Referral Ledger Back to TOC

The Referral Ledger records impression, comparison view, outbound click, WhatsApp start, call reveal, directions open, quote request, reserve intent, destination acknowledgement and outcome callback. Events are deduplicated, timestamped and assigned confidence. A WhatsApp click is an attributable action, not a sale.

Outcome states may include contacted, responded, quoted, unavailable, reserved, collected, purchased, abandoned and disputed. Merchant self-report, buyer confirmation, destination callback and sampled evidence carry different weights. Revenue products must use the event level they actually price; a cost-per-click product cannot be marketed as cost-per-sale proof.

30. Trust Passport Back to TOC

Trust is a set of expiring predicates, not one badge. Examples include phone OTP verified, branch evidence reviewed, marketplace seller status confirmed, market-roster match, company-registration evidence supplied, price refreshed today, recent response SLA, complaint resolved and collection proof sampled. Each predicate has source, date, reviewer, expiry and challenge path.

Shoppage should display what was checked and what was not. “Phone verified” is useful; “verified seller” without scope is misleading. Paid plans cannot buy a trust predicate. They may pay for a review service, but the outcome remains evidence-based.

31. Live Trust Seal - the 5th Icon Back to TOC

The strongest v3.0 distribution idea remains: a Shoppage icon and live seal next to Facebook, Instagram, LinkedIn or WhatsApp on merchant and marketplace websites. The seal is served dynamically, links back to the canonical page and can show factual states such as “offers refreshed today,” “open now” or “median response 22 minutes” when supported.

The seal must degrade to grey when data is stale and must never display stock counts pulled from an unauthorised or unreliable source. Removal should not be described publicly as evidence of wrongdoing; the value comes from positive live information, not coercive signalling.

32. Showcase Shorts Back to TOC

Showcase Shorts is a product-demonstration and referral format, not a separate entertainment network. Clips are attached to products, offers, merchants or markets and end in a Universal Link. Useful formats include compatibility demonstrations, new-stock walk-throughs, before-and-after installations, market navigation and price updates.

Natural stopping points, data controls, visible sponsorship and muted defaults are mandatory. Video cannot outrank fresh structured truth merely because it generates engagement. Merchants can embed a Shoppage Shorts carousel on their own websites while checkout remains at the merchant or marketplace destination.

33. Saved Needs and alerts Back to TOC

Buyers may save a product, specification, place or unresolved need and receive price, stock, new-seller or compatibility alerts. Alerts should be event-driven and frequency controlled. A saved need becomes a first-party demand asset and a direct-return channel, improving Owned Demand Share.

34. Guides and decision support [Back to TOC](#)

Shoppage should publish practical, local decision pages: inverter and battery compatibility, cable sizing, roofing quantity checks, genuine-versus-counterfeit indicators, product aliases, market guides and supplier-selection checklists. Guides route to comparison pages and requests, not to an editorial content business disconnected from commerce.

35. Language, aliases and input modes [Back to TOC](#)

Search must recognise brand abbreviations, common misspellings, local units, installer shorthand, Shona and Ndebele terms, and mixed-language phrases. Voice and image input should create a draft query and show uncertainty rather than silently inventing a model. Human-reviewed synonym sets become part of the moat.

36. Low-data and inclusive design [Back to TOC](#)

The core path must work on low-cost Android devices, intermittent connections and text-first mode. Search, contact, price, freshness, location and action buttons load before video or decorative media. No installed app or buyer registration is required for basic discovery. Offline queues may preserve merchant updates and buyer requests until connectivity returns.

Part IV - Marketplace federation [Back to TOC](#)

Turn existing marketplaces into indexed destinations, distribution partners and intelligence customers.

37. Why marketplaces should list on Shoppage [Back to TOC](#)

A marketplace has its own catalogue but still competes for discovery against Google, social media, retailer websites and other marketplaces. Shoppage can deliver incremental, category-specific Zimbabwe traffic and expose products through local aliases, compatibility and place signals that the marketplace may not maintain. Listing does not require surrendering checkout, seller contracts or payment economics.

The partnership is analogous to a store participating in a comparison-shopping service: product information is submitted or authorised, Shoppage matches it to canonical products, and users click through to the destination. Google's Comparison Shopping Service model and merchant feed APIs show that merchants and aggregators can coexist through structured feeds and destination links. [S13-S14]

38. Partner value proposition [Back to TOC](#)

- Incremental traffic from users comparing across the whole market rather than searching one marketplace.
- Feed-quality diagnostics: missing identifiers, duplicate variants, stale prices, broken links, weak titles and attribute gaps.
- Local query intelligence: Shona/Ndebele aliases, informal terminology, market areas and zero-result demand.
- Category opportunity maps showing demand with limited partner supply.
- Price-position and freshness benchmarks using rights-cleared, aggregated data.
- Embeddable comparison, live badge, price-history or market-intelligence widgets.
- Optional privacy-safe referral and outcome measurement for partner acquisition economics.

39. Marketplace integration modes [Back to TOC](#)

Mode	Use case	Control
Authorised crawl	Partner permits specified pages and crawl frequency	Robots/contract scope, rate limits, takedown
Scheduled feed	CSV, XML, JSON or merchant-feed format	Partner is source of price/availability truth
API/webhook	Large or frequently changing catalogue	Incremental updates, authentication and quotas
Affiliate network	Partner already exposes links and commissions	Affiliate terms govern attribution
Manual curated pilot	Early partner or limited category	Small set, human QA, no scale assumptions
Marketplace widget	Partner embeds Shoppage intelligence on its site	Scoped aggregate data; no competitor leakage

40. Partner Feed Contract Back to TOC

Every feed requires a written contract or accepted feed terms defining permitted fields, image and description rights, refresh frequency, destination URL rules, seller identity, availability meaning, currencies, returns/buyer-protection representation, retention, error handling and termination. Shoppage stores raw input separately from the processed canonical offer.

A feed record does not become public merely because it arrived. It passes validation, malware and URL checks, entity matching, prohibited-product rules and confidence thresholds. Partner corrections take precedence for partner-controlled fields, while Shoppage retains audit history and independent trust evidence.

41. Deep-link and destination rules Back to TOC

The referral should open the most specific valid destination: exact listing, product page, seller page or marketplace search. Links must not redirect through misleading intermediate pages, replace the partner brand or force Shoppage registration. Broken, unavailable or repeatedly mismatched destinations lose ranking and may be suspended.

42. Cross-marketplace ranking neutrality Back to TOC

Organic ranking is based on query relevance, product match, freshness, location, destination quality, disclosed total-cost context, trust and user preferences. Commercial payment is not an organic relevance factor. Sponsored listings occupy separately labelled inventory and are subject to the same truth and safety requirements.

Shoppage should publish a plain-language ranking policy and maintain an internal feature registry. Partner-specific confidential data cannot be used to advantage a competitor. Aggregate insights require minimum cohort sizes and disclosure controls.

43. Attribution callbacks Back to TOC

A partner may return privacy-safe events such as landing accepted, listing viewed, seller contacted, added to cart, order placed or cancelled. The core Shoppage model does not require transaction-level data, but richer callbacks improve optimisation and permit affiliate or performance contracts. Consent, data minimisation and retention rules apply.

44. Partner analytics Back to TOC

The partner dashboard shows referred sessions, categories, products, geographical demand, broken links, stale-feed rates, unmatched records, destination conversion where callbacks exist, and benchmark position. The dashboard must distinguish observed data from inferred demand and must never expose another partner's confidential transaction data.

45. Brand and anti-cannibalisation controls Back to TOC

Partners may fear price transparency or loss of direct traffic. Shoppage should allow agreed brand presentation, destination badges, seller-policy summaries and category participation. It should not permit partners to suppress legitimate cheaper offers merely to protect margin. The commercial answer is incremental demand and intelligence, not hidden exclusivity.

46. No-custody boundary Back to TOC

TRANSACTION RESPONSIBILITY

The contract of sale, payment, fulfilment, return, refund and warranty remains between the buyer and the destination merchant or marketplace under its terms. Shoppage provides discovery, comparison, referral and evidence. Any future transaction service must be a separately approved product with separate legal, financial and operational controls.

47. Marketplace pilot design [Back to TOC](#)

The first partner pilot should use three different destination types: one established marketplace, one formal retailer with a functioning website and one high-volume independent merchant or distributor. Each contributes 500-2,000 offers in one wedge category. The pilot measures match rate, freshness, click-through, destination quality, partner feedback and incremental demand over six weeks.

The pilot is successful only if the partner can identify traffic or insight it values, buyers experience fewer dead links than current alternatives, and Shoppage can maintain data without excessive manual labour. A signed logo is not a successful partnership.

Part V - Merchant, agency and field operating system

Back to TOC

Give each participant a reason to maintain truth and distribute Shoppage links.

48. Merchant claim workflow [Back to TOC](#)

Preloaded candidate profiles are not active shops. A merchant claims a profile through phone/email verification, identity or authority evidence, branch selection and source review. The claim process shows existing data, its source class and correction options. Disputed profiles enter a controlled review queue; suppression is fast where risk is credible.

49. Offer-ingestion modes [Back to TOC](#)

Mode	Best for	Operating rule
Snap-Extract-Confirm	Informal and micro merchants	AI drafts from image/PDF/voice; merchant confirms before publish
Spreadsheet/CSV	SMEs and distributors	Template validation and error report
Website or WooCommerce plugin	Existing online retailers	Authorised feed; destination remains merchant site
API/POS integration	Large merchants and marketplaces	Field ownership and conflict rules
Agent-assisted capture	Markets and low-digital-maturity merchants	Audit trail, merchant consent, quality sampling
Quick freshness confirm	Existing offers	10-second stock/price state update by web or WhatsApp

50. Offer freshness system [Back to TOC](#)

Freshness is a state machine, not a reminder. Each category and source type has a service level. A manually confirmed fast-moving item may expire after 24-72 hours; a manufactured-to-order service may remain active for weeks with a quote-required label; a marketplace API offer may be current for minutes. Expired offers are demoted or hidden automatically.

The public timestamp should say what happened: price updated, stock confirmed, merchant active or feed received. These are not interchangeable. Freshness history becomes a quality and defensibility asset.

51. Business Improvement Score [Back to TOC](#)

The score measures work that improves buyer outcomes: offer completeness, freshness, response rate, destination accuracy, correction history, resolved complaints and evidence quality. It does not reward logins, advertising spend or volume alone. The merchant receives specific actions such as “confirm five high-demand offers” or “respond to two requests older than one hour.”

52. Lead and request inbox [Back to TOC](#)

Merchants see Shoppage-originated calls, WhatsApp starts, quote requests, directions and matched needs in one queue. They can mark responded, unavailable, quoted, reserved, fulfilled or invalid. The system samples marking quality and compares merchant states with buyer or partner evidence.

53. Quote and reserve handoff [Back to TOC](#)

Shoppage may structure a request or intent without taking payment or becoming the seller. A quote response includes item, quantity, currency, validity, location, fulfilment method and destination terms. A reserve button routes to the merchant or partner process and records an intent event. The destination confirms or rejects the reservation.

54. Export and distribution cards [Back to TOC](#)

A merchant can generate WhatsApp Status, Facebook, Instagram, email and printable cards with product, current price, freshness, branch and Universal Link. The card is a distribution tool that sends traffic back through the canonical page for comparison and attribution. Merchants do not need advertising expertise to participate.

55. Reviews and reputation [Back to TOC](#)

Shoppage should avoid generic anonymous star ratings at launch. Early reputation should use structured outcomes: response speed, offer accuracy, destination reached, item available, quote honoured and complaint resolved. Transaction reviews from partner callbacks are labelled separately from referral-experience reviews.

56. Agency workspace [Back to TOC](#)

Agencies manage multiple clients with delegated permissions. The workspace includes freshness health, lead queues, Price Pulse, Demand Radar, missing-product opportunities, content cards, Live Seal installation and weekly client reports. The agency earns by maintaining commercial truth and acting on demand, not merely producing posts.

57. Agency commercial model [Back to TOC](#)

Launch with one agency pilot plan rather than complex tiers: up to 20 client profiles, five advanced-intelligence clients, branded reports and certification training. The indicative wholesale fee is US\$149 per month after a proof period. Agencies may resell Growth plans or charge managed-service retainers; Shoppage does not dictate their service price.

58. Market captains [Back to TOC](#)

Market captains onboard stalls, maintain rosters, capture evidence, distribute QR materials and escalate disputes. Their compensation is tied to accepted, durable work: claimed merchants still active after 30 days, refreshed offers, resolved requests and correction quality. Raw registration counts are not payable outcomes.

Controls include identity badges, territory, audit logs, conflict declarations, merchant complaint route, random rechecks and prohibition on charging merchants unofficial verification fees. Political or operator relationships are documented but do not override evidence standards.

59. First-party plugins and embeds [Back to TOC](#)

The initial plugin set is narrow: WordPress/WooCommerce Offer Feed, Live Seal, Showcase Shorts Embed, Price Pulse Ticker and Email Signature Generator. The plugins use standard product and offer fields compatible with Schema.org concepts such as Product, Offer and AggregateOffer. [S15, S17]

A public developer marketplace is deferred. Third-party integrations begin through reviewed partner agreements, scoped keys, quotas, audit logs and data-minimisation rules.

Part VI - Commerce Intelligence Graph, AI and technology [Back to TOC](#)

Build the national data asset without confusing volume with truth.

60. Core graph model [Back to TOC](#)

The Commerce Intelligence Graph links Organisation, Merchant, Branch, Marketplace, Seller-on-Marketplace, Market, Stall, Place, Product Family, Variant, Offer, Service, Compatibility Edge, Request, Referral Event, Trust Evidence, Campaign and Outcome. Every claim carries source, rights class, timestamp, confidence and field owner.

61. Authoritative-field doctrine [Back to TOC](#)

Different systems own different fields. A regulator or merchant-supplied document may own legal identity; Shoppage owns canonical product IDs and graph relationships; the merchant or authorised partner owns current price and availability; a website may own editorial brand copy; Shoppage owns referral history and its own trust reviews. Bidirectional synchronisation is prohibited until ownership and conflict rules are explicit.

62. Entity resolution [Back to TOC](#)

The 200,000-company file and other directories require deduplication across names, phones, addresses, coordinates, registration numbers, websites and source history. Resolution outputs are candidate links with confidence, not silent merges. High-impact merges require human review. A merchant can separate incorrectly merged branches or organisations.

63. Master product catalogue [Back to TOC](#)

The catalogue starts with the top demand entities in the launch wedge, not one million decorative records. A canonical variant includes brand, model, identifiers, specifications, images with rights, units, aliases, compatibility and category attributes. Merchant offers attach to the variant. Generic or ambiguous items remain families until sufficient evidence exists.

Google's merchant-product model and Schema.org's AggregateOffer provide useful external reference patterns for product inputs and multiple seller offers, but Shoppage should maintain a simpler local schema and map outward. [S14-S15]

64. Reference Library [Back to TOC](#)

Reference products support query understanding, education, compatibility and merchant enrichment. They are clearly labelled "reference - no current Zimbabwe offer confirmed" and are excluded from active-offer, active-merchant and revenue KPIs. Merchants may attach a local offer after validation.

65. Source Rights Register [Back to TOC](#)

Every source defaults to BLOCKED until its rights class is reviewed. The register records owner, access method, contractual basis, public-display permission, AI-use permission, image/text rights, derivative-data rights, required attribution, refresh, retention, geographic scope and suppression procedure. Technical accessibility is never treated as legal permission.

66. Offer and price history [Back to TOC](#)

Shoppage stores time-series states for permitted offers: price, currency, stock claim, freshness, destination and source. History powers alerts, anomaly detection and Price Pulse. It must preserve corrections without publicly exposing commercially sensitive partner data beyond agreed fields.

67. Demand graph [Back to TOC](#)

Search queries, filters, zero results, saved needs, structured requests, referrals and outcomes form the demand graph. Personal identity is separated from aggregate intelligence. Demand is segmented by location, category, time, urgency and fulfilment preference. This graph is strategically more valuable than raw page views because it reveals unsolved commerce.

68. AI use cases [Back to TOC](#)

- Extract draft products and offers from images, PDFs, messages, spreadsheets and voice notes.
- Normalise titles, units, variants and attributes; identify likely duplicates for review.
- Expand local aliases, spelling variations and multilingual queries.
- Match requests to compatible products, merchants, marketplaces and services.
- Draft export cards, captions, comparison summaries and merchant action recommendations.
- Prioritise stale offers, suspicious prices, copied scam images and support queues.
- Summarise evidence and uncertainty for human reviewers.

69. AI prohibitions [Back to TOC](#)

- AI may not invent local price, stock, location, verification, warranty or delivery claims.
- AI may not publish a merchant offer without confirmation or an authorised feed.
- AI may not make high-impact fraud, suspension or identity decisions without review and appeal.
- AI may not train on a source merely because the source was technically collected.
- AI may not personalise prices or conceal offers based on inferred vulnerability.
- AI may not merge organisations, people or branches at high confidence without evidence controls.

70. Evaluation harness [Back to TOC](#)

Before launch, maintain at least 500 Zimbabwe-specific queries across product aliases, misspellings, languages, locations, compatibility and zero-result cases. Each model release is tested for correct entity, local availability hallucination, unsafe category handling, destination accuracy and uncertainty. Behavioural learning-to-rank is gated until at least 5,000 resolved needs produce reliable labels.

71. Privacy and data protection [Back to TOC](#)

Shoppage will process merchant contacts, buyer requests, referral events and potentially location or communication metadata. It must determine the appropriate data-controller licence tier, appoint qualified responsibility, maintain processing records, security controls, data-subject routes, breach procedures and cross-border transfer review. Zimbabwe's 2024 regulations define licence tiers by number of data subjects and impose controller obligations. [S05]

72. Security baseline [Back to TOC](#)

- Multi-factor authentication for staff, agencies and high-risk merchant actions.
- Role-based permissions with branch, client and source scopes.
- Encryption in transit and at rest; secrets vault; signed feed/webhook requests.
- Immutable audit events for profile claims, offer changes, trust decisions and API access.
- Rate limits, bot controls, media scanning, URL reputation checks and abuse reporting.
- Daily backups, tested restoration, incident severity matrix and notification runbook.
- Data minimisation: do not store transaction or identity data merely because a partner can send it.

73. Recommended technology architecture [Back to TOC](#)

Start with a modular monolith rather than microservices. A practical stack is Next.js or equivalent PWA frontend; TypeScript backend; PostgreSQL for transactional and graph-edge records; OpenSearch or Typesense for search; Redis for caching and queues; S3-compatible object storage; an analytics warehouse after event volume justifies it; and provider-independent AI adapters.

Core modules are Identity and Permissions, Organisations and Places, Catalogue, Offers, Search, Requests, Referrals, Trust, Campaigns, Partner Feeds, Analytics and Rights. There is deliberately no Cart, Order, Payment, Settlement or Delivery domain in the core architecture.

74. Independence doctrine [Back to TOC](#)

Shoppage must retain canonical identity, web access, merchant maintenance, search, requests, saved needs, referrals and notifications without depending on Facebook, WhatsApp, Google or TikTok. External platforms are acquisition and communication adapters. Meta's WhatsApp catalog and Cloud API capabilities are useful integrations, but Shoppage's graph and account recovery must remain independent. [S16]

75. Reliability and performance [Back to TOC](#)

Target text-first first content under 1.5 seconds on representative low-end Android tests and compressed initial transfer under 200 KB for search results. Critical events queue offline and retry. Feed failures isolate by partner. Search and Universal Links receive the highest availability priority because they are the revenue path.

76. Build-versus-buy doctrine [Back to TOC](#)

Build the graph, matching, freshness, referral ledger, request network and local ranking because they are differentiating. Buy or use managed services for authentication, email/SMS delivery, object storage, observability, support ticketing and standard analytics where switching risk is acceptable. Do not build a video transcoding platform, payment gateway or general CRM.

Part VII - Go-to-market and distribution [Back to TOC](#)

Create density through partners, agencies, groups and physical markets without becoming platform-dependent.

77. Coverage-cell launch sequence [Back to TOC](#)

1. Choose 300 high-intent product/service entities and three Harare cells.
2. Recruit anchor destinations capable of covering at least 60% of those entities.
3. Integrate one marketplace, one formal retailer and a set of independent merchants.
4. Create current offers, compatibility edges, market pages and request routes.
5. Run closed buyer tests against Facebook/WhatsApp baseline tasks.
6. Open the cell publicly only when dead-end and freshness thresholds pass.
7. Replicate the operating playbook to the next cell; do not expand nationally by default.

78. Anchor-merchant acquisition [Back to TOC](#)

The first 30 anchor merchants should be chosen for catalogue breadth, response capability, branch relevance, willingness to refresh and reputation - not logo prestige alone. The offer is assisted onboarding, free foundation listing, weekly demand report and co-designed lead attribution. A partner who will not update availability should not dominate the launch result set.

79. Marketplace acquisition [Back to TOC](#)

The marketplace pitch is not “join our marketplace.” It is “let Shoppage index your authorised catalogue, improve its local discoverability and send you measured incremental traffic while you keep checkout.” The pilot proposal in Appendix F includes feed scope, data rights, brand presentation, attribution, reporting and commercial options.

80. Buyer acquisition [Back to TOC](#)

Initial demand acquisition targets existing high-intent behaviours: product queries in search, Facebook-group questions, WhatsApp referrals, market QR codes, contractor lists, diaspora project searches and agency client audiences. Campaigns route to a product, market or request page, not a generic homepage.

81. Facebook Groups doctrine [Back to TOC](#)

The 5,200-group file is a distribution and research map, not an audience Shoppage owns. No scraping or automated posting without permission. Priority is admin partnership: useful answer cards, weekly demand summaries, scam-awareness content, group-specific landing pages and tracked links. Group overlap means membership sums must never be treated as unique reach.

82. WhatsApp doctrine [Back to TOC](#)

WhatsApp remains the dominant merchant action rail: click-to-chat, status cards, quick freshness confirmation, request matching and service notifications. Shoppage should not attempt to replace merchant conversations. It creates a structured context before the chat and captures outcome states after it.

83. Search and SEO [Back to TOC](#)

SEO assets are canonical product pages, local offer comparisons, market pages, merchant pages, compatibility guides and structured data. Thin unclaimed profiles, duplicate product pages and automatically generated text are blocked from indexing. Merchant listing structured data can expose price and availability when the evidence is current. [S12]

84. Creator programme [Back to TOC](#)

Creators are paid fixed fees for accepted briefs during the attribution-learning stage. Requirements include merchant consent, disclosure, canonical links, factual claims, content-live period and fraud checks. Outcome commissions begin only when event definitions, duplication controls, partner callbacks, disputes and payout reconciliation are reliable.

85. Agency channel [Back to TOC](#)

Ten certified agencies managing 15-20 clients can activate more durable supply than a large temporary field team. Certification covers offer freshness, rights, market interpretation, request campaigns, trust evidence and reporting. Agencies are listed by city and category only after their clients meet quality thresholds.

86. Market and association partnerships [Back to TOC](#)

Market operators, trade associations, manufacturers, distributors and local authorities can provide rosters, category evidence, safety information and campaign access. Partnerships must define rights and governance; an institutional logo cannot substitute for merchant consent or current offer truth.

87. Diaspora distribution [Back to TOC](#)

Diaspora users often need supplier discovery, quotation and local fulfilment confidence rather than another checkout. Shoppage can compare destination options, route to established marketplaces or merchants, provide project request templates and record recipient/merchant confirmation where consented. Payments remain with the destination.

88. Expansion logic [Back to TOC](#)

Expand category when the current cell has high search success, healthy freshness and merchant retention. Expand city when the onboarding and market-captain playbook produces a useful cell at predictable cost. Expand to regional countries only when rights, language and source economics are repeatable. “National” is the eventual graph ambition, not a launch-stage claim.

Part VIII - Revenue model and unit economics [Back to TOC](#)

Monetise attributable value without forcing transaction custody.

89. Revenue principles [Back to TOC](#)

Revenue must align with value Shoppage can actually evidence. The business should not depend on taking a percentage of transactions it does not process. It can combine recurring software, paid distribution, qualified-action pricing, partner intelligence, managed activation and affiliate revenue where the destination supports it.

- Free supply improves consumer utility; monetisation cannot destroy coverage.
- Organic relevance and trust remain independent of payment.
- Labour-intensive activation is reported as service margin, not hidden in software gross margin.
- Performance fees use defined events and confidence thresholds.
- No merchant fee launches nationally before willingness-to-pay and retention are proven in a cell.

90. Launch pricing [Back to TOC](#)

Product	Indicative launch price	Included value
Free Foundation	US\$0	Claimed page, limited offers, contact actions, basic freshness and corrections
Growth Pilot	US\$19/month	More active offers, alerts, lead history, export cards, Live Seal and basic demand insights
Agency Pilot	US\$149/month	20 clients, five advanced intelligence slots, reports and delegated management
Partner Feed Pilot	Negotiated or free proof period	Feed QA, product matching, partner page and referral reporting
Managed activation	US\$25-US\$150 once-off by scope	Catalogue cleaning, branch setup, imports and evidence capture

91. Merchant subscription economics [Back to TOC](#)

The US\$19 plan is a hypothesis, not a universal truth. It is viable when the merchant receives attributable gross-profit opportunity materially above the fee and can maintain the page with little labour. Target month-3 retention above 75% among merchants who received at least five qualified actions. Merchants with no useful demand should not be pressured to remain paid.

92. Cost per click and qualified action [Back to TOC](#)

After the attribution gate, merchants or marketplaces may buy prepaid referral credits by category. Events include outbound product click, WhatsApp start, call reveal, quote request or directions. Pricing differs by intent

and category. Invalid, duplicate, bot and broken-destination events are excluded. A qualified lead product requires more information and consent than a click product.

93. Sponsored placements [Back to TOC](#)

Sponsored product, merchant, marketplace, market and guide placements are clearly labelled. Eligibility requires current data and policy compliance. Sponsorship changes visibility in allocated inventory but does not alter trust evidence or organic ranking. Category exclusivity should be short, transparent and never block legitimate organic competitors.

94. Agency revenue [Back to TOC](#)

Agency subscription provides recurring software revenue while agencies create a larger managed-service economy around Shoppage. Additional revenue may come from certification, co-branded market reports and campaign credits. Certification fees should recover training and review costs, not sell status.

95. Marketplace and retailer revenue [Back to TOC](#)

- Feed-quality and catalogue-matching subscription.
- Category Demand Radar and Price Pulse intelligence.
- Attributable CPC, affiliate or action-based referral agreements.
- Embeddable comparison, compatibility or market widgets.
- Private data-quality API and supplemental-feed service.
- Co-funded category or geographic campaigns.

96. Intelligence products [Back to TOC](#)

After sufficient density, Shoppage can sell aggregated, privacy-safe intelligence to merchants, agencies, manufacturers, distributors and institutions: price ranges, freshness, category gaps, search demand, geographic availability, compatibility demand and market activity. Minimum cohort sizes, confidence labels and source rights are mandatory.

97. Unit-economics definitions [Back to TOC](#)

Metric	Definition	Management use
Active merchant	Claimed merchant with current eligible offer or active service within 30 days	Real supply base
Qualified demand session	Session with product/service intent, location or action signal	Demand denominator
Attributable action	Valid outbound click, chat, call, quote, reserve intent or directions	Commercial activity
Resolved need	Buyer confirms suitable destination or high-confidence outcome event exists	North-star output
Revenue per active merchant	Recurring and usage revenue divided by active merchants	Value capture

Metric	Definition	Management use
Activation payback	Acquisition + onboarding cost divided by monthly contribution	Scalability
Fresh Offer Coverage	Top demand entities with required number of fresh local offers	Cell usefulness

98. Illustrative month-12 base economics [Back to TOC](#)

Revenue stream	Monthly base case	Assumption
Growth merchants	US\$15,200	800 paid x US\$19
Agency plans	US\$5,960	40 agencies x US\$149
Paid referrals and sponsorship	US\$8,000	Measured category demand; no sale claims
Partner/data products	US\$3,500	Marketplace, retailer and supplier pilots
Activation and services	US\$4,000	Reported separately from software
Total monthly revenue	US\$36,660	Illustrative target, not forecast

A disciplined month-12 cost base could range from US\$28,000 to US\$38,000 depending on local staffing, field operations, infrastructure, media and professional services. The base case therefore approaches operating break-even only if activation labour, paid acquisition and support remain controlled. Management should prefer a smaller profitable cell over a larger subsidised directory.

99. Twenty-four-month scenarios [Back to TOC](#)

Scenario	Month 12	Month 24
Downside	900 active merchants; 50k monthly intent; US\$12k revenue	1,800 merchants; US\$24k revenue; concentrate or pivot
Base	2,500 active merchants; 150k intent; US\$37k revenue	8,000 merchants; 700k intent; US\$110k-US\$130k revenue
Upside	4,000 active merchants; 250k intent; US\$60k revenue	15,000 merchants; 1.5m intent; US\$220k+ revenue

The scenario driver is not the number of preloaded profiles. It is the relationship among fresh offer coverage, qualified demand, attributable actions, paid conversion, agency leverage and activation cost. The board should reject any plan that reports raw catalogue size as commercial traction.

100. Capital requirement [Back to TOC](#)

A 12-18 month disciplined build and market-formation programme is likely to require approximately US\$250,000-US\$450,000 depending on founder compensation, engineering model and field intensity. A staged raise is preferable: proof capital for one cell, then expansion capital after the six launch gates pass.

Use	Indicative share	Control
Product and engineering	35%	Modular core; no cart/payment build
Data and catalogue operations	20%	Rights and quality measured
Market activation and partners	18%	Pay durable outcomes, not registrations
Demand acquisition and content	12%	Campaign-level attribution
Trust, legal, security and finance	10%	Pre-launch gates
Contingency	5%	Board-controlled

101. Break-even and kill logic [Back to TOC](#)

The business should not chase break-even by flooding search with low-quality sponsored inventory. Break-even is acceptable only with healthy repeat demand, paid retention, freshness and partner value. Kill or redesign a cell when, after two disciplined iterations, fewer than 25% of qualified needs resolve, merchant response stays below 40%, fresh coverage cannot exceed 50%, or activation payback remains above 12 months.

Part IX - Organisation, governance, compliance and execution [Back to TOC](#)

Make the operating model auditable, resource-constrained and capable of saying no.

102. Launch team [Back to TOC](#)

Role	Core accountability	Initial shape
CEO / commercial lead	Partners, capital, strategy, key accounts	Founder
Product and technology lead	Core platform, data and delivery	1 senior + contractors
Commerce data lead	Taxonomy, matching, rights and quality	1
Market activation lead	Cells, captains, agencies and merchants	1
Trust and support lead	Claims, corrections, complaints and abuse	1
Growth/content lead	Demand, SEO, groups, creators and reporting	1
Finance/legal support	Controls, contracts, compliance and board pack	Fractional initially

103. Governance cadence [Back to TOC](#)

- Daily: feed failures, stale offers, critical complaints, security and launch-blocking issues.
- Weekly: cell scorecard, request resolution, merchant activation, partner funnel, experiments and cash.
- Monthly: rights review, ranking audit, privacy/security, revenue quality, activation payback and source concentration.
- Quarterly: category/city expansion decision, model changes, capital plan and independent risk review.

104. Data operations [Back to TOC](#)

Data operations has explicit queues: source approval, entity resolution, product review, offer exceptions, image rights, merchant claims, corrections, suspicious price, stale feed and request matching. Each queue has age, owner, severity and sampling. Quality is measured by false merge, wrong variant, broken destination, stale offer and complaint rates.

105. Trust and safety operations [Back to TOC](#)

The trust team handles impersonation, fake location, copied images, prohibited goods, misleading price, bait-and-switch, harassment, fraudulent requests and evidence disputes. High-risk actions require dual review. Merchants and partners receive notice, reason and appeal except where safety or law requires immediate restriction.

106. Legal and compliance workstreams [Back to TOC](#)

- Data-controller licensing, DPO responsibility, records of processing, subject rights and breach response.
- Merchant/partner feed terms, content licences, image rights, brand use and takedown procedures.
- Consumer-facing referral terms that accurately allocate transaction responsibility.
- Advertising disclosure, sponsored-ranking policy and creator disclosure.
- Competition review for ranking, exclusivity, price intelligence and partner confidentiality.
- Electronic communications consent and WhatsApp/SMS notification controls.
- Restricted-product policy, sanctions screening where relevant and law-enforcement request process.
- Tax treatment of subscriptions, advertising, services and cross-border receipts.

107. Principal risks [Back to TOC](#)

Risk	Early warning	Primary control
Stale or false offers	High unavailable outcomes, complaints	Expiry, source weighting, auto-hide
No consumer habit	Low direct repeat and saved needs	Narrower cells, requests, alerts and guides
Merchant non-maintenance	Freshness falls after onboarding	Quick confirm, agency/captain support, ranking consequence
Marketplace resistance	Partners fear comparison or leakage	Pilot, brand controls, aggregate analytics, no custody
Data-rights challenge	Takedown or source dispute	Default BLOCKED register and contracts
Platform dependence	Traffic concentrated in one source	Owned Demand Share target and universal identity
Weak attribution	Merchants dispute value	Event definitions, deduplication and confidence
Field corruption	Fake verification or unofficial fees	Audit, complaints, outcome-based pay and rechecks
Security/privacy incident	Unauthorised access or exposure	Licence compliance, least privilege, incident plan
Overbuilding	Roadmap expands before cell proof	Binding kill list and gate review

108. First 30 days Back to TOC

1. Freeze the category and kill list in a signed decision register.
2. Audit the company, group and product files for rights, duplication, freshness and usable launch coverage.
3. Interview 30 buyers, 30 merchants, five marketplaces/retailers and five agencies using real buying tasks.
4. Define the top 300 wedge entities and compatibility schema.
5. Prototype search, comparison, Merchant Page, Marketplace Page, Request Network and Universal Link.
6. Draft merchant, marketplace, agency, creator and data-processing terms.
7. Secure three anchor destination pilots and one market/operator relationship.
8. Build the first 100-query evaluation set and manual baseline against WhatsApp/Facebook discovery.

109. Days 31-90 Back to TOC

1. Build the modular core, referral ledger and partner-feed pipeline.
2. Activate 100-150 merchants and three partner destinations in one cell.
3. Publish 5,000-10,000 fresh offers across the top demand set.
4. Run closed tests with at least 300 qualified needs.
5. Measure destination success, response, unavailable rate, request resolution and repeat use.
6. Install 30 Live Seals or website embeds and 300 QR assets.
7. Pilot two agencies and two market captains with audited compensation.
8. Conduct independent privacy/security and source-rights review before public launch.

110. Months 4-6 Back to TOC

Stabilise one cell. Improve resolution before adding categories. Target 600 active merchants, 35,000 fresh offers, 35,000 monthly intent sessions, 4,000 attributable actions, 30% resolved demand, 25% direct repeat and early paid conversion. Introduce limited sponsorship and Growth Pilot only after value interviews.

111. Months 7-12 Back to TOC

Replicate into two additional cells using the same category graph and operating playbook. Add 10-15 marketplace/retailer partners, 40 agencies, three formal market pages and a controlled creator programme. Launch paid referral credits only when attribution passes. Prepare partner intelligence and feed-quality products.

112. Months 13-24 Back to TOC

Choose among three paths based on evidence: deepen Home/Power/Building nationally; expand into adjacent high-intent categories such as auto parts or business supplies; or license the intelligence/referral stack to agencies and partners in another market. Do not pursue all three simultaneously.

113. KPI tree Back to TOC

The north-star metric is Resolved Demand Rate. It is driven by qualified demand, fresh offer coverage, correct matching, destination reachability, merchant response and request matching. Revenue is a downstream outcome; raw page views, profiles and catalogue records are diagnostic only.

Level	Metrics	Guardrails
North star	Resolved Demand Rate; resolved needs per active merchant	Complaint and false-resolution rate
Supply	Fresh Offer Coverage; active merchants; partner coverage	Stale, duplicate and wrong-variant rate

Level	Metrics	Guardrails
Demand	Qualified sessions; saved needs; direct repeat	Bot and low-intent share
Referral	Attributable actions; destination success; response	Broken links and unavailable outcomes
Revenue	MRR; paid conversion; revenue per active merchant	Churn, concentration and service-labour burden
Independence	Owned Demand Share; external platform concentration	Single-source traffic risk

114. Investor proof package [Back to TOC](#)

- Live cell with search and comparison that outperforms baseline buying tasks.
- Rights register and independent legal/privacy review.
- Partner agreements showing marketplaces list without surrendering checkout.
- Merchant cohort retention and attributable-value evidence.
- Freshness and response history, not screenshots of total profiles.
- Unit economics by merchant, agency, marketplace partner and field-activation channel.
- Product evaluation set with hallucination and wrong-offer rates.
- Board-ready cash plan, risk register and kill criteria.

115. Final verdict [Back to TOC](#)

THE MODEL

Shoppage should become Zimbabwe's neutral commerce discovery and intelligence layer - the place where a buyer compares every credible destination and where every merchant, marketplace, website, market and agency can distribute current offers without surrendering its own transaction.

The original referral strategy was directionally right. The 10x improvement is not to add a cart. It is to make referral commercially rigorous: rights-cleared data, current offers, canonical products, marketplace federation, transparent ranking, measurable handoffs, demand resolution, agency leverage, physical-market coverage and disciplined unit economics.

The business wins when an Ownai listing, Classifieds advert, retailer checkout, WhatsApp merchant and physical stall can all appear on one honest product comparison page, and the buyer can reach the best destination with less uncertainty. Shoppage then becomes more valuable as marketplaces multiply. It is infrastructure between commerce channels, not another channel demanding exclusive inventory.

Part X - Defensibility, network effects and scale [Back to TOC](#)

Separate real compounding assets from features that competitors can copy in weeks.

116. Moat doctrine [Back to TOC](#)

Shoppage should stop calling every feature a moat. A feature is useful but copyable. A defensible asset compounds through time, participation, rights or operational learning and becomes expensive to reproduce. Management should report each proposed moat with four tests: what unique input accumulates, why participants keep contributing, how a competitor could acquire the same input, and what measurable advantage appears in buyer resolution or merchant economics.

The most credible defence is a linked system rather than a single database: rights-cleared product and place identities; current offer histories; local compatibility and alias knowledge; marketplace and merchant integrations; demand and referral outcomes; physical market relationships; and a distributed agency/captain maintenance network. Each component reinforces the others.

117. Rights-cleared Commerce Intelligence Graph [Back to TOC](#)

A large dataset is not defensible when its use is disputed, stale or impossible to update. The moat begins when every material record has provenance and permitted use, merchants can correct it, authorised feeds update it, and historical changes remain auditable. A competitor may scrape public pages, but recreating contracts, field evidence, correction history and source-specific confidence is slower and legally riskier.

118. Zimbabwean aliases and commercial language [Back to TOC](#)

Search performance improves as Shoppage learns how Zimbabwean buyers, installers and merchants describe products: abbreviated model names, local unit conventions, brand nicknames, spelling variants, mixed Shona-English or Ndebele-English queries, market-area references and problem statements such as “magetsi backup for two fridges.” These mappings require local query and outcome evidence; a global catalogue cannot supply them reliably.

119. Compatibility graph [Back to TOC](#)

Compatibility is a strong wedge moat because mistakes are costly and local merchants often sell partial systems. Edges should cover voltage, current, connector, capacity, size, material, application, accessory and installation constraints. Each edge records evidence and review status. The graph improves search, alternatives, bundles and request matching without requiring Shoppage to sell the product.

For example, a buyer comparing a 24V battery should see compatible inverter classes, required cable and protection considerations, and local destinations with the relevant components. The system must distinguish engineering guidance from merchant-specific installation advice and include appropriate safety disclaimers.

120. Offer-freshness history [Back to TOC](#)

Competitors can copy a current price; they cannot instantly copy a merchant’s six-month record of update discipline, stock confirmations, unavailable outcomes and response behaviour. Freshness history influences ranking and merchant recommendations. It also identifies merchants whose advertised prices remain stable, whose feeds repeatedly break and whose manual confirmations are reliable.

121. Demand and zero-result graph [Back to TOC](#)

Every unresolved search, saved need and structured request teaches Shoppage where commerce is failing. A conventional marketplace sees only demand inside its catalogue. Shoppage sees demand across destinations and

therefore can identify missing suppliers, unmatched aliases, underserved areas and categories where a marketplace or distributor should add stock.

122. Referral-outcome graph [Back to TOC](#)

The outcome graph connects a query to comparison, destination, response and resolution with confidence levels. Over time, Shoppage learns which destinations actually respond to which needs, which offers are frequently unavailable and which channels generate useful demand. This is more difficult to reproduce than clicks because it requires partner callbacks, merchant cooperation, buyer confirmation and audit controls.

123. Physical-market topology [Back to TOC](#)

The hierarchy of market, zone, aisle, building, cooperative and stall is a locally acquired asset. It helps a buyer navigate commerce that is invisible to standard ecommerce feeds. The topology must be maintained through operator records, merchant claims and field evidence. The cost of keeping it current creates defensibility only when the resulting pages resolve demand and attract repeat use.

124. Marketplace federation network [Back to TOC](#)

Each marketplace integration improves the value of the comparison layer for buyers and makes Shoppage more relevant to the next partner. Partners benefit from shared canonical matching, feed diagnostics and demand insight. The network effect is not automatic: Shoppage must prevent confidential-data leakage and show that participation produces incremental traffic rather than merely exposing price differences.

125. Agency maintenance network [Back to TOC](#)

Agencies become a distributed operations network when Shoppage is embedded in their client reporting, content workflow and lead management. They lower merchant-acquisition and freshness costs while giving clients a reason to retain the agency. This becomes a switching cost only after agencies rely on historical intelligence and recurring client workflows, not merely after they complete certification.

126. QR and fifth-icon distribution infrastructure [Back to TOC](#)

Physical QR stickers, invoices, email signatures, website icons and product cards create direct entry points that do not depend on an algorithmic social feed. Each deployment is inexpensive, but thousands of active deployments create an owned distribution layer. The metric is not stickers printed; it is attributable sessions and merchant retention produced by installed assets.

127. Trust and correction history [Back to TOC](#)

A durable trust asset includes evidence, expiry, disputes, response, appeals and corrections. Merchants benefit when a long positive history follows them across campaigns and agency relationships. Buyers benefit because Shoppage can distinguish a newly claimed phone number from a destination with consistent offer and complaint behaviour. The system must allow rehabilitation and avoid permanent opaque scoring.

128. Data-quality operating system [Back to TOC](#)

Entity-resolution rules, reviewer guidance, sampling plans, rollback tools and category-specific validation are not glamorous, but they compound. A competitor can buy search technology; it still has to learn that two near-identical inverter names are different variants, that a district name is used inconsistently, or that an imported image is not evidence of local stock. Quality operations become a moat when error rates fall as scale rises.

129. Merchant workflow embedding [Back to TOC](#)

Quick freshness confirmations, lead-state updates, website plugins, agency reports and export cards make Shoppage part of daily commercial work. The goal is not to lock data hostage. Merchants retain export and correction rights, but they lose useful history, demand insight and distribution efficiency if they stop maintaining their profile. That is a legitimate product switching cost.

130. Independent identity and canonical links [Back to TOC](#)

Canonical Shoppage IDs survive changes in telephone numbers, social accounts, websites and marketplace listing IDs. Universal Links continue to resolve after campaigns are copied or destinations change. This independence allows accumulated trust, demand and product relationships to persist across volatile external platforms.

131. Price and availability time series [Back to TOC](#)

Price history is valuable for buyers, merchants, distributors and agencies when it is rights-cleared and confidence-labelled. It supports alerts, seasonality, suspicious-change detection and negotiation context. The data should not be sold as a definitive inflation index unless sampling, weighting and methodology support that use.

132. Request liquidity [Back to TOC](#)

A Request Network compounds when merchants know useful requests appear there and buyers know a missing product can still be found. Liquidity should be measured by response quality and resolution, not broadcast volume. Requests are routed to a small qualified set first; wider distribution follows only if the need remains unresolved.

133. Partner feed infrastructure [Back to TOC](#)

Reusable adapters, validation rules and mapping tools reduce the time to onboard the next retailer or marketplace. The infrastructure should accept common fields from Google-style product feeds, WooCommerce and simple spreadsheets while preserving Shoppage's own canonical model. Each adapter must have a named owner, test fixtures and rollback plan.

134. Features that are not moats [Back to TOC](#)

- A green badge without evidence history.
- A vertical video feed without exclusive useful content or demand.
- A generic AI chatbot.
- A WordPress plugin that only embeds a link.
- A large count of unclaimed profiles or reference products.
- A referral programme that pays before fraud and attribution controls.
- A dashboard showing metrics competitors can calculate from public data.
- A mobile app whose functions are available in the web product.

135. Moat scorecard [Back to TOC](#)

Asset	Month-12 proof	Failure signal
Rights-cleared graph	90% of active offers from contracted/merchant sources	Takedowns or unknown provenance
Freshness history	Six months of offer and outcome states	High unavailable rate despite "fresh" labels
Marketplace network	15 active partners with recurring feeds	Partners sign but do not maintain feeds
Demand graph	Meaningful zero-result and request resolution loops	Mostly navigational or bot queries

Asset	Month-12 proof	Failure signal
Agency network	40 agencies with retained active clients	Certifications without client usage
Physical distribution	35% Owned Demand Share	Traffic remains dependent on one platform
Compatibility graph	Higher resolution and fewer wrong-product complaints	Edges unused or unsafe

136. Defence against bulk extraction [Back to TOC](#)

Shoppage should provide useful public pages without exposing the entire intelligence graph as a free bulk dataset. Controls include scoped APIs, quotas, anomaly detection, cache and crawler policies, non-sequential identifiers, field-level permissions, canary records and contractual enforcement. Public facts cannot be made “owned” merely by aggregation, so the strongest defence remains service quality, update relationships and historical evidence rather than technical obstruction alone.

137. Regional replication [Back to TOC](#)

The model can extend beyond Zimbabwe because fragmented commerce, social selling and mixed merchant maturity exist elsewhere. However, the graph is local by design. A second country requires its own source rights, place topology, language, payment-context labels, marketplace relationships and field network. Software reuse is high; data and market-formation reuse is not.

Part XI - Detailed product and operating rules [Back to TOC](#)

Translate the strategy into deterministic workflows that engineering, data, trust and commercial teams can execute.

138. Search ranking feature stack [Back to TOC](#)

Ranking begins with eligibility. An offer is excluded for prohibited goods, expired freshness, broken destination, unresolved identity risk or poor product match. Eligible offers receive feature groups: lexical and semantic query relevance; canonical variant confidence; distance and place relevance; freshness; destination response and accuracy; trust evidence; price context; preferred fulfilment route; and diversity. The system prevents one marketplace or merchant from occupying every visible position when equivalent alternatives exist.

Initial weights are deterministic and reviewed weekly. Learning-to-rank may optimise only after the label set contains sufficient resolved needs and after fairness, sponsor separation and local-availability hallucination tests pass.

139. Sponsored and organic separation [Back to TOC](#)

Sponsored inventory has a separate candidate generator and visible label. It cannot rescue an expired offer, suppress a correction warning or create a trust predicate. Reports separate paid clicks from organic referrals. The auction, if introduced, uses category and geography eligibility plus quality thresholds; highest bid alone does not win.

140. Offer lifecycle [Back to TOC](#)

1. Input received from merchant, agent, website, marketplace or retailer feed.
2. Rights, schema, URL, media and prohibited-product validation.
3. Canonical product and branch matching with confidence.
4. Human review where confidence or risk is below threshold.
5. Publication with source, freshness and destination type.
6. Refresh by feed, merchant confirmation or authorised agent.
7. Outcome monitoring for unavailable, wrong price or broken destination.
8. Demotion, expiry, correction, suspension or deletion with audit history.

141. Source-confidence model [Back to TOC](#)

Confidence is field-specific. An API feed may be strong for price but weak for physical branch stock. A merchant confirmation may be strong for immediate availability but weak for technical specifications. A manufacturer source may be authoritative for product dimensions but irrelevant to a Zimbabwe warranty claim. Search and public labels should use the confidence of the particular field, not a single source score.

142. Zero-result workflow [Back to TOC](#)

When no active offer matches, the system first checks aliases and variant broadening, then nearby geography, compatible alternatives, recently confirmed destinations and the Reference Library. It then offers a structured request. The query enters a zero-result queue with reason codes: no local supply, stale supply, unmatched language, catalogue gap, restricted category, location ambiguity or search defect. Data and commercial teams own different reason codes.

143. Request broadcast controls [Back to TOC](#)

A request is sent first to three to five high-fit destinations based on category, location, response history and capacity. Merchants see only the information necessary to respond. If no response arrives, the system expands the

pool or asks the buyer for clarification. Repeated low-quality broadcasts reduce merchant trust and must be treated as a product defect.

144. Duplicate and impersonation controls [Back to TOC](#)

Profiles are screened for shared telephone numbers, copied images, near-identical names, conflicting locations and abnormal claim velocity. A duplicate may be a legitimate branch, agent or marketplace seller, so the system requests evidence instead of automatically banning. Impersonation complaints receive urgent suppression of high-risk contact actions while ownership is reviewed.

145. Conflicting marketplace and merchant data [Back to TOC](#)

The same seller may provide a direct offer and sell through a marketplace at a different price or policy. Shoppage keeps separate offers and destinations. It does not overwrite one with the other. If a partner feed conflicts with a seller's direct claim, each remains attributed to its source, and public comparison shows the relevant update time and destination terms.

146. Merchant inactivity and offboarding [Back to TOC](#)

An inactive merchant's offers expire first; the profile may remain as a corroborated supplier page if rights and accuracy permit. Paid status ends without deleting historical referral records required for audit. The merchant can export its own offer and lead data. Public pages show no active stock and retain a correction or claim route.

147. Marketplace-partner offboarding [Back to TOC](#)

When a marketplace feed ends, Shoppage removes or expires current commercial claims according to the contract, disables broken deep links and updates the Marketplace Page. Permitted aggregate history may remain, but no current availability is inferred. Buyers with saved needs are notified only if another active destination exists or the request route is available.

148. Complaint triage [Back to TOC](#)

Severity	Examples	Target action
Critical	Impersonation, dangerous product, serious fraud, exposed personal data	Immediate restriction and escalation
High	Repeated bait price, wrong destination, harassment, counterfeit allegation	Review inside four hours
Standard	Stale price, incorrect hours, duplicate branch, broken link	Acknowledge within one business day
Low	Description improvement, category correction	Queue by impact and evidence

149. Correction and suppression process [Back to TOC](#)

Every public record has a correction route. Requests capture the disputed field, claimant authority, evidence and urgency. High-risk contact or identity data can be temporarily hidden before final review. Decisions record reviewer, reason and source impact. A source-level error triggers a sample review of related records rather than isolated correction only.

150. Price anomaly controls [Back to TOC](#)

Prices are checked against the merchant's history, comparable variants, currency and category ranges. An anomaly generates a warning or review; it does not automatically replace the merchant price. Extreme bargains may be genuine. The public result can show "price requires confirmation" when evidence is weak or the destination repeatedly fails to honour similar offers.

151. Creator campaign workflow [Back to TOC](#)

1. Merchant or agency selects a current offer and objective.
2. Shoppage creates a written brief, disclosure and Universal Link.
3. Creator submits draft or publishes under agreed review rules.
4. System verifies link, factual claims and live period.
5. Fixed deliverable payment is approved after quality checks.
6. Referral events are reported separately from creator payment until performance gates pass.
7. Complaints or inaccurate claims affect creator eligibility and merchant campaign status.

152. Agency service-level agreement [Back to TOC](#)

A certified agency is expected to maintain client authority records, refresh priority offers, act on lead ageing, disclose sponsorship, protect client credentials and deliver a monthly outcome report. Certification is paused when a material share of managed clients becomes stale or complaints are ignored. Shoppage audits client consent and does not permit agencies to claim ownership of merchant profiles.

153. Market-captain compensation model [Back to TOC](#)

Illustrative compensation uses small components: accepted verified claim; day-30 active merchant; batch of refreshed high-demand offers; resolved correction; and market-roster quality score. No payment is made solely for photographing a storefront or entering a telephone number. A reserve is held against fraudulent or duplicate work, and captain income is visible in management reporting as activation cost.

154. Marketplace partner economics [Back to TOC](#)

The partner evaluates incremental gross profit or strategic value from referred traffic against integration and commercial cost. Shoppage supplies an incrementality framework: new users or sessions, category demand not previously captured, valid product clicks, destination conversion where callbacks exist, feed-quality improvements and cost per useful action. The comparison does not require disclosure of total marketplace GMV.

155. Merchant value proof [Back to TOC](#)

A merchant report translates activity into conservative commercial evidence: qualified actions, responses, quotes, unavailable outcomes, buyer-resolved events and self-reported completed transactions. It states confidence and avoids multiplying clicks by an assumed conversion rate. The paid-plan conversation asks whether the demonstrated gross-profit opportunity and workflow savings justify the fee.

156. Buyer repeat loops [Back to TOC](#)

- Price and availability alerts for saved needs.
- Restock and new-destination notifications.
- Request responses and compatible alternative alerts.
- Repeat contractor or household lists.
- Market events, new stall rosters and local guides.
- Service and replacement reminders where appropriate.
- Direct QR and referral sharing from completed searches.

157. Weekly operating review [Back to TOC](#)

The weekly review starts with buyer failures: top zero-result queries, unavailable referrals, wrong-product complaints and unresolved requests. It then reviews supply freshness, partner feeds, merchant response, direct-repeat demand, paid conversion, agency performance, source-rights issues and cash. Feature output is discussed only when it changed one of those outcomes.

Part XII - Commercialisation and launch playbooks

Back to TOC

Turn the strategy into partner conversations, campaigns, staffing capacity and disciplined cash execution.

158. First partner archetypes [Back to TOC](#)

The first partner list should balance catalogue breadth and integration learning: a national marketplace; a classifieds platform or large category publisher; a formal electronics or power retailer; a hardware distributor; a WooCommerce merchant; a solar installer; a building-material supplier; a market operator; an agency with 15-plus clients; and a diaspora-focused ecommerce destination. Diversity tests the federation model better than ten similar merchants.

159. Merchant acquisition offer [Back to TOC](#)

The launch offer is: “We will create your Shoppage identity, load your top high-demand offers, connect your existing website or WhatsApp, show you local demand and send you attributable enquiries. You keep the sale and customer relationship.” The merchant commits to accurate claims, fast confirmation and outcome marking. The free proof period has a defined review date.

160. Marketplace-partner pitch [Back to TOC](#)

The pitch is: “Shoppage is not asking you to move checkout. We index your authorised products beside other destinations, optimise them for Zimbabwean search language and compatibility, and send product-level traffic with transparent reporting. You receive feed diagnostics and demand gaps. After the pilot, we agree a subscription, referral or intelligence model only if the value is measurable.”

161. Agency pitch [Back to TOC](#)

The agency proposition is: “Replace static social reports with live commerce intelligence. Maintain client stock truth, show competitor and demand signals, create trackable content, and report leads and response. Shoppage gives you a recurring tool and a differentiated managed service.” Agencies should be recruited from the target cells, not only from prominent national directories.

162. Market-operator pitch [Back to TOC](#)

The operator proposition is visibility and navigation: a verified market page, searchable roster, category map, campaign links, correction route and aggregated demand report. Shoppage does not claim operator authority, collect rent or decide stall rights. Operator evidence is one source in the trust model.

163. Buyer launch campaign [Back to TOC](#)

The buyer campaign should demonstrate a hard task rather than announce a platform. Examples: “Compare 24V lithium batteries available today in Harare,” “Find three roofing suppliers near your project,” or “Send one solar request and receive matched destinations.” Each campaign has a cell page, measurable query set and request fallback.

164. Content operating calendar [Back to TOC](#)

Weekly rhythm	Content	Commercial purpose
Monday	Demand gaps and new requests	Activate merchant supply

Weekly rhythm	Content	Commercial purpose
Tuesday	Compatibility or buying guide	Improve search confidence
Wednesday	Market/merchant walk-through	Build place and trust evidence
Thursday	Price Pulse or restock alert	Drive repeat demand
Friday	Showcase Shorts and creator campaign	Distribute current offers
Weekend	Market guide and local collection options	Capture physical-shopping intent

165. SEO launch architecture [Back to TOC](#)

The first indexable set should be controlled: 300 canonical product/need pages, three cell pages, selected market pages, claimed merchant and marketplace pages, and 30 high-quality guides. Pages without fresh offers may remain noindex until they provide unique reference or request value. Canonicalisation prevents one product from producing dozens of thin location duplicates.

166. Diaspora pilot [Back to TOC](#)

A diaspora pilot focuses on one use case, such as solar equipment or building supplies for a Harare project. The user requests or compares local destinations, then completes the sale with the selected merchant or marketplace. The pilot measures destination confidence, recipient coordination and proof needs without Shoppage receiving funds.

167. Pricing experiments [Back to TOC](#)

1. Test free versus US\$19 Growth after merchants receive comparable action volume.
2. Test monthly subscription versus prepaid referral credits.
3. Test agency wholesale price against client count and advanced-intelligence slots.
4. Test marketplace subscription versus CPC using the same pilot evidence.
5. Test sponsor packages by category and cell with strict organic separation.
6. Test managed onboarding price against time, merchant activation and payback.

168. Channel attribution standards [Back to TOC](#)

Every campaign uses source, medium, campaign, content, product/market and partner identifiers. Offline QR has unique asset IDs. Agencies and creators use assigned links. Direct traffic is not automatically attributed to the last paid campaign. Reports include unattributed and low-confidence activity rather than forcing every session into a marketing story.

169. Commercial pipeline stages [Back to TOC](#)

Stage	Required evidence	Exit condition
Target	Fits cell/category and has useful supply or audience	Named decision maker

Stage	Required evidence	Exit condition
Discovery	Current channels, data, pain and rights understood	Qualified use case
Pilot design	Scope, feed/offers, roles and metrics agreed	Written pilot plan
Activated	Data live and destination tested	First valid referral
Value review	Traffic, insight, labour and economics reviewed	Commercial choice
Paid/partner	Contract and operating SLA active	Recurring review

170. Onboarding-capacity model [Back to TOC](#)

Capacity must be planned by onboarding lane. One trained agent may complete several simple claim-and-offer activations per day but only a few catalogue clean-ups. API and marketplace integrations consume engineering and data time. Management should report hours per activated merchant, offers accepted per hour, day-30 active rate and cost by lane. Volume targets are reduced when quality queues exceed SLA.

171. Staffing-capacity assumptions [Back to TOC](#)

The base team should be able to support one cell before adding headcount for national ambition. A commerce data lead with two reviewers can maintain a constrained wedge; an activation lead with agencies and captains can reach merchants; a trust lead can manage early claims and complaints; and one senior engineering lead can coordinate a modular build. Headcount gates depend on active supply and queue load, not planned profile count.

172. Cash-control rules [Back to TOC](#)

- Thirteen-week rolling cash forecast updated weekly.
- Founder, software, field activation and media costs reported separately.
- No annual software commitment without clear cancellation and owner.
- Partner and creator payments require accepted deliverables and audit trail.
- Campaign budgets stop automatically when destination quality falls.
- Capital expenditure and regulated-service experiments require board approval.
- Monthly actual-versus-plan review includes runway and committed liabilities.

173. Vendor selection [Back to TOC](#)

Technology vendors are assessed on data location, security, API portability, pricing at target volume, support, export, termination and dependency risk. AI providers must support model substitution and logging controls. Messaging providers must not become the only account-recovery path. No vendor is selected solely for a low introductory price.

174. Data quality sampling [Back to TOC](#)

Each week, reviewers sample active offers by source, merchant, category and freshness. They test product match, price, currency, destination, branch and public label. Requests and complaints oversample risky sources. Results produce source-level corrective actions and ranking changes. A high-volume feed with poor sample accuracy is paused even if it increases headline catalogue size.

175. Launch-day war room [Back to TOC](#)

Public launch requires named owners for search, feeds, Universal Links, claims, complaints, security, partner escalation and communications. Dashboards show errors, zero results, referral failures and abusive activity. Non-essential feature releases are frozen. A rollback can remove a source, category or campaign without taking down the whole service.

176. Post-launch triage [Back to TOC](#)

The first two weeks prioritise wrong results and dead ends over cosmetic improvements. Severity is based on buyer harm and demand volume. Search gaps lead to catalogue or alias work; unavailable offers lead to freshness changes; partner failures lead to feed or deep-link fixes; low merchant response leads to supply replacement or request-routing changes.

177. Metric definitions that prevent self-deception [Back to TOC](#)

- Profiles are not active merchants.
- Reference products are not available inventory.
- Feed records are not valid offers until processed.
- Clicks are not leads unless the product is sold as CPC.
- WhatsApp opens are not sales.
- Merchant-marked sales are not independently confirmed transactions.
- Group memberships are not unique reachable users.
- Page views are not qualified demand without intent criteria.
- Gross revenue is not software margin when activation services are included.

178. Investment narrative [Back to TOC](#)

The investment case is not that Zimbabwe lacks marketplaces. It is that commerce supply is fragmented across marketplaces, websites, social channels and physical markets, and no neutral local layer has organised product identity, live offers, compatibility, place, trust, demand and attribution across them. Shoppage can become the connective tissue and sell distribution plus intelligence without financing inventory or payments.

179. Board decision calendar [Back to TOC](#)

Timing	Decision	Evidence required
Day 30	Approve build and one-cell pilot	Rights audit, interviews, prototype and partners
Day 90	Approve public launch	Six gates and closed-test results
Month 6	Approve second cell and paid referrals	Resolution, retention and attribution economics
Month 12	Approve scale capital	Repeatable cell, partner revenue and payback
Month 18	Approve category or country expansion	Core profitability and operational capacity

180. Detailed 24-month roadmap [Back to TOC](#)

Months 0-3 establish rights, graph schema, product matching, one cell and closed demand tests. Months 4-6 stabilise freshness, requests, merchant response and initial monetisation. Months 7-12 replicate two cells, grow marketplace and agency partners, and introduce feed intelligence and controlled sponsored distribution. Months 13-18 deepen the strongest categories and build first-party plugins and partner APIs. Months 19-24 choose one expansion vector and prepare a separately costed regional or adjacent-category plan.

Appendices [Back to TOC](#)

Operating templates, commercial controls and evidence registers.

Appendix A - Binding decision register [Back to TOC](#)

Decision	Current position	Change authority
Category	Commerce intelligence, comparison and referral OS	Board/founder decision memo
Checkout	Not in core product	Separate legal and investment approval
Marketplaces	First-class supply partners and destinations	Commercial policy
Default search	Fresh local offers first	Product governance
Reference data	Separate library, no availability implication	Data governance
Ranking	Organic relevance independent of payment	Ranking committee
Launch geography	Harare cells before national expansion	Gate review
Launch wedge	Home, Power and Building	Cell evidence review
Core monetisation	Subscription, referral, sponsor, agency, data	Commercial review
Payments/fulfilment	Destination responsibility	Separate product approval

Appendix B - Minimum product and offer schema [Back to TOC](#)

Object	Required fields	Notes
Product variant	Canonical ID, title, brand/model, category, attributes, aliases	No seller price in master
Offer	Destination, branch, price/currency, availability state, update time, URL/action	Source owns current commercial truth

Object	Required fields	Notes
Merchant	Canonical ID, public name, contacts, branches, destination type	Claim and evidence state separate
Marketplace	Partner ID, seller model, checkout terms URL, feed state	Partner-level trust separate from seller
Place	Country, city, area, landmark, pin, market hierarchy	Evidence for containment links
Request	Need, quantity, location, urgency, contact consent	Private by default
Referral event	Source, product/offer, destination, action, timestamp, campaign	No sale inference from click
Trust evidence	Predicate, source, reviewer, date, expiry, challenge state	Never one undifferentiated badge

Appendix C - Freshness matrix [Back to TOC](#)

Offer class	Default freshness window	Expiry action
Fast-moving manual stock	24-72 hours	Demote then hide
Marketplace API stock	Minutes-hours by SLA	Flag feed failure
Price-only catalogue offer	Up to 7 days	Show “confirm availability”
Made-to-order product	14-30 days	Keep quote-required label
Service capability	30-90 days	Require periodic active confirmation
Market event offer	Event-specific	Remove at event close
Reference product	Not applicable	Never shown as active offer

Appendix D - Source Rights Register template [Back to TOC](#)

Field	Example value	Required control
Source owner	Marketplace Partner A	Named legal entity

Field	Example value	Required control
Access method	Signed JSON feed	No unauthorised crawl
Public display	Title, image, price, destination URL	Field-level permission
AI use	Extraction and matching only	No model training unless granted
Retention	Raw feed 90 days; history aggregated	Deletion schedule
Attribution	Partner brand and destination link	Public requirement
Suppression	Webhook/email; 24-hour urgent SLA	Logged action
Termination	30 days; immediate for abuse	Post-termination data rule

Appendix E - Merchant onboarding SOP [Back to TOC](#)

1. Create or locate candidate profile; show source and correction rights.
2. Verify claimant phone/email and authority for the business or branch.
3. Capture destination type, locations, categories and transaction channels.
4. Select ingestion method and confirm rights to product content.
5. Publish a small high-demand offer set first; validate matching and freshness.
6. Install Universal Links, QR and optional Live Seal.
7. Test one buyer action and merchant response end to end.
8. Train on quick freshness, lead states and complaints.
9. Review at day 7 and day 30; deactivate empty or abandoned claims.

Appendix F - Marketplace partner pilot proposal [Back to TOC](#)

Objective: index an authorised category feed, match products to the Shoppage graph, create a Marketplace Page and send measured incremental traffic while the partner retains checkout and transaction control.

Pilot element	Proposed scope	Success measure
Catalogue	500-2,000 offers in one wedge category	>85% valid import; >70% canonical match
Refresh	Daily feed or incremental API	<5% stale/broken offer rate
Referral	Product-level deep links with partner IDs	>90% destination success
Analytics	Queries, clicks, zero-result and feed errors	Partner identifies actionable insight

Pilot element	Proposed scope	Success measure
Term	Six weeks plus review	Continue only with measurable value
Commercial	Free proof period; CPC/data/subscription options	Mutually agreed post-pilot economics

Appendix G - Agency operating pack [Back to TOC](#)

- Client portfolio health and delegated permissions.
- Weekly Freshness, Response and Demand report.
- Price Pulse with source-confidence labels.
- Zero-result and Request Network campaign builder.
- Export cards, Live Seal, Shorts and email-signature assets.
- Lead ageing and merchant-response workflow.
- Certification audit and client complaint route.

Appendix H - Market captain SOP [Back to TOC](#)

1. Operate only in assigned territory with visible identification.
2. Obtain merchant consent before claim, media capture or offer publication.
3. Use evidence checklist for stall, zone and operator relationships.
4. Never collect unofficial verification or listing fees.
5. Submit work through named account and device; no shared credentials.
6. Respond to correction or complaint requests within SLA.
7. Accept random rechecks and payment reversals for fraudulent work.
8. Report conflicts, political/operator pressure and safety incidents.

Appendix I - Trust evidence dictionary [Back to TOC](#)

Predicate	Evidence	Expiry/limitation
Phone verified	OTP to merchant-controlled number	Does not prove identity or stock
Location reviewed	Photos, map, roster or visit	Recheck after move/complaint
Marketplace seller active	Partner feed or seller-status callback	Partner-specific only
Offer refreshed today	Merchant/feed timestamp	No guarantee it remains in stock
Responsive merchant	Median response from attributable leads	Requires minimum sample
Complaint resolved	Documented closure and buyer/partner evidence	Not proof of no other complaints

Predicate	Evidence	Expiry/limitation
Company evidence supplied	Registration/tax or authorised record	Does not guarantee service quality

Appendix J - Attribution event taxonomy [Back to TOC](#)

Event	Meaning	Commercial confidence
outbound_click	User leaves for exact destination	CPC eligible after validation
whatsapp_start	Prefilled chat link opened	Action, not response or sale
call_reveal	Phone action initiated	Action, not connected call
quote_submitted	Structured request delivered	Higher intent; consent required
merchant_responded	Destination records response	Medium; sampled
destination_callback	Partner returns event	High if signed and defined
buyer_resolved	Buyer confirms suitable solution	High but subject to fraud controls
purchase_confirmed	Supported transaction evidence	Only where partner/data rights permit

Appendix K - Partner feed fields [Back to TOC](#)

Field group	Core fields	Validation
Identity	partner_offer_id, seller_id, destination_url	Unique and stable
Product	title, brand, model, GTIN/MPN, category, attributes	Canonical match confidence
Commercial	price, currency, sale period, availability	Timestamp and meaning defined
Location	branch, city, collection, delivery claim	No implied national coverage
Media	image URLs, description, rights flag	Malware, licence and quality

Field group	Core fields	Validation
Policy	seller type, return-policy URL, protection type	Public representation agreed
Control	updated_at, delete flag, source version	Incremental history and rollback

Appendix L - Financial assumptions register [Back to TOC](#)

Assumption	Base value	Required test
Growth price	US\$19/month	Merchant interviews and paid retention
Agency price	US\$149/month	Five-agency pilot
Paid merchant share month 12	32% of active merchants	Cohort conversion
Monthly merchant churn	<5% paid cohort	Three-month evidence
Activation cost	<US\$25 self/agency; <US\$60 assisted	Channel accounting
Referral gross margin	>75% before field labour	Event and infrastructure costs
Owned Demand Share month 12	35%	Attribution source mix
Base month-12 revenue	US\$36,660	Not investor promise

Appendix M - Extended risk register [Back to TOC](#)

Risk scenario	Response trigger	Contingency
Marketplace terminates feed	Notice or feed shutdown	Remove current claims, retain permitted aggregate history, replace coverage
Merchant refuses price transparency	High anchor attrition	Allow quote/confirm states; do not fake comparison
Facebook restricts distribution	Traffic loss >20%	SEO, QR, saved needs, agencies and direct referrals

Risk scenario	Response trigger	Contingency
Currency disruption	Rapid divergence or policy change	Show source currency/time; suspend derived conversions
Mass stale inventory	Unavailable rate >20%	Shorten expiry, request mode, reduce public coverage
Data rights complaint	Credible claim received	Urgent suppression, preserve evidence, legal review
AI matching error cluster	Wrong-variant rate exceeds threshold	Rollback model, human review and test-set expansion
Field safety problem	Threat or incident	Suspend territory, remote onboarding, partner escalation

Appendix N - Experiment backlog [Back to TOC](#)

Experiment	Primary metric	Decision threshold
Search vs WhatsApp baseline	Time to suitable destination	Shoppage 30% faster
Two vs three fresh offers	Resolved Demand Rate	Meaningful lift justifies density cost
Live Seal installation	Direct referrals and merchant retention	Incremental attributable actions
Request broadcast size	Response quality and spam	Best resolution with limited recipients
Price Pulse report	Agency client retention	Clients act on at least one insight/month
Freshness reminder channel	Valid confirmations per cost	Choose web, WhatsApp or agent mix
Marketplace feed pilot	Destination success and partner value	Partner renews or pays

Appendix O - Public launch scorecard [Back to TOC](#)

Gate	Pass condition	Owner
Rights	100% public sources classified; no critical unresolved claims	Data/legal
Density	70% top 300 entities covered by two offers or request path	Commerce data
Demand superiority	Closed tests beat normal discovery process	Product
Attribution	90% valid actions retain required identity	Technology
Economics	At least 20 merchants/partners state measurable willingness to pay	Commercial
Security/privacy	Independent critical findings closed	Trust/legal

Appendix P - Competitive positioning matrix [Back to TOC](#)

Model	Owns transaction?	Shoppage relationship
Social commerce group/status	Usually no	Distribution source and merchant action channel
Classifieds	Usually buyer-seller contract	Indexed destination where authorised
Multi-vendor marketplace	Yes or orchestrates seller checkout	Marketplace partner and offer source
Retailer website	Merchant owns sale	Direct destination and feed partner
Physical market/stall	Merchant owns sale	Place/stall digital twin and referral
Price-comparison service	No, usually referral	Closest category analogue
Shoppage	No in core model	Neutral Zimbabwe commerce intelligence and referral layer

Appendix Q - Neutral ranking constitution [Back to TOC](#)

Purpose. Shoppage can only become a trusted comparison and referral layer if organic ranking is useful to buyers and commercially neutral. The ranking constitution is a binding product-control document, not marketing language. It applies to merchant sites, marketplace offers, retailer websites, physical stalls and Shoppage-sponsored inventory. No destination type receives an automatic preference merely because it is a marketplace, a formal retailer, an informal trader, a paying subscriber or a strategic partner.

Q1. Organic ranking objective [Back to TOC](#)

The organic system should maximise the probability that a buyer reaches a suitable, reachable and accurately represented destination. Suitability includes product identity, variant compatibility, location, price context, availability confidence, seller responsiveness and the buyer's stated fulfilment preference. Revenue is excluded from the organic relevance score. Commercial products may buy clearly labelled exposure, but they may not silently modify the organic score.

Q2. Ranking inputs [Back to TOC](#)

Input family	Permitted signals	Prohibited shortcuts
Query match	Canonical product, attributes, aliases, language and intent	Keyword stuffing or paid keyword preference
Offer quality	Freshness, price clarity, stock state, destination validity	Assuming feed volume means quality
Location	Branch, market, delivery claim and buyer preference	False national coverage from one address
Trust	Predicate-specific evidence, complaint handling and response history	One opaque universal trust score
Compatibility	Validated product edges and required accessories	AI-generated compatibility without evidence
User context	Explicit location, saved need and current-session choices	Sensitive profiling or undisclosed behavioural manipulation
Diversity	Useful seller and destination alternatives	Artificially filling results with weak offers

Q3. Destination neutrality [Back to TOC](#)

A marketplace offer and a direct merchant offer compete on the same buyer-value principles. A marketplace may rank higher because it provides a fresh price, reliable deep link, clear seller identity and stronger protection. A merchant may rank higher because it has verified nearby stock, better response history or a more suitable collection option. Shoppage must explain the material basis for visible ranking labels such as “freshest nearby offer,” “best confirmed price,” or “fastest responding destination.” It must not claim “best” where the evidence only supports “lowest displayed price.”

Q4. Sponsored placement [Back to TOC](#)

Sponsored results must be visually and programmatically labelled. They should meet minimum relevance, destination, freshness and safety thresholds before bidding or fixed sponsorship can apply. Payment buys placement inside a qualified sponsored surface; it does not buy verification, organic rank, compatibility claims or suppression of competitors. The organic result should remain available adjacent to or immediately after the sponsored unit so the buyer can compare.

Q5. Ranking governance [Back to TOC](#)

1. Publish a plain-language description of the principal ranking factors and sponsorship treatment.
2. Maintain an internal versioned ranking specification, feature dictionary and change log.
3. Run offline evaluation on Zimbabwe-specific query sets before material ranking releases.
4. Monitor destination-type exposure, partner concentration and unexplained rank shifts.
5. Provide a merchant and partner challenge route for factual errors, not a right to demand rank.
6. Require conflict disclosure from staff who can change ranking or approve sponsorship.
7. Escalate major changes affecting competition, sensitive categories or public trust to a ranking committee.
8. Preserve rollback capability and evidence for every production ranking version.

Q6. Anti-gaming controls [Back to TOC](#)

Gaming indicators include repeated false freshness confirmations, rotating destination URLs, duplicate sellers, copied images, misleading price fragments, fake branches, coordinated clicks, complaint intimidation and partner feeds that omit mandatory charges. Responses range from confidence reduction and offer expiry to partner-feed suspension and public enforcement labels. High-volume or paying partners receive no exemption. A partner may correct a feed rapidly, but repeated material misrepresentation changes source confidence and commercial eligibility.

Q7. Ranking audit scorecard [Back to TOC](#)

Audit question	Minimum evidence	Failure response
Does organic rank exclude payment?	Feature review and revenue-field test	Release blocked
Are sponsored units unmistakable?	Visual and accessibility QA	Campaign paused
Do top results resolve demand?	Sampled buyer outcomes	Reweight or reduce coverage
Are destination types treated consistently?	Exposure and quality analysis	Bias investigation
Can material changes be rolled back?	Version, model and configuration archive	Deployment blocked
Can merchants challenge factual inputs?	Logged correction route and SLA	Governance remediation

Appendix R - Marketplace and retailer listing agreement term sheet [Back to TOC](#)

This term sheet defines the commercial and data relationship for a marketplace, retailer, distributor or other catalogue destination that lists authorised offers on Shoppage. It is deliberately structured so the partner retains checkout, customer contract, payment, fulfilment, returns and seller-management responsibilities. Shoppage

supplies discovery, comparison, referral attribution, feed diagnostics and optional intelligence. Final agreements require local legal review and category-specific schedules.

R1. Service scope [Back to TOC](#)

- Ingest agreed product, seller, branch, price, availability, media, policy and destination fields through a file, API or approved connector.
- Match partner records to the Shoppage canonical catalogue while preserving the partner's own identifiers.
- Display authorised fields in search, comparison, Marketplace Pages, merchant pages and relevant market/category surfaces.
- Route users to the partner's deep link, app link, telephone, WhatsApp or other agreed destination.
- Provide reports on queries, referrals, broken destinations, unmatched products, stale records and unmet demand.
- Offer optional sponsored distribution or intelligence under separate pricing and labelling schedules.

R2. Responsibilities and transaction boundary [Back to TOC](#)

Party	Core responsibility	Express exclusion
Partner	Seller admission, offer truth supplied, checkout, payment, tax, fulfilment, returns and buyer contract	Shoppage does not become merchant of record
Shoppage	Processing, matching, public representation, referral routing, correction and reporting	No guarantee of sale or partner performance
Both	Security, lawful data sharing, incident response and agreed consumer communication	No implied joint venture or payment agency

The public experience must state where the buyer will complete the transaction and which destination's terms apply. A referral to a partner must not be represented as a Shoppage sale. Where a partner returns a purchase or conversion callback, Shoppage may use it only for the purposes and retention period stated in the agreement and consent model.

R3. Data rights schedule [Back to TOC](#)

The agreement should enumerate fields rather than use broad language such as "all catalogue data." For each field, the schedule states whether Shoppage may ingest, transform, publicly display, retain historically, aggregate, use for analytics, use for AI-assisted matching, share with the partner's sellers, or delete on termination. Product facts and brand assets may have different rights from seller data, prices and behavioural events. A technical ability to collect a field is not permission to use it.

R4. Feed service level [Back to TOC](#)

Control	Illustrative commitment	Remedy
Refresh frequency	Incremental updates hourly or daily by category	Freshness label, demotion or pause

Control	Illustrative commitment	Remedy
Delete/withdrawal	Delete flags processed within agreed period	Urgent suppression route
Destination success	Deep-link validity above agreed threshold	Redirect repair or feed suspension
Mandatory charges	Price representation includes disclosed compulsory charges	Correction and source-confidence review
Incident notification	Named contacts and defined security window	Joint response and temporary isolation
Schema change	Advance notice and test fixture	Versioned adapter and rollback

R5. Brand and seller presentation [Back to TOC](#)

Shoppage may identify the partner as the transaction destination and may display seller identity where authorised. Partner branding must not obscure seller responsibility or make sponsored rank appear organic. Shoppage controls its comparison layout, ranking labels and correction interface. The partner may require accurate trademark usage and may challenge factual errors, but it cannot require the removal of lawful competing offers merely because comparison is commercially uncomfortable.

R6. Attribution and reporting [Back to TOC](#)

Referral identity should use stable Shoppage event IDs and partner campaign parameters without exposing unrelated user data. Reports distinguish impressions, outbound clicks, WhatsApp starts, quote submissions, destination callbacks and supported purchase confirmations. The parties agree that a click is not a sale. CPC billing requires destination and fraud validation; CPA or revenue-share billing requires a signed event definition, duplicate rules, reversal window, reconciliation process and dispute evidence.

R7. Commercial structures [Back to TOC](#)

Model	Best use	Key control
Free proof pilot	Demonstrate incremental discovery and feed quality	Defined term and success criteria
Partner subscription	Marketplace Page, diagnostics and reporting	Feature and usage schedule
Validated CPC	High-intent outbound referral	Fraud, failed-link and self-click exclusions
Sponsored category/cell	Clearly labelled partner exposure	Relevance and concentration limits
Intelligence licence	Demand, pricing and feed-quality insight	Aggregation and confidentiality thresholds

Model	Best use	Key control
CPA/revenue share	Only with reliable transaction callback	Reversals and independent reconciliation

R8. Competition, neutrality and non-exclusivity [Back to TOC](#)

The standard agreement is non-exclusive. Shoppage may list competing marketplaces, direct retailers, informal merchants and physical stalls. The partner may use other comparison and acquisition services. Any temporary exclusivity must be narrow, paid, time-limited and approved against the neutrality doctrine. No partner may purchase a contractual right to manipulate organic ranking, suppress competitors, receive confidential competitor data or prevent a merchant from claiming its own profile.

R9. Consumer complaints and corrections [Back to TOC](#)

Shoppage receives complaints about inaccurate representation, dead destinations and referral conduct; the partner retains transaction complaints under its own terms. The agreement defines handoff, acknowledgement and evidence exchange. Shoppage may preserve a public correction or enforcement state where necessary to prevent recurring harm. Urgent claims involving impersonation, security, unlawful content or material price deception use a faster path than ordinary catalogue disputes.

R10. Termination and continuity [Back to TOC](#)

On termination, current partner offers are removed or converted to non-current reference records only where rights permit. Deep links are disabled or redirected transparently. Shoppage returns or deletes partner-confidential data under the schedule, while retaining legally required logs and permitted aggregated performance history. Canonical product identities and independently sourced facts are not automatically deleted merely because one feed terminates. Buyer-facing pages must never continue to imply live partner availability after the feed ends.

Appendix S - Commercial sensitivity, capital discipline and stop rules [Back to TOC](#)

The purpose of the financial model is to test whether the referral and intelligence system can finance its data quality, support and distribution work without hiding a field-services company inside a software gross margin. Revenue is built from multiple customers - merchants, agencies, marketplaces, sponsors and intelligence buyers - but each line must demonstrate incremental value. The model should be updated with cohort evidence monthly and formally rebased at each board gate.

S1. Revenue-engine sensitivities [Back to TOC](#)

Driver	Downside test	Management response
Paid merchant conversion	Below 15% after demonstrated leads	Reduce feature scope; test lead credits or agency bundle
Merchant churn	Above 8% monthly	Diagnose unresolved demand, service burden and unclear value
Agency productivity	Fewer than 10 maintained clients per agency	Simplify workflow or change wholesale economics
Marketplace CPC yield	Below infrastructure and support contribution	Use subscription/data model or terminate billing

Driver	Downside test	Management response
Sponsored demand	Concentrated in fewer than three buyers	Cap exposure and avoid fixed-cost expansion
Field activation cost	Above US\$60 without retention uplift	Shift to agencies, operators and remote claim flows
Resolved Demand Rate	Below 20% in dense launch set	Stop geographic expansion and repair supply quality

S2. Contribution model [Back to TOC](#)

Contribution must be calculated separately for merchant software, managed activation, agency wholesale, marketplace referrals, sponsorship, data products and creator campaigns. Direct costs include payment collection fees, messaging, search and AI usage, media delivery, customer support, field bounties, agency commissions, creator payments, feed engineering and complaint handling. Central product salaries are operating expenses, but category-specific manual work that scales with accounts should not be disguised as platform overhead.

S3. Illustrative operating cases [Back to TOC](#)

Case	Month-12 evidence profile	Decision
Downside	Weak resolution, <1,000 active merchants, <US\$12k MRR and high assisted cost	Narrow to one cell/category; suspend scale hiring
Base	Useful resolution, 2,000-3,000 active merchants, diversified US\$30k-US\$40k MRR	Replicate one proven cell at a time
Upside	Strong direct repeat, partner renewals, agency leverage and >US\$60k MRR	Fund controlled national/category expansion

These ranges are planning scenarios, not forecasts or valuation claims. The decisive variables are not total indexed records or aggregate group memberships. They are fresh offer density, resolved demand, repeat direct use, paid retention, partner renewal, source quality, activation cost and contribution after ongoing maintenance.

S4. Capital allocation order [Back to TOC](#)

1. Rights, privacy, security and source-governance controls required to operate lawfully and credibly.
2. Canonical catalogue, search relevance, freshness, Universal Links, Request Network and referral measurement.
3. One dense launch cell with category data and anchor destinations.
4. Merchant and partner activation that produces current offers and measurable buyer outcomes.
5. Agency and field tooling that lowers recurring maintenance cost.
6. Revenue products proven by willingness to pay.
7. Only then, geographic replication, adjacent categories, public APIs or regional expansion.

S5. Hiring discipline [Back to TOC](#)

The first permanent team should cover product/engineering, commerce data, partner integration, merchant operations, trust/rights and commercial activation. Specialist design, security review, legal work, field campaigns and content production can begin as controlled external engagements. Management should not hire a national field force, logistics team, payment team, large sales organisation or microservice platform group before the core referral economics are proven.

S6. Stop, pivot and double-down rules [Back to TOC](#)

Signal	Rule	Action
Demand superiority fails	Shoppage does not beat normal discovery in closed tasks	Stop public launch; fix or narrow wedge
Rights gate fails	Critical source cannot lawfully support public use	Remove source; redesign coverage
Partner model fails	Marketplaces see no incremental value after two pilots	Emphasise direct merchants/agency data or revise product
Merchant payment fails	Attributable value exists but subscriptions do not convert	Test referral credits, agency bundles and sponsor subsidy
Direct repeat grows	Owned Demand Share and saved-need usage rise consistently	Increase alert, QR and request investment
One cell reaches contribution	Repeatable activation and retention proven	Replicate the operating cell, not the national story

S7. Financing posture [Back to TOC](#)

Capital should fund evidence milestones rather than an 80-page feature list. A pre-seed or founder-funded phase should reach the rights audit, working prototype, anchor-partner pilots and closed demand test. The next tranche should fund one public cell, partner integrations and six months of cohort evidence. Scale capital is justified only when a cell demonstrates repeatable merchant activation, useful buyer resolution, improving Owned Demand Share and credible payback. Investor reporting should show cash runway, committed versus discretionary spend, revenue quality, concentration and the cost of maintaining every “active” merchant or offer.

S8. Board financial pack [Back to TOC](#)

- Cash, runway and committed obligations by month.
- Revenue by engine, customer type and cohort; no blended vanity ARR.
- Gross and contribution margin with field/agency/creator costs separated.
- Active merchants, paying merchants and maintained fresh merchants as different measures.
- Partner pipeline, pilot cost, renewal decision and integration maintenance burden.
- Acquisition cost and payback by self-serve, agency, marketplace, market captain and direct sales channel.
- Resolved Demand Rate, direct repeat and external-platform concentration alongside revenue.
- Top risks, kill-rule status and spending that can be paused within 30 days.

Appendix T - Source and evidence register [Back to TOC](#)

[S01] Shoppage v3.0 attached operating draft, generated 14 July 2026. Internal source; 37 actual pages and approximately 14,900 extracted words.

[S02] ZIMSTAT, Report of the 2024-2025 Economic Census / preliminary results.
https://www.zimstat.co.zw/wp-content/uploads/production/EconomicCensus/EC2025_PRESENTATION.pdf

[S03] Reserve Bank of Zimbabwe, National Payment Systems Fourth Quarter 2025 report.
https://www.rbz.co.zw/documents/nps/quarterly/2025/NPSD_FOURTH_QUARTER_REPORT_ACTIVITY_DECEMBER_2025.pdf

[S04] Reserve Bank of Zimbabwe, Monetary Policy Statement, 7 August 2025.
https://www.rbz.co.zw/documents/mps/2025/MPS_7_August_2025.pdf

[S05] POTRAZ, S.I. 155 of 2024 Cyber and Data Protection licensing regulations.
https://www.potraz.gov.zw/wp-content/uploads/2025/02/si-155-of-2024-Cyber-and-Data-Protection-Normal_240913_1250178.pdf

[S06] Cyber and Data Protection Act [Chapter 12:07]. <https://www.potraz.gov.zw/wp-content/uploads/2022/02/Data-Protection-Act-5-of-2021.pdf>

[S07] Ownai Marketplace public site. <https://www.ownai.co.zw/>

[S08] Classifieds.co.zw public site and Terms of Use. <https://www.classifieds.co.zw/> and <https://www.classifieds.co.zw/terms-of-use>

[S09] Shop@Zim public marketplace site. <https://shopatzim.co.zw/>

[S10] Gateway Stream online shop. <https://shop.gatewaystream.com/>

[S11] PriceRunner public comparison site. <https://www.pricerunner.com/>

[S12] Google Search Central, Merchant listing structured data. <https://developers.google.com/search/docs/appearance/structured-data/merchant-listing>

[S13] Google Merchant Center, Comparison Shopping Services. <https://support.google.com/merchants/answer/12652686>

[S14] Google Merchant API product management. <https://developers.google.com/merchant/api/guides/products/overview>

[S15] Schema.org AggregateOffer and Offer types. <https://schema.org/AggregateOffer> and <https://schema.org/Offer>

[S16] Meta for Developers, WhatsApp catalogs and Cloud API. <https://developers.facebook.com/documentation/business-messaging/whatsapp/catalogs/catalogs-overview/>

[S17] WordPress Plugin Developer Handbook. <https://developer.wordpress.org/plugins/>

[S18] PriceCheck public comparison site. <https://www.pricecheck.co.za/>

Research caution. Source access, product features, fees, law and market conditions can change. Re-verify critical sources immediately before launch, contracting or investor circulation. Internal dataset counts must be independently audited and should not be represented as active merchants, unique audience or rights-cleared inventory until that work is complete.