

SHOPPAGE

Zimbabwe Commerce Intelligence, Discovery, Media and Market Infrastructure

National Launch Master Business Model v1.1 - Liability-Controlled Preloaded Intelligence and Execution-Gate Edition

Date: 10 July 2026

Status: Binding strategy, preloaded-intelligence doctrine, mandatory pre-gates and controlled-launch engineering blueprint

Supersedes: Preloaded Intelligence Launch Master v1.0 as the operating master

Core decision: Shoppage will not launch as an empty marketplace waiting for merchants to create its value. It will launch as a Zimbabwe-owned commerce operating system preloaded with a rights-governed national reference graph: more than one million master-product identities, at least ten thousand claim-ready organisation and merchant profiles, a national place scaffold, known markets and market-in-market structures, historical commerce evidence, and cold-start recommendations. Merchant-specific prices, stock, fulfilment and current operating claims remain time-bound local offer truth. Shoppage must remain operational if Facebook, Google, TikTok or any other external distribution platform is unavailable. **The public launch clock does not start until four mandatory pre-gates pass: source rights, field activation economics, native-demand superiority and independent security/privacy/rights review.**

Document control and review outcome

This edition retains the v1.0 preloaded-launch doctrine and converts the adversarial review into binding execution controls. It incorporates the validated Commerce Intelligence Foundation, the 10 July 2026 implementation progress report and the binding decision that Shoppage will enter the market with substantial preloaded product, organisation, place, market and historical intelligence rather than an empty catalogue. The v0.7 document remains an evidence compendium, v0.9 remains the consolidated predecessor and v1.0 remains the preloaded-intelligence decision record, but **v1.1 is the current operating master.**

v1.1 binding corrections

1. The 12-week build begins only after four pre-mobilisation gates have documentary evidence.
2. Unclaimed profiles receive exact public copy, an immediate suppression route and a 24-hour action SLA.
3. Every source enters a rights register before public indexing, AI use or external claims.
4. The one-million-plus master catalogue is a reference library; default buyer commerce search prioritises fresh local offers and the Request Network.
5. Home, Power and Building receives a separate launch taxonomy, master, attribute and compatibility programme.
6. Launch recommendations are deterministic and hybrid; learning-to-rank requires a behavioural evidence gate.
7. Pemosites is deferred beyond month six and cannot consume launch engineering capacity.
8. Freshness is enforced by auto-hide and ranking/advertising eligibility, not reminder messages alone.
9. Facebook and other social channels receive strong manual export and authorised-import workflows without becoming critical infrastructure.
10. Launch pricing is simplified to a free foundation and one paid growth pilot while willingness-to-pay and activation economics are measured.
11. Verification badges are predicate-specific and budgeted; 10,000 claim-ready profiles are not represented as field-verified merchants.
12. Reference records and unclaimed profiles are excluded from active-inventory, active-merchant and revenue KPIs in investor and board reporting.

The review made five structural corrections:

1. Replaced repeated executive, final-decision, roadmap, risk, AI, investor and architecture sections with one binding version of each.
2. Separated **validated foundation evidence, current production truth, launch hypotheses and long-term ambitions.**
3. Reframed "fully functional in three months" as a **controlled, end-to-end commercial launch within a narrow geography and category wedge**, not national completeness.

4. Made payments, fulfilment, merchant websites, Pemosites and the platform-framework model explicit architectural decisions.
5. Converted the large catalogue of ideas into a prioritised portfolio with launch, expansion and post-six-month horizons.
6. Made Facebook, WhatsApp, TikTok, Google and other external platforms optional adapters rather than identity, content or demand foundations.
7. Kept the five named Shoppage domains internally while simplifying the public proposition to one brand and three jobs: Sell, Discover and Grow.
8. Replaced launch-stage creator commissions with fixed, verifiable content and merchant-digitisation bounties; outcome commissions begin only after attribution and fraud controls pass.
9. Reframed institutional ambition as National Impact Readiness, with explicit protections against political capture, privileged ranking and surveillance use.
10. Locked a preloaded-launch doctrine: master products, organisation profiles, national places and known markets are loaded before public launch.
11. Defined merchant claiming, correction, duplicate resolution and evidence-preserving profile ownership.
12. Added a Historical Commerce Archive and source-governance rules for directories, public social evidence, collaborators and purchased data.
13. Made a cold-start recommendation engine and a rights-cleared Zimbabwe commerce intelligence model part of launch architecture.
14. Converted "succeed against all odds" into a mission-critical execution doctrine with redundancy, manual fallbacks, weekly gates and explicit truth controls.

Binding interpretation rule

Where this document uses the words **must**, **will** or **non-negotiable**, the statement is a current product or governance decision. Where it uses **target**, **hypothesis**, **candidate**, **pilot** or **later**, the statement requires validation before it becomes a binding commitment.

Evidence and public-display ladder

- **Foundation validated:** a schema, pipeline, count or test has been reproduced. Use it for architecture, engineering and launch preparation.
- **Preloaded reference:** a rights-cleared master product, classification, place or historical reference exists. Use it for search, enrichment and recommendations with reference labels.
- **Sourced candidate:** one or more sources indicate an organisation, market, branch or relationship may exist. Use it for a claim-ready profile, internal review and cautious public discovery.
- **Corroborated profile:** an authoritative source or multiple independent sources support defined identity, location or contact predicates. It may be displayed publicly as an unclaimed profile with source and freshness labels.
- **Merchant-claimed:** an authorised representative controls current merchant-managed fields. It may power the Merchant Site, catalogue and operational updates subject to audit.
- **Verified current claim:** a specific predicate has current supporting evidence. It may receive public badges and ranking treatment within the stated scope.
- **Production truth:** a verified claim is fresh and operationally safe. It may affect search ranking, ads, recommendations, payments and fulfilment.

Shoppage will launch with substantial preloaded references and corroborated profiles. It will not collapse these states into one blanket "verified" badge. A directory entry, map point, Facebook page, model output or imported record verifies only the predicate supported by its source and date.

Contents

1. Executive verdict and scope
2. Zimbabwe ground reality
3. Vision, mission and brand architecture
4. Shoppage system architecture
5. Shoppage Central
6. Shoppage Showcase
7. Shoppage Ads
8. Shoppage Creators

9. Shoppage Markets
 10. Merchant sites and Pemosites
 11. Commerce Intelligence Graph and Preloaded Launch Inventory
 12. Defensibility and network effects
 13. Daily utility and responsible habit formation
 14. AI-native operating model
 15. Trust, safety and consumer protection
 16. Payments and transaction ownership
 17. Fulfilment and last-mile delivery
 18. Data governance, rights and compliance
 19. Revenue model and pricing architecture
 20. Competitive strategy and platform substitution
 21. Shoppage Framework and ecosystem
 22. Go-to-market and field operations
 23. International quality and technical architecture
 24. Months 0-3 controlled launch roadmap
 25. Months 4-6 expansion roadmap
 26. Metrics, economics and investor proof
 27. Governance, organisation and decision rights
 28. National impact readiness and institutional engagement
 29. Risk register and decision gates
 30. Final business model and immediate actions
 31. Appendix A: prioritised 95-aim portfolio
 32. Appendix B: strategic extension portfolio
 33. Appendix C: core data entities and events
 34. Appendix D: source and evidence notes
 35. Appendix E: holistic review summary
 36. Appendix F: mandatory pre-public-launch gates
 37. Appendix G: source-rights register specification
 38. Appendix H: unclaimed-profile policy and public copy
 39. Appendix I: coverage-cell and native-demand test standard
 40. Appendix J: independence kill-switch acceptance tests
 41. Appendix K: creator-bounty operating procedure
 42. Appendix L: v1.1 engineering backlog and kill list
-

1. Executive verdict and scope

1.1 What Shoppage is

Shoppage is a local commerce-intelligence, discovery, media and operating platform. It contains five internal product domains with a reusable framework beneath them. These domains organise teams, permissions, APIs and economics; they are not five separate consumer products that users must learn.

Internal domain	Primary function	External promise
Shoppage Central	Merchant operating system and authoritative Merchant Site	Sell and maintain commercial truth
Shoppage Showcase	Search, geolocal discovery, requests and shoppable short video	Discover what is available, nearby and credible
Shoppage Ads	Local commerce advertising and attribution	Grow through measurable local outcomes
Shoppage Creators	Content production, community distribution and later affiliate commerce	Grow through trusted local voices

Internal domain	Primary function	External promise
Shoppage Markets	Digital operating layer for malls, markets and clusters	Sell and discover within real commercial places

The **Shoppage Framework** supplies shared identity, places, products, offers, search, trust, advertising, payments, fulfilment, intelligence, AI, billing and audit capabilities to the five domains and later approved extensions. Public communication uses one Shoppage brand and three verbs: **Sell, Discover and Grow**.

1.2 What Shoppage is not

During the first six months Shoppage is not:

- a national fulfilment fleet;
- a deposit-taking wallet or unlicensed escrow provider;
- a general entertainment social network;
- an unrestricted public plugin store;
- an unqualified claim that every preloaded organisation is currently trading or every master product is currently sold in Zimbabwe;
- an owner of every merchant transaction;
- a custom website agency disguised as software;
- a political advertising platform;
- a broad marketplace launching every category simultaneously.

1.3 Controlled-launch definition

“Fully functional in months 0-3” means the following end-to-end flows work in a controlled Harare launch:

- a merchant can discover or claim a preloaded profile, resolve ownership, create an authoritative Merchant Site, publish products and confirm price and stock;
- a buyer can search the preloaded master catalogue and national organisation/place graph, distinguish reference products from current local offers, browse short videos, compare local options, submit a request and contact a merchant;
- a merchant can buy a simple promotion and see attributable results;
- a creator can complete a contracted content, catalogue or merchant-onboarding bounty, while tagged links collect measurement data for later commission models;
- a market can have a verified page, roster and campaign surface;
- payment initiation and fulfilment options can be orchestrated through partners;
- evidence, provenance, source rights, correction, merchant claiming, moderation and audit trails are operational;
- the platform can measure search-to-contact, merchant response, freshness, recommendation quality, profile claims and verified outcomes.

It does not mean nationwide depth, every product category, mature transaction protection, public APIs or automated intelligence across the entire economy.

1.4 Launch wedge

The recommended launch wedge remains **Home, Power and Building** in Harare:

- solar, batteries and inverters;
- electrical and plumbing supplies;
- hardware and building materials;
- selected appliances and furniture;
- installation, repair and improvement services.

This wedge has high buyer intent, meaningful local availability differences, quote-based purchasing, compatibility questions, strong video-demonstration value, sufficient merchant margins for advertising, and diaspora relevance.

Shoppage will retain a reference library of more than one million master-product identities. The validated seed is primarily global food data and therefore cannot be treated as evidence that the launch wedge is ready. Before public launch, Shoppage must create a **separate Home, Power and Building launch corpus** with:

- at least 20,000 category-relevant product, variant and service identities from rights-cleared sources or merchant-authorised enrichment;
- complete attribute schemas for the top 2,000 expected demand entities;
- compatibility relationships for batteries, inverters, solar components, electrical protection, plumbing fittings and selected appliance parts;
- Zimbabwean aliases, abbreviations, units, common misspellings and local-language synonyms;
- merchant-offer attachment rules and human review queues.

The default buyer experience is **Local Commerce Search**: fresh branch-linked offers, active merchants, services, market pages and the Request Network. The wider million-product corpus is exposed through a clearly named **Reference Library** or enrichment workflow. Reference-only items must not dominate normal buyer results or imply that they are sold in Zimbabwe.

1.5 Launch and month-three operating envelope

Target area	Binding launch/month-three target	Interpretation
Master reference library	1,000,000+ indexed master identities	Separate reference search partition; source/licence preserved; local availability never inferred
Launch-wedge corpus	20,000+ relevant product, variant and service identities	Home, Power and Building attributes, aliases and compatibility prioritised
Organisations/merchants	10,000+ claim-ready profiles	Source-labelled candidates/corroborated predicates; never counted as active merchants until claimed and operating
Activated merchants	50-100 claimed, responsive merchants	Release depends on activation economics and coverage cells, not profile volume
Local offers	3,000-5,000 fresh, branch-linked offers	Each has price/availability evidence and category freshness state
Places	National place scaffold loaded	Designated-source coverage; correction and reconciliation remain continuous
Markets/clusters	Publicly evidenced markets loaded; 5-10 operational pages	Market-in-market containment requires reviewed evidence
Creators	20-40 trained and approved	Fixed-bounty model; no outcome commissions at launch
Commerce clips	250-500 catalogue-linked clips	Moderated, rights-cleared, canonical-link compliant
Recommendation engine	Rules-based and hybrid cold-start recommender	Learning-to-rank is gated until sufficient resolved demand exists
Merchant claiming	Ownership, correction, suppression, merge and dispute live	24-hour action SLA for material unclaimed-profile complaints
Payment rails	At least two practical provider/rail options	Shoppage does not hold customer funds
Fulfilment modes	Collection, merchant delivery and partner quote	Asset-light orchestration
Freshness	At least 80% of visible offers inside category window	Stale offers auto-hide and lose ad eligibility
Usable coverage cells	Minimum three claimed merchants, 50 fresh offers and >60% response	Ads cannot be sold in cells below usable status

These are operating targets. Reference scale, profile scale and place coverage are front-loaded assets; commercial success is measured by claimed profiles, fresh offers, resolved needs, repeat direct use, merchant economics and revenue.

2. Zimbabwe ground reality

The model must work with current Zimbabwean commerce rather than imported assumptions. Public evidence describes e-commerce as predominantly urban and informal, with Facebook Marketplace and WhatsApp central to product discovery and negotiation, and cash, cash on delivery and mobile payments common in place of card-first website checkout.[1]

Digital 2026 estimates 6.54 million individual internet users, 2.60 million social-media user identities and 16.2 million active cellular connections in late 2025. These figures support a meaningful mobile opportunity but do not justify designing as though every connection is an active online shopper or every user has stable broadband.[2]

UNCTAD identifies diaspora demand, weakly structured logistics, informal service providers, last-mile gaps, a need for pickup centres, digital-skills constraints and an incomplete e-commerce legal environment.[3]

2.1 Design consequences

Ground reality	Required Shoppage response	Failure avoided
Commerce starts on social channels	Shareable product cards, clips and links that return to Shoppage truth	Demanding immediate channel abandonment
Trust is local and relationship-based	Progressive verification, locations, response history and proof-backed reviews	Anonymous-listing marketplace behaviour
Connectivity and data are uneven	Low-data mode, manual location, compressed media and resilient forms	Video-heavy experience that excludes users
Payments are mixed	Cash, collection, mobile money, transfer and provider checkout	Mandatory card-first funnel
Logistics are fragmented	Quote, collection and partner orchestration	Premature national fleet
Merchant digital capability varies	Assisted onboarding, voice/image ingestion and managed campaigns	Self-service-only adoption model
Currency conditions change	Multi-currency data model and dated exchange assumptions	Hard-coded currency logic
Public business data is fragmented and uneven	Preload several sources, preserve evidence, enable merchant claims, field verification and public corrections	Pretending one directory or scrape is national truth

2.2 Inclusion requirements

Shoppage should be smartphone-first, not smartphone-exclusive in its operating model. Core inclusion measures are:

- text-first and low-data browsing;
- manual city, suburb, landmark and market selection;
- agent-assisted onboarding and verification;
- shareable links that work without an installed app;
- English at launch with progressive Shona and Ndebele search, prompts and support;
- voice-note and image-assisted merchant ingestion;
- clear price dates and currency labels;
- no forced buyer registration before search or contact.

3. Vision, mission and brand architecture

3.1 Vision

Make Zimbabwean commerce visible, searchable, trusted, measurable and locally valuable.

3.2 Mission

Shoppage will organise merchant identity, places, products, services, offers, demand, content and commercial outcomes into a continuously improving local commerce graph. It will use that graph to help buyers resolve needs, help merchants grow, help creators earn, help markets digitise and help local institutions understand commerce without selling personal data.

3.3 Positioning

For buyers: find what is available, where it is, who can be trusted and how to obtain it.

For merchants: publish once, maintain commercial truth, reach local demand and measure every meaningful action.

3.4 Brand rules

Shoppage is the master brand. Central, Showcase, Ads, Creators and Markets are internal product domains and specialist workspaces, not five competing consumer brands or separate public companies. One identity, catalogue, trust system, billing relationship and analytics model must operate across all domains.

Public product language is deliberately simpler:

- **Sell:** create and maintain a Merchant Site, catalogue, market presence, leads, payment and fulfilment options.
- **Discover:** search, browse nearby availability, watch useful commerce media, compare options and post requests.
- **Grow:** promote products, commission content, understand demand and improve merchant performance.

A buyer should never be required to decide whether to open “Showcase” or “Markets”; the buyer searches Shoppage. A merchant should not need to understand an internal domain map before publishing or promoting an offer. Domain names remain useful in architecture, investor and partner materials, and in specialist workspaces where the distinction improves operations.

Recommended digital structure:

- [shoppage.co.zw](#) - public discovery, requests, market pages and Merchant Sites;
- [manage.shoppage.co.zw](#) - merchant Sell workspace;
- [grow.shoppage.co.zw](#) - advertising, creator and growth workspace, with role-specific views;
- [developers.shoppage.co.zw](#) - later controlled developer portal;
- market pages at [shoppage.co.zw/markets/{slug}](#).

Separate ads and creators subdomains may exist operationally, but they should not be the primary public acquisition message. Routes may share one deployment during the first three months. Domain, company-name and trademark due diligence are required before major brand investment.

4. Shoppage system architecture

4.1 Shared horizontal engines

The five internal domains depend on common engines. Duplicating these engines inside each domain would weaken the moat and increase cost.

Shared engine	Responsibility	Non-negotiable control
Identity and organisations	People, merchants, branches, roles and claims	Tenant isolation and auditable authority
Places and markets	Geography, service areas, markets and containment	Evidence-based relationships, not proximity guesses
Catalogue and offers	Masters, variants, services, prices	Product identity separated from

Shared engine	Responsibility	Non-negotiable control
	and stock	seller offer
Search and discovery	Keyword, semantic, geolocal and request matching	Relevance and freshness before paid influence
Trust and safety	Verification, reports, complaints and reputation	Human appeal for high-impact decisions
Content and media	Images, short video, captions and rights	Every commerce clip links to structured objects
Advertising and attribution	Campaigns, placements, spend and outcomes	Transparent labels and relevance floors
Payments and fulfilment	Provider initiation, quotes, statuses and evidence	No unlicensed custody or hidden off-platform handoff
Intelligence and AI	Enrichment, recommendations, forecasting and copilots	Evidence, evaluation and human review
Billing and entitlements	Plans, limits, credits and partner fees	Clear invoices and auditable balances

4.2 Primary participants

- buyers and household purchasers;
- informal and formal merchants;
- branches, stalls and service providers;
- creators, affiliates and community operators;
- markets, malls and commercial clusters;
- brands, distributors and manufacturers;
- payment, fulfilment, POS and service partners;
- Shoppage field, moderation, support and data teams;
- approved powered businesses and subsidiaries.

4.3 Core journeys

Merchant journey

Preload and resolve organisation/branch -> expose claim-ready profile -> authorised representative claims or corrects it -> verify relevant predicates -> create authoritative Merchant Site -> attach masters/variants and current offers -> confirm freshness -> publish content -> receive leads -> promote -> record outcomes -> improve from intelligence.

Buyer journey

Search, browse or ask -> choose location and constraints -> inspect product, merchant and freshness -> compare total outcome -> contact, request, reserve or initiate payment -> choose collection/delivery -> receive proof and support -> save need or review outcome.

Creator journey

Apply and verify -> select approved merchant/product -> create or upload content -> disclose commercial relationship -> tag canonical objects -> publish -> generate attributable action -> pass fraud/dispute window -> receive payout.

Market journey

Resolve operator and boundary -> create market structure -> onboard merchants/stalls -> publish market page -> maintain roster and hours -> run market campaigns -> measure traffic, demand and merchant outcomes.

4.4 Independence doctrine

Shoppage must operate as a complete core product without Facebook, WhatsApp, TikTok, Google or any other external platform. External platforms can accelerate acquisition and distribution, but they cannot be required for identity, canonical content, search, requests, leads, merchant maintenance or outcome history.

Identity is owned by Shoppage

- Every user, merchant, branch, creator and partner has a canonical Shoppage identifier.
- Phone or email OTP and Shoppage credentials are the default login paths; social login is optional convenience only.
- Facebook Page IDs, WhatsApp business numbers, Google Business IDs and other platform identifiers are stored as revocable **external_identity** records with source, scope, confidence and last-verification state.
- No external-platform identifier may be a database primary key or the only recovery route.
- Identity evidence follows data minimisation. National ID or equivalent evidence is requested only when a defined risk, regulatory or transaction rule requires it, preferably through an approved verifier without retaining unnecessary raw documents.

Content is owned by Shoppage

The Merchant Site, product, variant, offer, market, creator asset and campaign record are canonical. A Facebook post, WhatsApp share, TikTok clip, Google feed or other external representation is a render or referral of the Shoppage record. External content must link back to a stable Shoppage object whenever platform rules permit. Imported observations never silently overwrite merchant-authorised current truth; they create evidence or review events. Merchant claims likewise do not erase older evidence, regulatory facts, complaint records or source provenance.

Demand is owned by Shoppage

Direct demand must be created through capabilities external social platforms cannot reliably provide together:

1. **Freshness truth:** visible last-confirmed price and stock, category-specific expiry, one-tap merchant reconfirmation and automatic ranking downgrade for stale offers.
2. **Operational proximity:** open-now state, realistic distance, branch availability, accepted payment methods, collection and delivery feasibility, response performance and local landmarks.
3. **Structured Request Network:** buyers state an exact need, location, timing and constraints; Shoppage routes it to eligible merchants, records responses and learns unresolved demand.

External adapters are optional translators

Where a platform permits compliant APIs, Shoppage may read or publish through a scoped adapter. Where it does not, Shoppage provides user-initiated sharing and export assets. The launch Facebook/WhatsApp distribution kit must generate:

- 1080 x 1080 product and offer cards;
- price, currency, location and "Confirmed Today" overlays when current;
- QR code and canonical Shoppage Universal Link;
- English, Shona and later Ndebele copy-ready captions;
- UTM/referral identifiers and campaign attribution;
- safe fallback copy when price or stock is not confirmed.

Merchant-authorised Page catalogue APIs, merchant-provided CSV files and user-initiated exports are approved source patterns. Unauthorised scraping, credential sharing or automation that violates platform terms is not a core growth strategy.

Core domains: Identity, Organisation, Product, Offer, Request, Lead, Content, Outcome

Optional adapters: Facebook, WhatsApp, TikTok, Google, POS, ERP, payments, fulfilment

Rule: core events may be exported to adapters; critical core flows never require adapter availability

Kill switch: disabling any adapter must not block login, publish, search, request, lead or maintenance

4.5 Independence acceptance criteria

Before controlled launch:

- the application boots and core buyer/merchant flows pass with all social adapters disabled;
- no social login or external ID is required to recover an account;
- every outbound social object resolves to a canonical Shoppage URL or documented export artifact;
- adapter timeout, revocation and policy-change tests do not corrupt core data;
- the installable PWA can display previously cached Merchant Sites, products, saved needs and operating information while temporarily offline; merchant draft changes can queue locally and sync safely when connectivity returns;
- full server search, payment status and live availability are never falsely presented as current while offline;
- merchant data and media can be exported in a documented format;

- an external-platform concentration dashboard tracks referral dependence and direct/owned return traffic.

5. Shoppage Central

Shoppage Central is the merchant operating system and must create standalone value before Shoppage has national buyer traffic.

5.1 Merchant Site as authoritative commercial reference

Every activated merchant receives a Shoppage Merchant Site that can serve as the reference website for the business. It includes:

- verified business and branch profile;
- products, variants, services, prices and freshness;
- images, short videos, offers and collections;
- maps, hours, delivery zones and collection options;
- payment methods and transaction options;
- WhatsApp, call, request-quote and reservation actions;
- policies, warranties, FAQs and trust signals;
- structured metadata for public search engines;
- QR code and shareable links;
- analytics and campaign attribution.

A Merchant Site is not a decorative profile. It is the first-party publishing surface for commercial truth.

5.2 Merchant workspace

Workspace	Core functions	Launch control
Today	Leads, stale stock, price checks and urgent actions	Daily action queue
Catalogue	Products, variants, services, media and imports	Draft, review and audit history
Offers	Price, currency, availability, branch and validity	Freshness and evidence state
Content	Clips, cards, captions, collections and QR	Rights and moderation checks
Leads	Calls, WhatsApp, RFQs, reservations and notes	Source attribution and status
Ads	Simple objectives, budget, geography and reports	Spend caps and relevance checks
Creators	Invitations, bounty contracts, assets and later commissions	Contract, disclosure and release-gate state
Intelligence	Demand, gaps, response and improvement actions	Confidence and privacy controls
Settings	Staff, roles, branches, policies and integrations	Least privilege and audit log

5.3 Merchant ingestion

Merchants should be able to publish through:

- manual mobile entry;
- spreadsheet or CSV;
- authorised POS, ERP or distributor feed;
- image or PDF extraction with confirmation;

- voice note or WhatsApp-assisted draft;
- agent-assisted field capture;
- later API or plugin integration.

AI may draft titles, attributes, categories and variants, but the merchant or authorised operator confirms commercial claims before publication.

5.4 Freshness system

Each visible offer has:

- observed or confirmed time;
- price-valid-until where appropriate;
- stock state and confidence;
- branch or service area;
- evidence source;
- stale-after rule by category;
- merchant reconfirmation path;
- customer correction/report path.

Freshness states should be **confirmed**, **recent**, **aging**, **stale**, **unknown** or **disputed**. Search and ads must downgrade or exclude stale and disputed offers.

5.5 Merchant improvement loop

Central should tell a merchant what to do next, not merely show charts. Examples:

- confirm 14 offers that will expire tomorrow;
- reply to three high-intent requests;
- add a missing inverter capacity attribute;
- create a 20-second demonstration for a frequently viewed product;
- extend delivery to an area with unresolved demand;
- stop an ad whose stock is no longer available;
- compare response time with the category benchmark.

6. Shoppage Showcase

Showcase is the buyer front door: commercial search, nearby discovery, requests, collections and low-data short-form commerce video.

6.1 Buyer surfaces

Surface	Buyer job	Launch requirement
Search	Find a product, variant, service or supplier	Typo tolerance, synonyms, geo and freshness
Nearby	See credible availability around a place	Manual location and realistic radius
Showcase feed	Discover useful product content	Catalogue-linked clips only
Product page	Understand exact item and offer	Variant, price date, stock and vendor
Merchant Site	Evaluate the seller	Branches, policies, trust and contact
Market page	Browse a mall, market or cluster	Verified roster, map and content
Ask/Request	Describe an unmet need	Structured routing and response tracking

Surface	Buyer job	Launch requirement
Saved needs	Return when availability or price changes	Consent-based alerts
Compare	Compare total outcome	Price, trust, location, warranty and delivery

6.2 Geolocal commerce

Geo must affect search, requests and advertising, but location precision must be proportionate and consent-based. The graph should support:

- country, province, city, suburb, growth point and landmark;
- market, mall, building, section, stall and branch;
- pickup points and verified collection points;
- merchant delivery zones and service radii;
- routes and corridors for peri-urban and rural fulfilment;
- “open now”, “available today” and realistic travel/delivery filters;
- coverage cells by category and place.

GPS is optional. Manual location is mandatory.

6.3 Short-form commerce video

Showcase video is commerce media, not general entertainment. Every clip must connect to a product, service, merchant, market, collection, offer or request.

Launch formats:

- product demonstration;
- new-stock or price-drop alert;
- store or market walk-through;
- installation or before-and-after proof;
- comparison or buying guide;
- creator review with disclosure;
- frequently asked question;
- availability update.

Operational rules:

- normally 10-60 seconds;
- muted autoplay only where data settings allow;
- captions and thumbnails required;
- low-data mode and conservative preloading;
- content rights recorded;
- current price/stock overlay or stale-content warning;
- clear creator/advertiser disclosure;
- reporting on views, saves, catalogue opens, contacts and outcomes.

6.4 Native demand wedge

Showcase must earn direct use before it attempts broad media scale. The first buyer promise is: **find the credible nearby option that can serve me now, or route my request to someone who can**. Freshness truth, operational proximity and the Request Network are therefore launch-critical; the video feed is supportive discovery, not the primary retention mechanism.

Direct-use tests should measure whether a buyer can resolve a need faster and with less uncertainty than posting in a group, scrolling old listings or calling several shops.

7. Shoppage Ads

Shoppage Ads is the mechanism for moving a measurable share of Zimbabwean commercial advertising spend from broad global attention platforms into local, outcome-oriented commerce media.

7.1 Advertising doctrine

- Sell commercial intent and attributable action, not vanity engagement.
- Protect organic relevance; payment cannot make an irrelevant result useful.
- Promote only eligible merchants, products, services and content.
- Display paid labels clearly.
- Use simple objectives for SMEs and deeper controls for larger advertisers later.
- Report the buyer action that followed the ad.
- Exclude political advertising during the initial operating period.
- Maintain an independent ad-integrity function with authority to reject revenue.

7.2 Launch products

Ad product	Best use	Initial charging model
Sponsored product	High-intent category or keyword demand	Budget/click or fixed package
Sponsored Merchant Site	Store or branch visibility	Fixed duration or impression package
Sponsored clip	Product demonstration and launch	Budget/view with action reporting
Nearby deal	Local availability and store traffic	Area-duration package
Market sponsorship	Market, mall or commercial-cluster promotion	Contracted package
Request matching	Qualified buyer request access	Qualified lead or subscription
Managed campaign	Merchants lacking internal capability	Service fee plus media budget
Creator campaign	Creator-produced catalogue-linked content	Fixed bounty at launch; verified-outcome commission only after gates

7.3 Simple campaign workflow

1. Select objective: visibility, catalogue visit, WhatsApp lead, call, quote, reservation or verified sale later.
2. Select the product, Merchant Site, market or clip.
3. Choose location, category, duration and budget.
4. Receive automatic quality, stock, policy and reach checks.
5. Publish automatically for low-risk trusted advertisers or enter review.
6. Apply frequency limits and organic-result safeguards.
7. Report spend, reach, actions, merchant response and verified outcomes.
8. Provide recommendations without silently changing budget or targeting.

7.4 Substitution metrics

- **Commercial Discovery Share:** share of covered-category buyer needs resolved through Shoppage.
- **Merchant Advertising Wallet Share:** share of a merchant's measurable acquisition budget allocated to Shoppage.
- **Local Advertising Value Retention:** advertiser revenue remaining in Zimbabwe after direct imported infrastructure, processing and software costs.

Shoppage should not claim that every advertising dollar remains in Zimbabwe. The credible claim is that more campaign capability, data, jobs, creator payouts, partner revenue, tax and merchant relationships are retained locally.

7.5 Cold-start advertising rule

Ads cannot succeed without useful organic demand. During months 0-3, revenue should come from managed campaigns, market packages, onboarding bundles and clearly limited sponsored placements. Large-scale self-service advertising should expand only after search volume, repeat users and outcome measurement are credible.

8. Shoppage Creators

Creators are a core distribution, education, merchant-digitisation and content-production network. Facebook group administrators, WhatsApp community operators, product reviewers, technicians, videographers, market walkers, campus promoters and specialist communities already hold trusted local attention. Shoppage should convert that capability into structured commerce without building a launch-stage affiliate-fraud engine.

8.1 Creator roles

Role	Value created	Initial earning method
Commerce creator	Product demonstrations and reviews	Contracted content bounty
Market creator	Market walk-throughs and merchant discovery	Market content package
Specialist creator	Technical buying education	Contracted guide/video fee
Catalogue creator	Merchant media and product-data capture	Accepted-work service fee
Community operator	Relevant distribution to a trusted audience	Approved post/distribution bounty
Affiliate curator	Collections and tagged recommendations	Measurement pilot, then qualified-outcome commission

8.2 Tier 1 - months 0-3: fixed, verifiable bounties

Launch compensation is based on accepted work, not clicks, views or unverified sales. Example packages include one catalogue-linked video, a defined number of compliant community posts, a merchant onboarding, a market walk-through or a product-data capture assignment.

A bounty is payable only when:

- the merchant, products, deliverables, rights and fee are agreed in advance;
- the creator uses the canonical Shoppage Universal Link or QR asset;
- sponsored disclosure is present;
- the content remains available for the contracted verification period, normally at least 48 hours unless the platform removes it;
- Shoppage verifies the deliverable and there is no rights, fraud or safety breach.

Tagged links and events are collected from day one for learning, but they do not automatically determine creator pay during the initial period. This keeps creators funded while attribution quality is being proven.

8.3 Tier 2 - months 4-6: verified lead commissions

Outcome-based commissions begin only after unique-link attribution, duplicate detection, merchant response workflows, dispute handling and payout reconciliation pass release gates. Eligible commission events may include a structured qualified lead, accepted quote, completed reservation, verified collection or confirmed transaction where supported.

A WhatsApp click alone is not a sale and should not automatically trigger commission. The platform should combine a tagged Shoppage session, a consented or structured lead event, duplicate-phone and device checks where lawful, merchant response state and sampling/audit. Off-platform merchant assertions remain lower-confidence evidence unless independently supported. High-risk categories are excluded until specialist policies and review capacity exist.

8.4 Canonical-link and distribution rule

Every paid creator post must use a stable Shoppage Universal Link, not make an external social shop or page the only destination. Links must survive copying between Facebook, WhatsApp, SMS and browsers, retain privacy-safe attribution, open the relevant Merchant Site/product/request and degrade gracefully on low-data devices.

8.5 Safeguards

- sponsored and affiliate disclosure;
 - merchant approval of promoted objects and terms;
 - music, image, customer-consent and content rights;
 - anti-self-referral, duplicate-lead and collusion controls;
 - review windows before outcome-based payout;
 - category restrictions and expert review for high-risk claims;
 - creator appeals and merchant dispute handling;
 - tax and contractor treatment confirmed before scale;
 - follower count, raw views and superficial engagement are never sufficient payout events.
-

9. Shoppage Markets

Markets turns physical commerce clusters into searchable, measurable and promotable digital operating units.

9.1 Supported market types

- malls and shopping centres;
- formal and informal markets;
- CBD retail clusters;
- industrial and hardware districts;
- trade fairs and temporary events;
- growth points and rural service centres;
- distributor and franchise networks;
- diaspora-oriented supplier collections.

9.2 Market page

A useful market page includes:

- verified operator or governance status;
- location, boundary and navigation;
- current merchant/stall roster;
- categories, products and services;
- opening hours and event dates;
- market clips and offers;
- collection and delivery information;
- complaints and correction path;
- campaign and sponsorship inventory;
- market-level analytics.

9.3 Market evidence

Market-inside-market, building, section, store and stall relationships cannot be inferred from nearby coordinates alone. Operator records, municipal evidence, field verification, address evidence or reviewed polygon relationships are required.

9.4 Market captains

A Market Captain may onboard merchants, maintain rosters, capture media, coordinate campaigns and escalate disputes. The role requires:

- defined authority and territory;
 - training and identification;
 - audit trails and quality score;
 - separation between verification and personal commercial interest;
 - anti-bribery and complaint controls;
 - compensation tied to accepted work and durable activation, not raw registrations.
-

10. Merchant sites and Pemosites

10.1 Binding decision

Every merchant receives a Shoppage Merchant Site. Merchants with effective existing websites may keep them and connect them through feeds, APIs, embeds and attribution links. Shoppage should not force unnecessary migration.

10.2 Source-of-truth model

For every field, the integration must define the authoritative system. Examples:

- legal identity may come from verified merchant/regulator evidence;
- product master may be Shoppage-controlled;
- price and stock may be merchant POS-controlled;
- editorial corporate content may be website-controlled;
- trust and response history remain Shoppage-controlled.

No bidirectional sync should launch without conflict rules, timestamps, ownership and recovery procedures.

10.3 Pemosites

Pemosites remains a candidate premium product line using the Shoppage core, not a launch workstream and not a separate operating company.

Binding deferral: no Pemosites custom-domain programme, custom templates, migration service or bespoke corporate-site work may enter the months 0-6 roadmap. The launch team builds only the reusable Shoppage Merchant Site engine.

Pemosites may enter planning after month six only when all of the following pass:

- at least 50 activated wedge merchants;
- at least 80% offer freshness for four consecutive weeks;
- core search, Request Network and merchant maintenance reliability meet target;
- a separate product owner and budget are assigned;
- requests can be served through reusable templates rather than custom engineering;
- work cannot reduce delivery capacity for Shoppage core.

When admitted, Pemosites must reuse Shoppage identity, products, branches, offers, analytics and permissions. It must never create a parallel merchant or product database.

10.4 Spinout gate

Consider a separate legal entity only after evidence such as:

- at least 100 paying sites or equivalent contracted recurring revenue;
- repeatable acquisition and onboarding;
- sustainable gross margin after support labour;
- standardised product rather than bespoke agency dependence;
- separate management capacity;
- clear IP, data-processing and cross-selling agreements.

Months 0-3 build the shared Merchant Site engine. Months 4-6 may pilot Pemosites with 10-20 businesses.

11. Commerce Intelligence Graph and Preloaded Launch Inventory

The strongest moat is not the interface. It is a governed, continuously improving evidence graph that connects organisations, branches, places, markets, products, variants, offers, content, demand and outcomes.

11.1 Validated foundation and binding launch expansion

The attached implementation package validates:

- 6,101 organisation-source records and 5,281 distinct normalised names;
- 5,899 ISIC-mapped records;
- exactly 1,000,000 unique global food-product identities;
- 950,334 checksum-valid GTINs and 466,276 family keys;
- 25,210 place/market nodes and 25,209 parent edges;
- 118 explicitly tagged markets or malls;
- a versioned merchant-offer feed contract, idempotent observation store, exact-GTIN matching and review queue.

The binding v1.1 launch expansion is:

- retain and productionise the one-million-plus reference catalogue in a separate rights-aware search partition;
- build the 20,000+ launch-wedge corpus and compatibility graph;
- ingest and resolve at least 10,000 claim-ready organisation profiles from approved sources;
- load the national place scaffold and publicly evidenced market inventory;
- create claim, correction, suppression, duplicate and dispute workflows;
- acquire merchant-authorised local offers through field onboarding, distributors, manufacturers, POS/ecommerce feeds, CSV and Merchant Site maintenance;
- deploy deterministic/hybrid cold-start recommendation and request routing;
- maintain a source-rights ledger and deletion-propagation system before public indexing.

The package proves only the reproduced foundation counts. It does not prove 10,000 current trading merchants, one million Zimbabwe offers or complete official market containment. External claims and board reporting must preserve this distinction.

11.2 Stable entities and changing claims

Stable entity	Time-varying claim	Example evidence
Organisation	Currently trading	Merchant, field or transaction evidence
Branch	Operating at a location	Merchant, operator or field verification
Master product	Sold by a merchant	Rights-cleared merchant/distributor feed
Variant	In stock at a branch	Current offer observation
Place/market	Contains a store or stall	Operator, municipal or reviewed field evidence
Creator	Authorised to promote an offer	Campaign contract and merchant approval
Merchant	Responds reliably	Measured lead history
Offer	Price is valid	Timestamped merchant/provider evidence

11.3 Source and provenance requirements

Every production claim should preserve:

- source system, source type and authority;

- predicate scope;
- source identifier and permitted reference;
- observed/retrieved time;
- evidence hash or immutable reference;
- reuse rights and restrictions;
- confidence and verification state;
- reviewer and decision reason;
- validity and supersession history;
- unresolved conflicts.

11.4 Coverage cells

A coverage cell represents a place-category combination. It tracks:

- verified merchants and branches;
- active markets;
- fresh offers and variant depth;
- price and stock recency;
- response rate and verified outcomes;
- buyer demand and unresolved requests;
- trust and complaint signals;
- coverage state: empty, thin, usable, strong or authoritative.

Coverage cells guide field operations, creator tasks, category launches, advertising availability and honest national reporting.

11.5 Intelligence products

Buyer intelligence

- Shoppage Ask and saved needs;
- Price Pulse with date, currency and comparable context;
- Stock Truth with confidence and confirmation time;
- Compatibility Graph for parts, accessories and systems;
- local substitution and alternative suppliers;
- total-outcome comparison including delivery, warranty and trust.

Merchant intelligence

- Business Improvement Copilot;
- demand gaps by area;
- response and conversion diagnostics;
- catalogue completeness and freshness;
- pricing and assortment benchmarks;
- campaign and creator recommendations;
- branch and delivery-zone opportunities.

Market, brand and institutional intelligence

- category demand and stock-out patterns;
- market and cluster performance;
- reseller and distribution visibility;
- price and availability trends;
- merchant digitisation and local-value scorecards;
- privacy-safe aggregated planning data.

Shoppage monetises intelligence, not raw personal data or merchant-confidential records.

11.6 Preloaded intelligence inventory

Shoppage launches with distinct inventories that must never be conflated:

Inventory	Launch content	Public treatment
Master reference library	1,000,000+ product identities and global metadata	Explicit Reference Library; not default evidence of Zimbabwe availability
Launch-wedge commerce corpus	20,000+ relevant products, variants, services and compatibility edges	Used in local search only where a fresh offer, active service or useful Request path exists
Organisation graph	10,000+ organisations, branches and facilities	Unclaimed reference or claimed Merchant Site; exact source and freshness labels
Place and market graph	National places, markets and supported containment edges	Searchable geography; uncertain relationships remain review-labelled
Historical commerce archive	Time-stamped approved observations	Intelligence and entity resolution; no verbatim reuse without rights
Current local offers	Branch-linked price, currency, stock, fulfilment and evidence	Primary buyer commerce truth; expiry and correction enforced

Every record receives a canonical ID, source ledger, rights class, observation date, confidence, review state, current-state status and correction path. Search and AI may use incomplete records internally, but public interfaces must communicate uncertainty and cannot imply a commercial relationship that has not been established.

11.7 Merchant claiming, correction and profile ownership

Every preloaded merchant or company profile should be claimable. The claim workflow must support:

1. phone, email-domain, document, branch, association, field-agent or authorised-partner evidence;
2. organisation/branch resolution before ownership is granted;
3. owner, manager and delegated staff roles;
4. duplicate merge and mistaken-identity review;
5. correction and suppression requests from merchants and the public;
6. appeals when a claim is rejected or challenged;
7. audit history for every field changed after claim;
8. export of merchant-managed data and revocation of staff access.

Every public unclaimed profile must display this or equivalent approved wording prominently:

Unclaimed reference profile. This page was assembled from the dated public or licensed sources shown below and is not currently managed by the business. Information may be incomplete or outdated. **Claim this business or request a correction/removal.**

Mandatory controls:

- show source name/class and last-observed date for each public predicate where practical;
- no “verified merchant”, “active seller”, stock, price, ad or payment badge without a current supporting state;
- no paid promotion from an unclaimed profile;
- acknowledge every correction/removal request automatically;
- within 24 hours, correct or temporarily suppress materially harmful or identity-sensitive disputed fields while review continues;
- resolve ordinary corrections within three business days where evidence is available;
- provide an escalation path and immutable audit log;
- measure complaint volume, action time, correction time and reversal rate.

A merchant claim gives authority over current trading fields, catalogue and offers. It does not permit deletion of lawful regulator records, historical evidence, complaints, source provenance or other parties' rights. Claim authority and historical evidence must remain separate.

11.8 Historical Commerce Archive

Shoppage may build historical intelligence only from source classes that pass the Source Rights Register:

- merchant-supplied Facebook, Instagram, WhatsApp, POS, ecommerce and website exports;
- platform-authorised APIs and approved business integrations;
- collaborating websites and platforms under written data-sharing agreements;
- purchased directories with explicit provenance, permitted-use, update and deletion warranties;
- public regulator, association, mall and market records where reuse is lawful;
- field observations, price tags, catalogues, receipts and campaign outcomes collected under contributor terms;
- merchant, distributor and manufacturer feeds.

Public Facebook groups and posts are not an automatic source class. Collection requires a documented lawful basis, platform/contract permission where applicable, purpose limitation, personal-data review and a deletion route. Automated collection without permission is prohibited as a production dependency.

The archive should store structured observations and source references rather than copying whole posts. Copyrighted text, images and video require separate rights. Personal conversations, private groups, minors' data, sensitive information and unrelated social interactions are excluded. Source withdrawals and corrections must propagate to public views, derived features and future model corpora according to the deletion plan recorded in the register.

11.9 AI recommendation and local commerce model

A recommendation capability is required at launch, but launch production is **rules-based and hybrid**, not an unsupported claim of mature machine learning.

Launch recommendation signals

- query and explicit buyer need;
- category, brand, attributes, variant and reviewed compatibility;
- distance, delivery area, open-now state and collection feasibility;
- fresh local offer, price/currency and stock confidence;
- merchant verification scope, response and complaint state;
- product completeness and content usefulness;
- saved needs, exclusions and user-controlled preferences;
- seasonality, market context and diaspora/local mode;
- transparent sponsored eligibility after organic relevance floors.

Launch decision order

1. filter out stale, unsafe, ineligible or geographically impossible results;
2. apply deterministic category, attribute, compatibility and location rules;
3. rank fresh local offers and responsive merchants;
4. route thin-coverage needs into the Request Network;
5. use semantic retrieval and model reranking only inside defined confidence limits;
6. expose explanations such as “nearby”, “confirmed today” or “matches 48V requirement”.

Evaluation corpus and release gate

Before launch, Shoppage must maintain at least 500 reviewed Zimbabwe commerce queries/needs across English, local terminology, misspellings, locations and launch categories. Release thresholds include:

- zero unsupported stock, price or verified-status claims;
- locality and compatibility error rates below approved thresholds;
- measured precision@5 and request-routing success above a simple keyword baseline;
- human review of high-risk compatibility and substitution recommendations;
- deterministic fallback when model services fail.

Learning-to-rank or behavioural personalisation may enter production only after at least **5,000 resolved needs** or an approved equivalent evidence threshold, plus offline evaluation, staged rollout and rollback controls.

Shoppage Local Commerce Intelligence Model

Shoppage should build a local model stack, not train a frontier foundation model from scratch during launch. It comprises taxonomy and aliases, multilingual embeddings, semantic retrieval, entity resolution, compatibility rules/models,

recommendation/ranking, fraud/moderation and retrieval-augmented copilots. All training, fine-tuning and evaluation corpora require a manifest, rights class, provenance, removal path and leakage tests. Model output is never evidence by itself.

11.10 No-empty-platform launch rule

The platform launches preloaded, but reference scale cannot overwhelm buyer truth.

Before public launch:

- default buyer search ranks fresh local offers, active services, claimed Merchant Sites, operational markets and Request options;
- reference-only results appear in a clearly separated Reference Library or an explicitly labelled fallback section;
- a reference-only product must never appear to be available locally;
- category pages are opened publicly only when their coverage cells reach usable status or the Request Network provides a credible alternative;
- unclaimed merchants can be claimed, corrected or suppressed without requiring an engineer;
- recommendations degrade to deterministic rules and Request routing when data is sparse;
- source, rights and deletion controls are operational;
- the system never manufactures stock, prices, operating status, ownership or market containment from inference.

The rule is: **preloaded breadth supports intelligence and enrichment; fresh local truth earns the primary buyer surface.**

12. Defensibility and network effects

A competitor can copy a storefront or video feed. It cannot reproduce the following assets overnight:

1. **Preloaded and verified commerce graph:** one-million-plus masters, claim-ready organisations, national places, markets, time-stamped claims, sources, conflicts and correction history.
2. **Fresh local offer network:** branch-level price, stock, service and fulfilment observations.
3. **Merchant operating workflow:** daily actions, staff permissions, lead history and improvement data.
4. **Trust and outcome history:** verification, response, complaint and fulfilment performance.
5. **Geo and market relationships:** operators, captains, stalls, clusters, routes and coverage cells.
6. **Creator attribution graph:** content rights, tagged objects, campaigns and verified outcomes.
7. **Buyer demand graph:** searches, saved needs, requests, recommendation responses and unmet local demand.
8. **Advertising performance graph:** local creative, location, product and outcome intelligence.
9. **Framework ecosystem:** integrations and subsidiaries that reuse and improve the same core.
10. **Historical Commerce Archive:** multi-source observations and trend history that are expensive to acquire, clean, license and maintain.
11. **Merchant claim network:** ownership, corrections, branches, staff authority and continuing first-party refresh.
12. **Local model advantage:** Zimbabwe-specific entities, language, compatibility, demand and ranking evaluation built on rights-cleared data.

12.1 Compounding loops

Supply loop

More activated merchants -> more useful offers -> better buyer resolution -> more leads -> stronger merchant retention -> more first-party data.

Demand loop

More buyer searches and requests -> better demand insight -> targeted merchant acquisition -> deeper coverage -> more resolved searches.

Trust loop

More verified outcomes -> stronger reputation -> higher conversion -> more measurable outcomes -> better risk models.

Creator loop

More attributable products -> better creator earnings -> more useful content -> more discovery -> more merchant participation.

Framework loop

More approved extensions -> more reuse and observations -> stronger core graph -> lower marginal cost for new businesses -> more extension demand.

12.2 Copy-resistance rules

- Never reduce the graph to a scrapeable public directory.
 - Keep public APIs scoped and rights-cleared.
 - Preserve evidence and change history internally.
 - Give merchants export rights without exposing other merchants' confidential data.
 - Use contracts, rate limits, anomaly detection and canary fields against bulk extraction.
 - Do not claim proprietary ownership over third-party data beyond legal rights.
 - Make service quality, field relationships and operational learning as important as code.
-

13. Daily utility and responsible habit formation

The product should become a daily habit because it resolves real tasks, not because it manipulates attention.

13.1 Buyer return reasons

- stock and price alerts;
- saved needs and matched requests;
- nearby availability and open-now updates;
- trusted deal and new-stock feeds;
- creator buying guides;
- repeat household or business lists;
- service reminders, compatibility and replacement needs;
- market events and local commercial news.

13.2 Merchant return reasons

- stale stock and price actions;
- new leads and response queues;
- campaign and creator results;
- product-demand signals;
- unresolved buyer requests;
- trust, complaint and fulfilment actions;
- daily Business Improvement Copilot recommendations;
- branch and market performance.

13.3 Responsible engagement

Shoppage should not use infinite-scroll coercion, deceptive notifications, fake scarcity, hidden sponsorships, unreadable cancellation or fabricated social proof. Notification frequency must be controllable. Video feeds should have natural stopping points and data controls.

13.4 High-value product systems

Priority systems include:

- Shoppage Ask / Request Network;
- Price Pulse;
- Stock Truth;

- Compatibility Graph;
- Merchant Reputation Passport;
- Verified Collection Network;
- Universal Merchant Link and QR;
- WhatsApp and voice ingestion;
- Diaspora Purchase Concierge;
- B2B Procurement Exchange;
- Made in Zimbabwe Discovery;
- Proof-backed reviews;
- Commerce Academy;
- Business Improvement Score;
- Community Correction Network;
- Commerce Continuity View.

These are sequenced by evidence and roadmap gates, not launched simultaneously.

14. AI-native operating model

AI is a productivity, quality and intelligence layer. It is not permission to publish unverified claims or generate fragile software.

14.1 Launch AI capabilities

Merchant operations

- extract draft products from images, PDFs, spreadsheets and messages;
- clean titles and descriptions;
- suggest categories, attributes and variants;
- identify duplicate products and offers;
- generate captions, cards and price overlays;
- prioritise stale or incomplete records;
- assist support and onboarding agents.

Buyers and discovery

- interpret natural-language and voice requests;
- expand local synonyms and brand nicknames;
- match requests to qualified merchants;
- recommend compatible or locally available alternatives;
- produce cold-start recommendations from the preloaded graph before behavioural scale exists;
- distinguish global reference matches from current Zimbabwe offers;
- summarise differences with uncertainty shown.

Trust and operations

- detect suspicious pricing, duplicate scam images and anomalous behaviour;
- prioritise moderation and resolution queues;
- identify data-quality regressions;
- draft, but not decide, high-impact cases.

14.2 AI data and model architecture

The AI layer must use the Commerce Intelligence Graph as a governed retrieval and feature system. Product text, social history, purchased directories, partner data and merchant feeds enter separate provenance partitions. Training and retrieval access is granted by rights class rather than by technical availability.

The model stack should use provider-independent adapters, an internal model registry, reproducible evaluation sets, shadow tests and fallbacks. High-volume tasks such as classification, extraction and embeddings should use the cheapest model that

meets measured quality. Sensitive or rights-restricted data should not be sent to third-party models without an approved processing basis and contract.

Core launch evaluations include entity-resolution precision, category/attribute extraction, multilingual query recall, recommendation relevance, geographic validity, offer freshness, hallucination rate, bias across merchant size/location and human-review workload.

14.3 Human decision boundaries

Human approval and appeal are required for:

- identity-verification rejection;
- serious fraud labels;
- account suspension or permanent removal;
- disputed payment or fulfilment liability;
- high-impact entity merges;
- unsafe-category publication;
- material public or government reporting.

14.4 AI evaluation

Each production model requires:

- defined task, permitted inputs and prohibited outputs;
- source-rights and privacy clearance for training/evaluation data;
- Zimbabwe-specific evaluation set and simple non-AI baseline;
- acceptance threshold for relevance, locality, compatibility and unsupported claims;
- error taxonomy by language, category, location and user group;
- human-review coverage for high-impact decisions;
- cost, latency and availability budget;
- deterministic fallback when unavailable;
- staged rollout, drift/abuse monitoring and rollback path;
- named owner and signed release record.

At launch, recommendation and classification models are advisory. AI cannot publish a new merchant, price, stock state, legal status, market containment or verification badge without an approved evidence path.

14.5 AI-assisted development controls

- modular architecture and typed contracts;
 - tests written with or before generated code;
 - human review for authentication, payments, permissions and data migrations;
 - dependency and licence scanning;
 - no secrets in prompts or repositories;
 - architecture decision records;
 - staging, feature flags and rollback;
 - production observability from the first release.
-

15. Trust, safety and consumer protection

Trust is a product system, not a badge.

15.1 Progressive verification

Verification is predicate-specific and must be financially and operationally sustainable.

Level	Evidence	Public meaning
Sourced reference	One approved source	A dated source reported this fact; not merchant-managed
Corroborated identity/location	Authoritative source or multiple independent sources	Defined name/location/contact predicates are supported
Claimed representative	Ownership/authority workflow passed	An authorised user manages current merchant fields
Remote verified	Phone/email/domain/document or partner evidence	Only the listed predicates were checked
Field verified	Agent visit, geotagged evidence and review	Physical presence was checked on the stated date
Regulated/authorised	Current regulator, licence, bank or brand evidence	The specific licence/relationship was checked

The 10,000+ launch profiles are not assumed to be field verified. Field verification is reserved for launch merchants, high-risk categories, paid verification products and cases where it materially improves trust. Its unit cost must be measured separately.

Every badge must state what was checked, when, by whom/which evidence class and when it expires. “Verified” cannot imply general endorsement, product authenticity, solvency or guaranteed fulfilment beyond the stated scope.

15.2 Reputation

Merchant reputation should combine:

- time on platform;
- last active and freshness behaviour;
- response rate and time;
- accepted/declined request history;
- verified collection or fulfilment outcomes;
- complaint severity and resolution;
- policy and warranty adherence;
- data accuracy corrections.

15.3 Proof-backed reviews

Reviews should be weighted by evidence such as verified contact, quote, reservation, payment, collection or fulfilment. Unverified comments may be displayed separately or used as reports, not treated as equal transactional proof.

15.4 Category safety

High-risk categories require stricter controls or exclusion. Policies must cover counterfeit goods, unsafe products, regulated goods, prohibited services, misleading health/financial claims, age-restricted goods, stolen goods and intellectual-property complaints.

15.5 Disputes and complaints

The system needs:

- clear report types;
- evidence upload;
- response deadlines;
- merchant and creator appeal;
- neutral review ownership;
- severity-based escalation;
- outcome and remediation history;
- regulatory or law-enforcement referral protocol where legally required.

16. Payments and transaction ownership

16.1 Binding decision

Shoppage owns the commercial workflow and evidence record. Licensed providers process and settle money. The vendor remains the seller and merchant of record for ordinary vendor sales unless a later, explicitly contracted model states otherwise.

Shoppage will not initially:

- hold customer funds;
- pool merchant settlements;
- operate a deposit wallet;
- store card numbers, CVVs or mobile-money PINs;
- promise escrow without a licensed structure;
- become merchant of record by default.

Zimbabwe already has domestic online-payment infrastructure, including Zimswitch Online and Paynow merchant/e-commerce services.[4][5]

16.2 Payment methods

- cash on collection;
- cash on delivery where a merchant/partner supports it;
- mobile money;
- bank transfer or ZIPIT-style transfer;
- provider-hosted card or online checkout;
- payment link;
- deposit or reservation payment;
- later buyer-protection or escrow-like products through licensed partners.

16.3 Shoppage responsibilities

- display accepted methods and conditions;
- initiate provider payment;
- create order/reservation reference;
- receive signed status callbacks;
- reconcile status and exception records;
- issue Shoppage transaction evidence and merchant receipt links;
- preserve dispute and refund evidence;
- calculate platform fees where contracted;
- prevent ads or content from falsely implying payment protection.

16.4 Provider-adapter standard

Every adapter must support idempotency, signed webhooks, replay protection, timeout handling, reconciliation, refund/status mapping, sandbox testing, incident ownership and provider exit procedures.

16.5 Ad credits

Shoppage may sell prepaid promotional credits for platform advertising, provided they are clearly limited-purpose, non-withdrawable and accounted for as service credits. They must not be marketed or operated as a general stored-value wallet.

17. Fulfilment and last-mile delivery

17.1 Binding decision

Physical fulfilment is partner-led and merchant-led, but it cannot be invisible or completely off-platform. Shoppage orchestrates options, quotes, statuses and evidence without initially owning a national fleet or warehouse network.

17.2 Supported modes

- buyer collection from branch or stall;
- verified collection point;
- merchant delivery;
- courier or delivery partner;
- installer or technician fulfilment;
- long-distance bus/courier route;
- buyer-arranged bulk haulage;
- scheduled market or event pickup.

17.3 Shoppage responsibilities

- collect dimensions, weight, category and handling constraints where relevant;
- show service area and delivery assumptions;
- request and compare quotes;
- capture selected provider and reference;
- display dispatch, pickup and delivery states;
- collect proof of collection, delivery or installation;
- record failed attempts, cancellations and complaints;
- maintain merchant/partner performance histories;
- separate delivery fee, platform fee and merchant price.

17.4 Category-specific fulfilment

Category	Appropriate launch approach	Later depth
Small parcels/electronics	Standard rates, tracking and proof	Automated partner selection
Building materials	Quote, vehicle type, distance and offloading	Route/haulage marketplace
Solar/installations	Product plus technician scheduling	Certification and job milestones
Furniture/appliances	Scheduled delivery and condition proof	Reverse logistics and installation
Informal markets	Verified collection desk or captain coordination	Shared market fulfilment hub

17.5 Asset-ownership gate

Owned vehicles, warehouses or dark stores require a specific zone/category business case with sustained density, positive fully loaded economics, unacceptable partner failure and dedicated operational leadership. They are not default strategic ambitions.

18. Data governance, rights and compliance

18.1 Data protection

Shoppage requires independent Zimbabwe counsel before public indexing of unclaimed business profiles or processing at scale. Counsel must review data-controller obligations, data-subject rights, lawful bases, direct marketing, cookies, cross-border processors, children, sensitive data, breach response and the relationship between business-directory facts and personal data.

Core controls:

- purpose limitation and data minimisation;
- transparent notices for claimed and unclaimed profiles;
- role-based access and least privilege;
- consent or another documented lawful basis where required;
- access, correction, objection, suppression and deletion workflows;
- 24-hour action SLA for materially harmful disputed unclaimed-profile fields;
- retention and deletion schedules;
- cross-border processor register and contracts;
- encryption, logging, incident response and breach records;
- privacy impact assessment for historical social data and AI corpora;
- vendor, plugin and partner data-processing agreements.

18.2 Rights and licensing

No source may enter public search, external claims, model training or partner APIs until it has a completed Source Rights Register record and an accountable owner.

The register must record:

- source owner, dataset and version/hash;
- acquisition method and source class;
- licence, contract or permission evidence;
- public display, indexing, derived-claim, analytics, redistribution and AI-training permissions;
- attribution, share-alike, field-level and media restrictions;
- personal-data classification and lawful basis;
- update cadence, expiry and termination risk;
- correction/takedown channel and response obligation;
- deletion-propagation plan covering indexes, caches, derived features, embeddings, training snapshots and partner exports;
- legal-review status, owner and next review date.

Purchased directories and collaboration feeds require written warranties; “bought” does not mean rights-cleared. Approved social-source methods are platform-authorized APIs, merchant/user-provided exports, partner agreements and specifically cleared observations. Public visibility alone does not authorise unrestricted automated collection.

The Open Food Facts-derived reference requires explicit ODbL/database-content and image-licence review. Source-derived, merchant-proprietary and Shoppage-generated data must remain technically separable where licence or deletion obligations require it.

Before public launch, Shoppage must pass a deletion-propagation drill that removes a test source from public search, caches, current-state materialisations, recommendation features and downstream exports without corrupting unrelated records.

18.3 Merchant and buyer data

- merchants own their original content subject to the licence granted to operate Shoppage;
- merchants can export their own catalogue and transactions;
- Shoppage owns platform-generated models, aggregate analytics and system records subject to law and contracts;
- one merchant’s confidential performance cannot be exposed to another;
- personal data is not sold;
- aggregate intelligence uses minimum cohort sizes and suppression controls;
- partner access is purpose-limited, logged and revocable.

18.4 Consumer and advertising compliance

Before launch, counsel should review:

- Consumer Protection Act requirements;
- pricing and discount disclosures;
- electronic contracting and records;

- platform/vendor liability statements;
- returns, warranties and refund terms;
- creator and sponsored-content disclosure;
- prohibited and regulated goods;
- courier and postal licensing implications;
- tax invoices, commissions and creator payments;
- terms for market captains and field agents.

19. Revenue model and pricing architecture

Shoppage needs multiple revenue streams, but each must reinforce platform utility rather than tax immature demand.

19.1 Revenue hierarchy

Revenue stream	Timing	Economic character
Assisted onboarding and catalogue services	Launch	Early cash, labour-intensive
Managed campaigns and content production	Launch	Early revenue and learning
Sponsored products/sites/clips/markets	Launch with demand limits	Scalable media revenue
Merchant subscriptions	Launch and expansion	Recurring software revenue
Qualified request/lead products	Launch in selected categories	High-value performance revenue
Creator campaign administration	Launch/expansion	Service and platform fee
Pemosites and premium Merchant Sites	Months 4-6	Recurring plus setup revenue
Intelligence products	Launch internally; Months 4-6 external products	High-margin, rights-sensitive and evidence-dependent
Framework/API/plugin fees	Private beta then later	Platform revenue
Payment/fulfilment orchestration fees	Later where contractually justified	Transaction-linked
Native transaction commission	Later	Requires strong trust and operations

19.2 Pricing hypotheses

Launch pricing is intentionally simple and remains a pilot hypothesis.

- **Free foundation:** Merchant Site, up to 20 active offers, QR/canonical links, contact actions, basic claim/correction tools and basic analytics.
- **Growth Pilot:** target test price of approximately US\$19 per month or an approved local-currency equivalent; expanded catalogue, Request Network access, freshness automation, ad eligibility, clips, staff access and improved analytics.
- **Assisted onboarding/content:** separately priced service with labour and gross margin tracked independently.
- **Managed advertising:** fixed packages sold only in usable coverage cells.
- **Enterprise/market/brand:** contracted pilots after core reliability.

Do not introduce multi-tier pricing complexity during the first controlled cohort. The price changes when evidence changes.

A merchant-paid product is viable only when:

- attributable or strongly evidenced value exceeds the fee;
- activation and support labour produce an acceptable payback;
- churn, freshness and response behaviour remain healthy;
- merchants understand what is included without sales explanation.

The field gate tests willingness-to-pay, gross profit per converted lead and activation cost. Global or South African SaaS benchmarks are not evidence of Zimbabwean willingness-to-pay.

19.3 Unit-economics discipline

Track separately:

- software gross margin;
- managed-service gross margin;
- creator payouts and media cost;
- field activation cost;
- support cost by merchant tier;
- payment and fulfilment pass-through;
- infrastructure cost per active merchant, search and video minute;
- customer acquisition cost and payback;
- revenue concentration.

Do not hide labour-intensive onboarding inside “software” margin.

20. Competitive strategy and platform substitution

20.1 Current alternatives

Zimbabwean buyers and merchants already use Facebook Marketplace, Facebook pages/groups, WhatsApp, Google Search/Maps, Classifieds.co.zw and marketplace sites such as Ownai, ShopAtZim and Hypermarket. Current public sites show that Classifieds offers broad category listings and free ads, Ownai offers a multi-category cart and comparison experience, and Hypermarket positions around multi-vendor products, verified sellers and Paynow checkout.[7][8][9]

Shoppage must not claim competitors have no value. It wins by solving a different combination of problems better.

20.2 Outclass strategy

Existing channel strength	Structural gap	Shoppage advantage to build
Facebook reach and communities	Weak structured catalogue, freshness and attribution	Merchant truth, product graph and measurable actions
WhatsApp conversion and relationships	Poor discovery, comparison and analytics	Searchable catalogue and structured handoff
Google broad search and maps	Limited branch-level local stock and trust history	Local offer, freshness, market and response intelligence
Classifieds breadth	Listing-centric, limited merchant operating system	Recurring merchant workflow and evidence graph
Marketplace checkout	Transaction-focused but limited media/intelligence depth	Search, video, ads, creators, markets and platform framework
TikTok/Reels entertainment	Attention not necessarily linked to current local offer	Every clip linked to merchant/product/market and action

20.3 Google Shopping substitution

Shoppage should aim to resolve Zimbabwe-specific shopping questions with information generic search often lacks:

- exact variant and compatibility;
- current branch availability;
- local price date and currency;
- merchant response and verification history;
- collection and delivery feasibility;

- market/stall context;
- local alternatives and requests.

The objective is category-by-category commercial search leadership, not replacing general web search.

20.4 Google/Meta advertising substitution

Merchant budget shifts only if Shoppage demonstrates better cost per qualified local action. Shoppage should baseline external ad spend for participating merchants, run controlled campaigns, report comparable outcomes and calculate wallet-share movement.

External platforms remain optional acquisition channels during the transition. Cards, clips and links should be optimised for Facebook, WhatsApp, TikTok and Google while merchant identity, canonical content, requests, leads and outcome data remain on Shoppage. No external channel may become a critical-path dependency or the only measurable source of buyer demand.

20.5 Social-platform independence strategy

Shoppage should win with external platforms where they provide reach, but survive without them. Acquisition may include user-initiated Facebook/WhatsApp sharing, permitted feeds, creator distribution and public-search indexing. The product must continue to acquire and retain users through Merchant Sites, QR infrastructure, saved needs, price/stock alerts, the Request Network, market pages, direct notifications and referrals based on resolved needs.

Management should review **Owned Demand Share** and **External Platform Concentration** monthly. A rising share of direct, saved, QR, request and notification-originated sessions is evidence that Shoppage is becoming a platform rather than a social-selling utility.

20.6 Competitive-response system

- monthly competitor capability review;
 - quarterly merchant and buyer interviews;
 - public feature/request intake;
 - two-week delivery cycles for measurable improvements;
 - feature flags and experiments;
 - “adopt, adapt or reject” decision log;
 - no feature matching without a clear user or economic outcome.
-

21. Shoppage Framework and ecosystem

21.1 Framework decision

Shoppage is both a product family and a reusable commerce framework. The framework allows approved subsidiaries, partners and independent businesses to build specialised applications without recreating identity, product, place, trust, payment, fulfilment and intelligence systems.

21.2 Commerce kernel

The kernel provides:

- identity, organisations, roles and consent;
- place, market and branch graph;
- catalogue, variants and offers;
- search, recommendations and saved needs;
- trust, verification and disputes;
- leads, RFQs, reservations and orders;
- advertising, content and creator attribution;
- payment and fulfilment orchestration;
- billing and entitlements;
- analytics, evidence and AI services.

21.3 Extension types

Extension type	Example	Governance
First-party module	Price Pulse or Merchant CRM	Core roadmap and shared P&L
First-party vertical	Construction procurement	Dedicated product owner, shared kernel
Product line	Pemosites	Shared data and internal cost centre
Subsidiary	Proven specialised business	Admission gates and licence agreement
Certified plugin	POS, ERP, payment or fulfilment adapter	Review, scopes, support and certification
Powered application	Independent market or distributor portal	API contract, consent and usage billing
White-label deployment	Association or institutional commerce portal	No database fork; central identity and evidence rules

21.4 Platform rules

- no direct third-party database access;
- scoped credentials and least privilege;
- signed webhooks and idempotent writes;
- merchant consent for plugin access;
- versioned APIs and deprecation policy;
- audit of every partner action;
- data export and termination procedure;
- certification and reliability score;
- no arbitrary third-party code inside core production processes;
- no public app store in months 0-6.

21.5 Subsidiary admission gate

A candidate becomes a subsidiary only when it has:

- distinct customer and value proposition;
- evidence of demand and repeat usage;
- identifiable revenue and cost structure;
- dedicated accountable leader;
- compliance and risk assessment;
- strategic reason not to remain a module;
- platform licence and data-sharing agreement;
- board-approved capital plan.

21.6 Ecosystem economics

Possible framework revenue includes API usage, plugin certification, revenue share, platform licence, white-label subscription, managed integration, premium support and intelligence access. These should begin after core reliability and merchant consent are proven.

22. Go-to-market and field operations

22.1 Supply-first launch

The first problem is useful supply, not app downloads or imported profile counts. Shoppage must complete a pre-mobilisation field economics test before the 12-week public-launch clock begins.

The test cohort is the first 30 merchants in the launch wedge. It records:

- acquisition channel and signed interest/participation evidence;
- staff hours and cash cost to claim, verify and activate;
- products/offers loaded and percentage accepted without rework;
- merchant training time and support contacts;
- 7- and 14-day freshness retention;
- response rate and median response time;
- willingness-to-pay and reasons for refusal;
- gross-margin/value evidence from leads or requests.

Proceed when the cohort demonstrates a credible path to payback within six months or another board-approved threshold. If it does not, change onboarding method, pricing, category or scope before scaling.

Merchant acquisition channels include direct field sales, distributors/brands, market operators, associations, creators, payment/POS partners, referrals and managed catalogue days.

22.2 Merchant activation standard

A merchant counts as activated only when:

- authorised representative is established;
- branch/location and contact path are usable;
- minimum catalogue quality is met;
- freshness obligations and auto-hide rules are accepted;
- at least one buyer action can be completed;
- response workflow is demonstrated;
- policies and payment/fulfilment methods are stated;
- staff know how to maintain the account;
- at least 80% of launch offers remain fresh at the 14-day checkpoint.

Profiles loaded from directories are never included in activation counts. Board and investor dashboards must separately report reference profiles, claimed profiles, activated merchants and transacting/responsive merchants.

22.3 Buyer acquisition

Buyer acquisition is organised so that social reach accelerates Shoppage without becoming its oxygen.

Owned and physical distribution

- QR codes at shops, stalls, markets, receipts, quotations, packaging and collection points;
- memorable Merchant Site and market links;
- Request Network flyers and category guides;
- saved needs and consented stock/price alerts;
- market events, referrals and repeat-use loops based on resolved needs.

Partner distribution

- market/mall operators, creators, distributors, brands, installers, payment and fulfilment partners;
- trade associations, campuses and specialist communities.

External-platform acquisition

- 1080 x 1080 Shoppage export cards and copy-ready captions;

- user-initiated Facebook, WhatsApp, TikTok and other sharing;
- permitted platform feeds or publishing adapters;
- public search indexing of Merchant Sites and buying guides;
- paid external acquisition only when measured against direct alternatives.

Native-demand superiority test

Before public launch, run at least 30 representative buyer needs against Shoppage and the incumbent social workflow. The test must measure:

- time to first credible option;
- time to three credible options or an honest no-supply result;
- relevance, locality, freshness and response;
- buyer effort and completion intent.

The public launch gate requires Shoppage to resolve at least 70% of tested needs and demonstrate a materially faster or lower-effort result than the comparison workflow in the selected wedge. The board approves exact thresholds before recruitment communications.

22.4 Field network

Field agents and Market Captains need a controlled work system:

- assigned territory and task;
- source/evidence capture and merchant consent;
- duplicate detection and quality review;
- accepted-work payment;
- code of conduct, identity and complaint process;
- fraud controls and geotagged proof where appropriate;
- productivity, first-pass quality, activation cost and retention metrics.

Field verification is a distinct paid/controlled task. Its cost must never be hidden inside “data acquisition”. No plan may imply that all 10,000 profiles have been physically checked.

22.5 1,000-entity programme

The 1,000-entity programme should build a traceable, reviewed commerce register, not produce a marketing number. It should include merchants, branches, markets, creators, distributors, service providers, fulfilment partners and institutions relevant to launch categories. Each record needs source, predicate, date, verification state and owner.

23. International quality and technical architecture

23.1 Quality standard

Global quality means the core experience is fast, reliable, accessible, secure and understandable on common Zimbabwean devices and networks.

Initial service targets:

Quality area	Controlled-launch target	Expansion target
Availability	99.5% monthly for core buyer/merchant flows	99.9% after operational maturity
Web performance	p75 LCP below 2.5s on representative 4G tests	Continuous device/network monitoring
Search latency	p95 below 500ms for common queries	Category-specific relevance SLAs

Quality area	Controlled-launch target	Expansion target
Crash-free sessions	At least 99.5%	At least 99.8%
Accessibility	Core flows audited toward WCAG 2.2 AA	Full design-system compliance
Data freshness	Category-specific visible SLA	Automated partner/merchant health
Security	No known critical production vulnerabilities	Continuous scanning and formal testing
Support	Severity and customer-class response targets	7-day operational coverage where justified
Offline resilience	Cached sites/items and safely queued drafts after install	Selective local indexes where justified

Targets require measurement; they are not guarantees before instrumentation is live.

23.2 Technical architecture

Recommended controlled-launch stack:

- Next.js App Router PWA for buyer and simplified merchant experiences;
- Refine or equivalent for dense internal/enterprise administration;
- NestJS modular monolith as authoritative business layer;
- Prisma and PostgreSQL/Supabase-managed infrastructure;
- Meilisearch for catalogue and geo-faceted search across one-million-plus masters, organisations, markets and local offers, benchmarked against actual data;
- BullMQ/Redis for imports, indexing, notifications, media and AI jobs;
- managed object storage and a dedicated short-video service;
- vector/embedding and model-feature stores implemented conservatively within PostgreSQL/pgvector or a justified managed service;
- model registry, evaluation harness, corpus manifests and rights-aware retrieval controls;
- structured logs, metrics, traces, error tracking and security monitoring;
- pnpm/Turborepo monorepo and shared contracts;
- optional channel-adapter layer for social sharing, public feeds and external identities, isolated from critical core services.

23.3 Portability

Provider and channel choices must be adapter-based. The platform needs documented export and migration paths for database, object storage, search, video, messaging, AI, payments, hosting and external distribution channels. Supabase, Vercel, Railway, Cloudflare Stream or Mux may be pragmatic providers, but no provider should become an unexamined strategic dependency.

23.4 Production data architecture

The supplied SQLite assets are for reproducible validation and scale testing. Production requires PostgreSQL multi-tenancy, evidence/history tables, materialised current-state views, review queues, search indexes, access controls, backups and migration testing.

23.5 Security baseline

- MFA for platform administrators and high-risk merchant roles;
- row-level and service-layer tenant isolation;
- least-privilege service accounts;
- secrets manager and key rotation;
- signed uploads and malware scanning;
- secure webhook validation;
- dependency and container scanning;
- audit logs for data, permissions, payments and moderation;
- tested backup restore;

- incident response and breach drills;
- penetration testing before material payment/transaction scale.

24. Months 0-3 controlled launch roadmap

24.0 Mandatory pre-mobilisation gates

The 12-week execution clock starts only after the executive owner signs evidence for all four gates.

Gate	Required evidence	Stop condition
Source rights and deletion	Completed source-rights ledger for every launch dataset; counsel review; deletion-propagation drill passes	Any material source lacks permitted-use or deletion path
Merchant activation economics	Signed/qualified first-30 pipeline; measured activation cost/time; credible six-month payback path	No repeatable activation method or economics
Native demand superiority	Controlled Request/Search test demonstrates faster/lower-effort credible resolution than incumbent social workflow	Shoppage does not materially improve buyer outcome
Security, privacy and operating readiness	Independent security/privacy/rights review; incident, support and manual fallbacks documented	Critical defect or unresolved high-risk compliance issue

Minimum mobilisation finance: approved cash/runway must cover the 12-week build, at least three months of post-launch operation and a 25% contingency. Scope reduction authority belongs to the single accountable executive owner and weekly launch room.

AI-assisted development can shorten implementation, but field operations, merchant activation, rights, moderation and adoption remain critical-path work. The plan uses four parallel workstreams: product/engineering, commerce data, field/go-to-market and governance/operations.

24.1 Weeks 1-2: scope, kernel and launch cohort

Product and engineering

- lock domain model, evidence states, public-display rules and event taxonomy;
- create monorepo, CI/CD, environments, auth, tenant isolation and observability;
- implement canonical IDs, external-identity mappings and adapter feature flags;
- convert independence acceptance criteria into automated and manual test cases;
- implement unclaimed-profile banner, correction/suppression intake and SLA monitoring.

Commerce data

- import the reference catalogue into segregated source/licence partitions;
- create the Home, Power and Building corpus, attributes and compatibility review queue;
- ingest 10,000+ profiles only from sources that passed the rights register;
- load place and market records with evidence/review states;
- build deletion propagation and source quarantine tooling.

Field and commercial

- activate first 20-30 anchor merchants from the pre-gate pipeline;
- select 3-5 coverage cells;
- run the simplified Free/Growth Pilot pricing test;
- contract payment, fulfilment and fixed-bounty creator pilots.

Governance

- publish unclaimed-profile, privacy, correction and source policies;
- approve prohibited categories and advertising/content rules;
- establish incident and complaint ownership;
- open the weekly truth-and-gates launch room.

24.2 Weeks 3-4: Merchant Sites and offer truth

- launch merchant claim, ownership-resolution, correction, duplicate-merge and onboarding flow;
- launch Merchant Site templates;
- product-reference attachment, local master/variant creation, branch and offer management;
- image/CSV assisted ingestion;
- price, stock, evidence and category freshness states with 7-day default auto-hide, 72-hour fast-moving override and 14-day made-to-order/service state;
- ranking and ad-eligibility loss for stale offers;
- canonical Universal Links, QR assets, 1080 x 1080 export cards and copy-ready multilingual social captions;
- offline-cached Merchant Site shell and queued merchant drafts;
- internal review queues;
- field capture app/workflow;
- first 500-1,000 usable offers.

24.3 Weeks 5-6: buyer discovery and requests

- keyword, semantic, filter and location search across masters, organisations, markets and local offers;
- product, Merchant Site and market pages;
- manual location and nearby discovery;
- Ask/Request flow and merchant routing;
- WhatsApp/call/quote attribution;
- save/share and basic alerts;
- freshness, operational-proximity and Request Network acceptance tests as the native demand wedge;
- trust and report controls;
- rules-based/hybrid recommendation alpha evaluated against the reviewed 500-query corpus;
- first controlled buyer tests.

24.4 Weeks 7-8: Showcase, creators and markets

- catalogue-linked short-video upload and playback;
- captions, rights and moderation;
- creator applications, fixed-bounty contracts, canonical-link rules and deliverable verification;
- tagged-link measurement without launch-stage performance commissions;
- market pages, rosters, maps and captains;
- content and market campaign analytics;
- low-data and data-saver modes.

24.5 Weeks 9-10: Ads, payment and fulfilment orchestration

- fixed-package sponsored products, sites and clips;
- managed campaign workflow;
- ad eligibility, labels, spend caps and reports;
- provider-hosted payment initiation and callbacks;
- collection, merchant delivery and partner quote options;
- proof/status and dispute evidence;
- no Pemosites work; protect the core Merchant Site engine and launch reliability.

24.6 Weeks 11-12: hardening and controlled launch

- load, security, accessibility and recovery tests;
- search relevance, source-state display, recommendation and AI evaluation;
- content, trust and payment/fulfilment drills;
- merchant training and support playbooks;
- analytics and executive scorecard;

- pilot campaign and fixed creator-bounty payouts;
- social-adapter kill-switch and external-platform outage drill;
- public launch to bounded geography/category;
- day-30 improvement backlog.

24.7 Month-three release gates

- all four pre-mobilisation gates remain valid and no material assumption has deteriorated;
- no critical security defects and tested backup/rollback;
- source-rights ledger complete for every public launch source;
- deletion-propagation and source-quarantine drill passes;
- 1,000,000+ reference products indexed in segregated partitions;
- 20,000+ launch-wedge corpus identities and approved attribute/compatibility coverage;
- 10,000+ claim-ready profiles loaded with banner, claim, correction and suppression paths;
- national place scaffold and reviewed market inventory queryable;
- at least three usable coverage cells;
- 50+ activated merchants or a board-approved evidence-based threshold;
- at least 3,000 fresh local offers or sufficient coverage-cell density;
- 80% freshness within category windows and auto-hide functioning;
- native-demand superiority test passes;
- rules/hybrid recommender passes the evaluation corpus with no unsupported truth claims;
- unclaimed-profile 24-hour action SLA demonstrated in drills;
- ads sold only in usable cells and clearly attributable;
- payment/fulfilment responsibilities documented;
- operating team resolves incidents without engineering improvisation;
- core login, claim, correct, suppress, publish, search, request, lead and maintenance flows pass with all social adapters disabled;
- at least 20 creator bounties have been processed manually before commission engineering is considered.

24.8 Explicit exclusions

- national fleet or warehouse;
- Shoppage wallet or fund custody;
- every category and province;
- public plugin marketplace;
- unrestricted public APIs;
- fully automated verification decisions;
- mature loyalty, escrow, refunds and returns engine;
- native mobile apps unless PWA evidence proves insufficient;
- Pemosites, custom domains, bespoke corporate websites or public template marketplace;
- manual perfection of every one-million-plus reference product before launch; enrichment is automated and prioritised by demand, risk and local relevance;
- unauthorised scraping or automated publishing that violates external-platform rules;
- performance-based creator commissions before attribution, deduplication and dispute gates pass;
- full offline search, live stock or payment-status claims when the device is disconnected.

25. Months 4-6 expansion roadmap

Expansion is conditional on launch evidence, not calendar alone.

25.1 Entry gates

- repeat buyer use in covered categories;
- merchant response and freshness above target;
- measurable campaign outcomes;
- support and moderation capacity;

- payment and fulfilment incident rate within limits;
- positive unit-economics trajectory for at least one revenue stream;
- reliable evidence and rights controls;
- creator link attribution, deduplication, dispute and payout reconciliation pass before outcome commissions are enabled.

25.2 Expansion priorities

- deepen Harare local-offer truth while maintaining national reference discovery;
- increase merchant profile claim, correction and verified-current-operation rates;
- tune the local recommendation/ranking models using measured outcomes and explicit feedback;
- replicate to Bulawayo or the next strongest evidence-based city;
- add selected electronics/auto/agriculture/service verticals only after taxonomy and partner readiness;
- expand creator marketplace and market sponsorships;
- introduce verified lead commissions only for approved categories and creators after the Tier 2 release gate;
- pilot Pemosites with 10-20 businesses;
- introduce self-service ads to proven merchant cohorts;
- launch saved needs, Price Pulse and deeper compatibility;
- add more payment and fulfilment adapters;
- private developer portal and TypeScript SDK;
- three to five certified integrations;
- one powered market, distributor or procurement application;
- privacy-safe aggregate intelligence pilots;
- diaspora purchase and proof-of-fulfilment pilot;
- formal evaluation of one vertical as a product line or subsidiary.

25.3 Expansion rule

A new city or category receives investment only when it has:

- merchant anchors;
- field/partner owner;
- product taxonomy and safety policy;
- sufficient buyer demand or requests;
- payment/fulfilment feasibility;
- support capacity;
- measurable coverage-cell plan;
- budget and stop criteria.

26. Metrics, economics and investor proof

26.1 North-star metric

Resolved Local Commerce Needs: buyer needs that reach a credible outcome through a fresh, attributable Shoppage path.

A resolved need may be a verified merchant contact, accepted quote, reservation, collection, fulfilment or transaction, with confidence levels kept separate.

26.2 Core metric tree

Domain	Core metrics	Guardrails
Supply	Activated merchants, fresh offers, variant depth, coverage cells	Accuracy, rights and merchant retention
Demand	Searches, recommendations, requests, repeat users, resolved needs, owned-demand share	Failed-search, irrelevant recommendations, complaint and external-platform concentration

Domain	Core metrics	Guardrails
Conversion	Search-to-product, product-to-contact, response, quote acceptance	Duplicate/fake leads
Ads	Spend, cost per qualified action, wallet share, repeat advertisers	Organic relevance and complaint rate
Creators	Accepted bounties, useful content, tagged-link learning; later attributable outcomes	Rights, fraud, disclosure, duplicate leads and payout disputes
Markets	Active merchants, roster freshness, market traffic/actions	Operator disputes and stale rosters
Operations	Review backlog, support time, freshness SLA, incidents	Staff workload and quality
Economics	Revenue, gross margin, CAC, payback, retention, concentration	Service labour and imported costs

26.3 Investor proof milestones

By month 3

- controlled product works end to end with social adapters disabled;
- curated supply and buyer actions exist;
- foundation-to-production data path is operational;
- at least one paid revenue stream has real customers;
- no material compliance or security blocker;
- roadmap is evidence-driven.

By month 6

- repeat merchant and buyer behaviour with increasing owned/direct demand share;
- multiple revenue streams with separated margins;
- advertising outcome evidence;
- category/city replication playbook;
- platform reuse through Pemosites or one powered application;
- improving coverage and data quality;
- credible capital plan for the next 12-18 months.

26.4 Investor safeguards

- do not count reference records as active inventory;
- do not count merchant registrations as activated supply;
- separate GMV, leads and verified sales;
- disclose related-party and partner concentration;
- separate software and service margins;
- maintain board-visible data-rights and security incidents;
- gate subsidiaries and capital allocation;
- track imported infrastructure exposure;
- preserve PlatformCo IP and data-contract ownership.

27. Governance, organisation and decision rights

27.1 Initial operating units

- Product and engineering;
- Commerce data and intelligence;
- Merchant success and field operations;

- Buyer growth and content;
- Ads and creator partnerships;
- Trust, safety and support;
- Finance, legal and administration.

A small team may combine roles, but accountability must remain explicit.

27.2 Minimum launch team

Indicative roles:

- product/venture lead;
- technical lead;
- two to four full-stack engineers;
- product/design lead;
- data/search engineer or strong backend equivalent;
- field/merchant activation lead;
- two to six agents or captains depending on target;
- trust/support operator;
- ads/creator commercial lead;
- fractional finance/legal/security specialists.

The plan should be adjusted to actual funding. Targets must be reduced before quality controls are removed.

27.3 Decision rights

- Product scope: product lead with executive escalation.
- Architecture/security: technical lead; security veto on critical risk.
- Verification/trust: independent trust owner with appeal process.
- Ad integrity: separate from sales quota.
- Data rights/privacy: designated owner and counsel.
- Payment/fulfilment provider acceptance: cross-functional risk review.
- Subsidiaries and white-label deals: board approval.
- Public claims and government reporting: evidence owner and executive sign-off.

27.4 Operating cadence

- daily launch operations review;
- weekly product, data and field scorecard;
- fortnightly release cycle;
- monthly unit-economics and risk review;
- quarterly strategy, competition and portfolio review;
- formal incident post-mortems.

27.5 Mission-critical execution doctrine

Shoppage is a compulsory strategic programme in the sense that leadership is committed to solving the problem, not in the sense that risks, laws or evidence can be ignored. Success cannot be guaranteed by a document; it is made more probable through disciplined execution.

Non-negotiable operating rules:

- one accountable executive owner and one weekly cross-functional launch room;
- a single binding backlog tied to v1.1 gates, not parallel wish lists;
- manual operational fallback for claims, ingestion, search correction, leads, payments and fulfilment;
- no single-source dependency for data, distribution, hosting, AI, payments or logistics;
- minimum cash/runway and contingency thresholds before expansion;
- weekly truth review covering claims, source rights, security, buyer harm and merchant outcomes;
- rapid scope reduction when a feature threatens the core search, claim, offer or request flows;
- field operations treated as a first-class product capability;

- unresolved blockers escalated with an owner, decision date and fallback path;
- no fabricated metrics, “all Zimbabwe” claims or hidden uncertainty to protect the narrative.

The answer to obstacles is not excuses and not recklessness. It is redundancy, evidence, iteration, field execution and focus.

28. National impact readiness and institutional engagement

National-scale credibility cannot be manufactured through endorsements and must not substitute for product-market evidence. Shoppage should become institution-ready through measurable, non-partisan public value, independent governance and transparent safeguards.

28.1 Public-value case

- MSME and informal-merchant digitisation;
- locally retained advertising capability and income;
- creator, agent and market employment;
- improved discovery of Zimbabwean products and suppliers;
- price, availability and market transparency;
- diaspora access to verified local commerce;
- better market and local-authority information;
- digital skills and Commerce Academy;
- export-ready supplier discovery;
- commerce-continuity information during shortages or disruption.

28.2 Safeguards

- political neutrality and no political ads in the initial platform;
- no exclusive government data access, ownership or privileged data lane;
- no government or sponsor ranking privilege in organic search;
- no surveillance backdoor or repurposing of buyer behaviour for political profiling;
- no exclusive procurement preference implied;
- aggregate public reporting, not personal surveillance;
- documented lawful-access process;
- independent correction, appeal and transparency reporting;
- public methodology for impact metrics;
- separation of commercial, regulatory and public-interest data products.

28.3 Independent Impact Scorecard and briefing pack

The primary national-scale artifact is an independently auditable Impact Scorecard, not a request for endorsement. It should report:

- active and retained MSMEs digitised;
- fresh local offers and coverage cells;
- percentage of covered searches resolved with confirmed stock or a credible request response;
- merchant advertising-wallet share and net local advertising value retained;
- creator, agent and market income;
- buyer response, complaint and trust outcomes;
- geographic and language inclusion;
- privacy, correction and transparency indicators.

The briefing pack may include a two-page impact case, live merchant/buyer/creator demonstration, 100-day scorecard, privacy and governance note, partnership menu, local-value-retention methodology and clear evidence boundaries. It should be usable by investors, DFIs, associations, municipalities, regulators, ministries and national leadership without promising political patronage or exclusive access.

29. Risk register and decision gates

Risk	Primary mitigation	Stop/escalation gate
Platform scope exceeds capacity	Narrow wedge, explicit exclusions and weekly gates	Reduce targets before removing controls
Stale or false inventory destroys trust	Freshness SLA, downgrade/hide and correction network	Suspend ads and visibility for repeat failure
Merchant activation is labour-heavy	Assisted tools, distributor feeds and agent economics	Reprice or narrow cohort if activation cost is unsustainable
Buyer demand is insufficient	Requests, creator distribution, SEO and market QR	Do not scale ad sales without organic resolution
Ads harm relevance	Eligibility floors and independent integrity	Pause product/advertiser with complaint or relevance breach
Fraud/scams	Progressive verification, anomaly detection and human review	Category/merchant suspension and appeal
Payment liability confusion	Provider-hosted processing and clear merchant-of-record terms	Disable provider/flow until reconciliation is reliable
Fulfilment failure	Mode-specific evidence and partner performance	Remove unreliable partner or service area
Data-rights breach	Source inventory, licences and access controls	Quarantine affected dataset and notify stakeholders
AI publishes unverified truth	Draft-only controls and evaluation	Roll back model/feature
Pemosites becomes an agency distraction	Shared templates and spinout gates	Stop custom work that cannot be productised
Framework opens too early	Internal-first APIs and private beta	No public ecosystem until security/reliability gates pass
Subsidiary proliferation	Admission tests and separate P&Ls	Board rejects unfunded or unproven venture
Political capture or endorsement dependence	Neutral governance, Impact Scorecard and no privileged access/ranking	Decline conditions that compromise independence
External-platform dependency	Canonical Shoppage identity/content/demand, adapters and kill-switch tests	Disable adapter and remediate if a core flow depends on it
Creator commission fraud	Fixed launch bounties, staged attribution and dispute windows	Do not enable outcome commissions until gates pass
Infrastructure cost or lock-in	Adapters, budgets and export tests	Migrate or redesign before scale

| Unclaimed-profile liability | Exact banner, source labels, suppression and 24-hour action SLA | Suppress disputed fields/profile and investigate | | Reference library overwhelms local truth | Separate search partitions and coverage-cell gates | Remove reference-only results from default buyer surfaces | | Merchant activation economics fail | First-30 cohort and payback gate | Redesign onboarding, price or wedge before mobilisation | | Native demand does not beat incumbent workflow | Controlled comparison test | Do not public launch; improve Request/search/distribution | | Behavioural AI launched too early | Rules/hybrid launch and 5,000-resolved-needs gate | Disable learning-to-rank/personalisation | | Source deletion cannot propagate | Source lineage and quarantine tooling | Block source from production and public claims |

29.1 Binding decision gates

- **Pre-mobilisation gate:** source rights, activation economics, native-demand superiority and independent security/privacy/rights evidence all pass before the 12-week clock starts.
- **Public launch gate:** end-to-end flows, usable coverage cells, fresh supply, support and liability controls pass.
- **Ads scale gate:** repeat direct demand and attributable outcomes exist; no ads in non-usable cells.
- **AI learning gate:** at least 5,000 resolved needs or approved equivalent evidence, plus offline/staged evaluation.
- **Creator commission gate:** attribution, deduplication, fraud, dispute and payout controls pass after fixed-bounty evidence.
- **Transaction-depth gate:** payment/fulfilment reconciliation and dispute capacity pass.

- **City/category gate:** coverage, taxonomy, owner, demand, partners and budget are ready.
- **API/ecosystem gate:** stable contracts, consent, security, audit and revocation pass.
- **Pemosites gate:** month six completed and core merchant/freshness/reliability thresholds pass.
- **Subsidiary gate:** demand, economics, leadership and risk justify separation.
- **Capital gate:** milestones, runway and unit-economics evidence support investment.

30. Final business model and immediate actions

30.1 Final model

Shoppage is Zimbabwe's independent, preloaded commerce operating system: one public product for selling, discovering and growing, launched with a million-plus master catalogue, a claim-ready national organisation and place graph, known markets, historical intelligence, cold-start recommendations and a structured Request Network, then continuously converted into fresh merchant-owned offers and verified commercial truth.

Its durable advantage is the combination of:

- authoritative merchant operations;
- fresh local offers;
- structured product and variant intelligence;
- place and market relationships;
- buyer needs and unresolved demand;
- trust and outcome history;
- attributable creator and advertising data;
- payments and fulfilment evidence;
- reusable platform infrastructure.

30.2 Immediate executive actions

1. Approve v1.1 as the binding master and archive v1.0 as the preloaded-intelligence decision record.
2. Appoint one accountable executive launch owner with scope-reduction and pause authority.
3. Freeze the Harare Home, Power and Building wedge and public launch exclusions.
4. Complete the Source Rights Register for every proposed launch source.
5. Commission independent Zimbabwe counsel and independent security/privacy review.
6. Implement and publish the unclaimed-profile policy, banner, correction/suppression flow and 24-hour action SLA.
7. Run the first-30 merchant activation economics cohort and secure a qualified launch pipeline.
8. Run the native-demand superiority test against current social workflows.
9. Build and execute the social-adaptor-disabled independence drill.
10. Create the 500-query AI evaluation corpus and lock the rules/hybrid launch baseline.
11. Approve simplified Free/Growth Pilot pricing and separate assisted-service economics.
12. Contract pilot payment, fulfilment and fixed-bounty creator partners.
13. Approve minimum runway and 25% contingency before starting the 12-week clock.
14. Convert this document and the accompanying gate workbook into the engineering and operations backlog.

Appendix A. Prioritised 95-aim portfolio

The 80 aims remain valuable, but they are no longer described as simultaneous mandatory deliverables. They are classified as **Launch (0-3 months)**, **Expansion (4-6 months)** or **Later option (after six months)**. Every aim requires an owner and proof metric before it enters a sprint.

Aim	Horizon	Proof metric
1. Become the default local product-search habit in covered categories	Expansion	Repeat direct search and unaided recall
2. Capture measurable merchant	Expansion	Percentage of attributable spend

Aim	Horizon	Proof metric
advertising-wallet share		moved
3. Retain more advertising value locally	Expansion	Net local value-retention ratio
4. Operate independently while using external platforms as optional feeders	Launch	Core-flow kill-switch pass and owned-demand share
5. Own commercial intent category by category	Launch	Resolved-need share in selected cells
6. Build Zimbabwean commercial keyword intelligence	Launch	Query success and synonym coverage
7. Establish a transparent local ad standard	Launch	Advertiser comprehension and complaint rate
8. Build a local commerce-media partner network	Later	Active publishers and partner revenue
9. Resolve needs faster than social-group posting	Launch	Median time to credible option
10. Show stock confidence instead of false certainty	Launch	Freshness visibility and error reports
11. Support exact variant discovery	Launch	Variant match and correction rate
12. Provide useful price intelligence	Expansion	Coverage and alert usefulness
13. Offer available-today discovery	Launch	Same-day availability actions
14. Automate saved needs	Expansion	Alert-to-action conversion
15. Route buyer requests	Launch	Request response and fill rate
16. Support voice-first search	Expansion	Successful voice-query resolution
17. Support English, Shona and Ndebele progressively	Expansion	Query coverage and satisfaction by language
18. Design for low data and unstable networks	Launch	Page/media data budget and completion rate
19. Compare total buyer outcome	Expansion	Comparison-to-action conversion
20. Build consent-based personal commerce memory	Later	Repeat-use benefit and opt-out rate
21. Give every merchant a credible mobile storefront	Launch	Activated Merchant Sites
22. Reduce publishing time through assisted ingestion	Launch	Median time to publish accepted offer
23. Turn every product into reusable content	Launch	Content assets per active offer
24. Create a daily merchant action queue	Launch	Action completion and freshness improvement
25. Measure lead quality, not only volume	Launch	Accepted qualified-action rate
26. Preserve customer relationship history	Launch	Lead follow-up and repeat-action rate
27. Deliver actionable merchant analytics	Launch	Recommended-action completion
28. Add inventory/POS integrations	Expansion	Synced merchants and feed health
29. Build merchant team and branch management	Expansion	Multi-user/branch retention

Aim	Horizon	Proof metric
30. Create merchant benchmarks	Expansion	Benchmark coverage and action lift
31. Build a canonical product and variant graph	Launch	Resolution accuracy and variant completeness
32. Capture branch-level local offers	Launch	Fresh local offers per coverage cell
33. Maintain evidence and provenance	Launch	Claims with complete evidence metadata
34. Maintain correction and review queues	Launch	Backlog SLA and overturned-decision rate
35. Separate candidates from verified production truth	Launch	Public-claim audit pass
36. Build coverage cells	Launch	Cells reaching usable/strong state
37. Map compatibility and substitution	Expansion	Supported category compatibility coverage
38. Capture unmet demand	Launch	Structured failed searches and requests
39. Build price and availability history	Expansion	Historical observations and alert accuracy
40. Build merchant and market performance history	Expansion	Outcome-linked records
41. Create a Zimbabwe business knowledge layer	Launch	Verified capabilities beyond ordinary directories
42. Create correction and verification networks	Expansion	Accepted correction quality
43. Protect graph from scraping and poisoning	Launch	Abuse blocks and anomaly response
44. Monetise aggregate intelligence without personal-data sale	Later	Approved products and privacy audits
45. Make self-service advertising understandable	Launch	Campaign setup completion
46. Guarantee relevance floors for sponsored results	Launch	Sponsored relevance and complaint rate
47. Measure creator links without tying launch pay to fragile attribution	Launch	Tagged-link completeness and fraud-learning quality
48. Build creator marketplace matching	Expansion	Match acceptance and outcomes
49. Pay creators through accepted-work bounties first, then verified outcomes	Launch	Bounty acceptance, rights compliance and payout disputes
50. Build local creative intelligence	Expansion	Campaign improvement from recommendations
51. Launch market and area sponsorships	Expansion	Sponsorship renewal and local actions
52. Enable brand-to-stockist campaigns	Expansion	Campaigns linked to fresh authorised stock
53. Use controlled Shoppage ad credits	Expansion	Credit activation and paid conversion
54. Build campaign experiments	Expansion	Valid tests and lift measurement

Aim	Horizon	Proof metric
55. Establish independent ad integrity	Launch	Review independence and complaint outcomes
56. Digitise markets as operating units	Launch	Useful market pages and roster freshness
57. Create Market Captain roles	Launch	Accepted work and merchant retention
58. Deploy QR infrastructure	Launch	QR sessions and attributable actions
59. Map commercial corridors and informal landmarks	Expansion	Search/fulfilment success outside street addressing
60. Build verified collection points	Expansion	Successful collections and incident rate
61. Promote market events and live discovery	Expansion	Event traffic and merchant outcomes
62. Replicate province by province using gates	Later	City/category gate pass rate
63. Build progressive verification	Launch	Predicate-level verification coverage
64. Use proof-backed reviews	Expansion	Reviews with verified evidence
65. Detect suspicious pricing and scam content	Launch	Precision, recall and human-review outcomes
66. Create category-specific safety policies	Launch	Policy coverage and incident rate
67. Support realistic payment paths	Launch	Payment-option completion and reconciliation
68. Add buyer protection gradually	Later	Reserve/process readiness and claim outcomes
69. Integrate delivery as a quoted network	Launch	Quote acceptance and fulfilment evidence
70. Measure fulfilment reliability	Expansion	Delivery/collection performance
71. Set international performance budgets	Launch	SLO dashboard and target compliance
72. Build accessible interfaces	Expansion	WCAG audit and assisted-user testing
73. Use a modular shared platform	Launch	Core reuse across five layers
74. Maintain provider and social-channel portability	Launch	Tested export, kill-switch and adapter paths
75. Create an AI evaluation harness	Launch	Model release-gate compliance
76. Keep humans in high-impact decisions	Launch	Human-review and appeal coverage
77. Use AI to reduce cost, not conceal weak operations	Launch	Cost/quality improvement without error growth
78. Run continuous competitor response	Expansion	Evidence-based releases and impact
79. Build observability from day one	Launch	Incident detection and diagnosis time
80. Treat security and privacy as product quality	Launch	Security/privacy control completion
81. Launch with one-million-plus	Launch	Indexed masters and source

Aim	Horizon	Proof metric
master products		partitions
82. Load at least 10,000 claim-ready organisation profiles	Launch	Profiles, source evidence and duplicate rate
83. Make every preloaded merchant profile claimable	Launch	Claim completion, disputes and time to resolution
84. Maintain national place and known-market breadth	Launch	Designated-source coverage and correction rate
85. Separate reference products from local offers visibly	Launch	False-availability incident rate
86. Build a Historical Commerce Archive	Launch/Expansion	Rights-cleared observations and usable trend depth
87. Ingest collaborator and purchased data safely	Launch	Contract coverage and rights exceptions
88. Deploy cold-start recommendations before behavioural scale	Launch	Relevance, coverage and fallback pass rate
89. Build the Shoppage Local Commerce Intelligence Model	Expansion	Evaluated local retrieval/ranking quality
90. Preserve removal and correction through derived models	Launch	Deletion-propagation SLA and audit pass
91. Measure merchant-profile claim conversion	Launch	Claimed/eligible profile rate
92. Build product compatibility and substitution intelligence	Launch/Expansion	Successful alternative/compatibility outcomes
93. Avoid single-source dependence in every core graph	Launch	Source concentration by domain
94. Use AI to enrich at scale without inventing facts	Launch	Human acceptance and hallucination rate
95. Maintain manual fallback for every mission-critical flow	Launch	Tested fallback coverage

Appendix B. Strategic extension portfolio

The framework supports many future businesses, but the portfolio is not a simultaneous roadmap.

Merchant infrastructure

Pemosites; POS and stock sync; digital catalogue/QR ordering; branch/franchise tools; CRM; quotations/invoices; bookings; content services; data migration; Business Improvement Copilot.

Commerce networks

B2B procurement; distributor portals; construction and installer network; agricultural commerce; local services; logistics/collection exchange; diaspora procurement; export-ready suppliers; Made in Zimbabwe; repair and compatibility network.

Physical markets

Market digital twins; stall/tenant tools; market advertising; events; visitor/demand analytics; collection coordination; loyalty/offers; shopping trails; municipal dashboards; vacancy and merchant-mix intelligence.

Media and advertising

Partner-publisher inventory; creator storefronts; live shopping; deal alerts; sponsored requests; brand/distributor campaigns; area sponsorships; product launches; content-production network; privacy-safe retail-media measurement.

Intelligence

Price/availability; demand gaps; merchant benchmarking; site selection; substitution/compatibility; brand distribution; informal-cluster mapping; consumer needs; supply-chain performance; public-value measurement.

Platform and developer

Commerce APIs; embeddable widgets; payment/fulfilment adapters; themes/plugins; POS/ERP connectors; white-label deployments; identity services; trust APIs; product resolution; approved AI commerce tools.

Portfolio ideas are scored on user pain, strategic fit, data contribution, revenue, capital intensity, regulatory risk, execution capacity and reuse of the core.

Appendix C. Core data entities and events

Core entities

- Person, user, role and consent;
- organisation, legal entity, merchant and branch;
- place, market, building, section, stall and service area;
- category, master product, variant, compatibility and service;
- merchant product, offer, price, inventory observation and promotion;
- media asset, clip, collection and content right;
- creator, affiliate agreement and attribution link;
- search, saved need, request, lead, quote and reservation;
- campaign, placement, spend, ad credit and outcome;
- payment intent, provider event, reconciliation and refund state;
- fulfilment quote, provider, milestone and proof;
- verification, complaint, review, dispute and appeal;
- source, evidence, licence, confidence, review and supersession;
- plan, entitlement, invoice, partner application and API credential.

Minimum event standard

Every event should capture event name, actor, organisation, object, session, timestamp, source surface, location precision, campaign/creator attribution, consent context and relevant outcome state. Sensitive values should not be placed in analytics events when a reference identifier is sufficient.

Appendix D. Source and evidence notes

1. U.S. International Trade Administration, **Zimbabwe - eCommerce**, updated March 2026: social-media-led e-commerce, Facebook Marketplace/WhatsApp, cash on delivery and mobile payments.
2. DataReportal, **Digital 2026: Zimbabwe**: late-2025 mobile connections, internet users and social-media user identities.
3. UNCTAD, **Zimbabwe eTrade Readiness Assessment**: urban/informal and diaspora-driven commerce, logistics, pickup-centre, legal and skills gaps.
4. Zimswitch official site: Zimswitch Online, recurring payment and national-switch positioning.
5. Paynow official site: merchant services and e-commerce payment solutions.
6. POTRAZ official site: Data Protection Authority, Cyber and Data Protection Act and implementation guidance.

7. Classifieds.co.zw public site: broad categories, USD/ZiG display and free-ad functionality.
 8. Ownai public site: multi-category catalogue, cart, wish list and compare functions.
 9. Hypermarket public site: multi-vendor positioning, trusted-vendor claims and Paynow checkout.
 10. Shoppage Commerce Intelligence Foundation v0.4 and Implementation Progress Report dated 10 July 2026.
 11. Google Merchant Center local inventory documentation; YouTube Shopping affiliate documentation; TikTok Shop affiliate documentation, used as global pattern references rather than Zimbabwe availability claims.
 12. Cloudflare Stream, Mux and Supabase documentation, used as technical options subject to cost, rights and portability review.
-

Appendix E. Holistic review summary

The v0.7 model contained strong strategic material, but its append-only structure created avoidable investor and implementation risk. The review identified and corrected the following:

- repeated executive and final-decision statements;
- four overlapping implementation roadmaps;
- repeated AI, investor, risk, security and architecture sections;
- unqualified “national”, “complete” and “fully functional” language;
- possible confusion between global product references and Zimbabwe offers;
- launch-wedge mismatch with the food-product seed;
- unclear early ad-market cold-start economics;
- payment and fulfilment ownership ambiguity in earlier sections;
- merchant-site and Pemosites overlap;
- platform-ecosystem ambition that could open too early;
- 80 aims and 60 opportunities without sufficient horizon separation;
- accessibility metadata gaps in the original document;
- excessive direct formatting and non-semantic headings;
- raw URL citation labels;
- need to date and qualify current statistics and competitor features.

The v0.8 consolidation resolved these issues. The v0.9 adversarial pass made platform independence testable, simplified external brand language, staged creator compensation realistically and removed endorsement-dependent wording. The v1.0 pass locked the preloaded-intelligence launch, merchant claiming, historical archive, local recommendation model and mission-critical execution doctrine. The full block-level audit register is maintained separately for traceability.

Appendix F. Mandatory pre-public-launch gates

The four pre-gates are pass/fail controls, not presentation milestones. Each requires an evidence pack, accountable signatory and dated decision.

F.1 Source rights and deletion propagation

Required evidence:

- completed source register for 100% of public launch records;
- licence/contract/legal basis attached or referenced;
- public display, search, derived analytics and AI permissions separated;
- deletion/takedown path tested;
- one full source quarantine drill completed;
- counsel opinion and unresolved-risk list.

F.2 Merchant activation economics

Required evidence:

- first 30 merchants qualified or signed;
- cost and labour per activated merchant;
- 14-day freshness and response behaviour;
- willingness-to-pay evidence;
- six-month payback path or approved alternative threshold;
- operational plan for field and remote activation.

F.3 Native-demand superiority

Required evidence:

- at least 30 representative buyer needs;
- comparable incumbent social workflow baseline;
- time, relevance, locality, freshness and effort measures;
- at least 70% resolution and a material buyer-outcome advantage;
- documented failure analysis and corrections.

F.4 Independent security, privacy and operating readiness

Required evidence:

- independent review completed;
- critical/high findings resolved or formally accepted with controls;
- incident, complaint and manual fallback playbooks;
- backup, rollback and access-control tests;
- privacy/public-profile notices and agreements ready.

Appendix G. Source-rights register specification

The operational register is maintained in the accompanying workbook. Minimum fields are:

Field group	Required fields
Identity	Source ID, owner, dataset, version, hash, URL/reference
Acquisition	Method, date, source class, collector, automation/manual
Rights	Licence/contract, evidence link, public display, indexing, derived claims, analytics, redistribution, AI training, media use
Personal data	Classification, lawful basis, sensitive/minor flags, jurisdiction
Obligations	Attribution, share-alike, field restrictions, retention, expiry
Operations	Refresh cadence, correction channel, takedown SLA, deletion-propagation plan
Governance	Owner, legal review, status, next review date, residual risk

No blank permission is interpreted as permission. The default state is **blocked** until approved.

Appendix H. Unclaimed-profile policy and public copy

H.1 Public banner

Unclaimed reference profile. This page was assembled from the dated public or licensed sources shown below and is not currently managed by the business. Information may be incomplete or outdated. **Claim this business** or **request a correction/removal**.

H.2 Display restrictions

- no active seller, verified merchant, stock, price, payment or fulfilment claims without current evidence;
- no merchant-created marketing language attributed to the business;
- no paid ads or lead charges against an unclaimed profile;
- personal mobile numbers and residential addresses receive heightened review;
- disputed material fields are suppressed within the 24-hour action SLA.

H.3 SLA

- immediate acknowledgement and case number;
- within 24 hours: correct or temporarily suppress materially harmful/identity-sensitive disputed fields;
- within three business days: ordinary evidence-based correction target;
- complex ownership disputes: status update every three business days;
- appeal and audit history retained.

Appendix I. Coverage-cell and native-demand test standard

A coverage cell is **category x geography x fulfilment mode**.

State	Minimum conditions	Commercial treatment
Seeded	Reference taxonomy/profile data only	Reference Library and Request Network; no ads
Usable	3 claimed merchants, 50 fresh offers, >60% response within agreed window	Public category discovery; limited managed campaigns
Strong	10 claimed merchants, 250 fresh offers, >75% response, 30 resolved needs	Self-service ads and stronger recommendation surfaces
Defensible	Repeat direct demand, merchant retention, price/stock history and creator/market distribution	Expansion and investor proof

No cell may be called “covered” from reference records alone.

Appendix J. Independence kill-switch acceptance tests

Minimum staging tests:

1. Disable Facebook, WhatsApp, TikTok and Google adapters through feature flags.
2. Create/recover a buyer account using Shoppage-owned identity.
3. Claim an organisation and complete ordinary ownership review.
4. Publish a product and fresh local offer.
5. Search and filter local offers.
6. Submit a buyer request and route it to eligible merchants.
7. Record merchant response and buyer contact/lead.
8. Correct/suppress an unclaimed profile.
9. Export a social card and canonical link without API dependency.

10. Queue an offline merchant draft and sync safely.
11. Export merchant-managed data.
12. Re-enable adapters without duplicate events or data corruption.

Any failure of steps 2-8 blocks public launch.

Appendix K. Creator-bounty operating procedure

K.1 Launch compensation

- fixed fee per accepted deliverable;
- written brief, merchant consent and disclosure copy;
- canonical Shoppage link required;
- content remains live for the contracted period;
- verification after 48 hours;
- payout through an approved rail after acceptance;
- fraud, duplicate and complaint checks.

K.2 Minimum manual evidence before commission code

- at least 20 accepted/payed bounties;
- verified click and lead instrumentation stable;
- duplicate-phone and self-referral controls tested;
- merchant outcome marking quality measured;
- dispute and payout procedures completed.

Appendix L. v1.1 engineering backlog and kill list

L.1 P0 backlog

- source-rights service and source partitioning;
- deletion propagation/quarantine;
- unclaimed-profile policy UI and SLA queue;
- claim/correction/suppression/duplicate workflows;
- local-offer freshness auto-hide and ad eligibility;
- coverage-cell computation;
- Request Network and routing;
- social export-card generator;
- independence feature flags and acceptance suite;
- rules/hybrid recommendation and evaluation harness;
- board metric separation for reference, claimed, activated and fresh supply.

L.2 Kill list for months 0-6

- Pemosites and bespoke websites;
- public plugin marketplace or unrestricted APIs;
- owned fleet, warehouse or wallet;
- creator outcome commissions before gates;
- behavioural personalisation before evidence threshold;
- default buyer search dominated by reference-only products;
- unauthorised social scraping;
- field verification claims across all preloaded profiles;
- multi-tier subscription complexity;
- city/category expansion before strong coverage cells.

Appendix M. v1.1 review disposition

The adversarial review is substantially accepted. Its central correction is that preloaded scale is valuable only when rights, public-display language, buyer surfaces and deletion controls convert potential liability into controlled intelligence. The following refinements were made:

- the requested “24-hour correction SLA” is implemented as a 24-hour action/suppression requirement with a three-business-day ordinary correction target;
 - AI is not removed from launch, but bounded to rules/hybrid recommendations with a 500-query evaluation set and a 5,000-resolved-needs learning gate;
 - the one-million reference corpus remains a strategic asset but is separated from default buyer commerce search;
 - Pemosites is deferred beyond month six;
 - pricing is simplified to two launch plans rather than assumed to be permanently fixed at US\$19.
-

End of Shoppage National Commerce Model - National Launch Master v1.1